

 0xflotus fix: small error (#770)cde0294 · 4 years ago 

 .github/workflows	Fix broken links (#762)	6 years ago
 admin	Update authors (and separate roles into groups fo...	9 years ago
 figures	Update diagram for data transfer (#485)	9 years ago
 translations	Russian Translation (#1) (#744)	6 years ago
 .travis.yml	Whitelist swagger.io due to bug	8 years ago
 AUTHORS.md	Update credits.	8 years ago
 CODE_OF_CONDUCT.md	Create CODE_OF_CONDUCT.md	9 years ago
 CONTRIBUTING.md	Fix signup link.	10 years ago
 LICENSE.txt	Added cc-by-4.0 license, issue #535	9 years ago
 README.md	fix: small error (#770)	4 years ago



# The Open Guide to Amazon Web Services

---

[Chat](#) [Slack](#) ⇔ Join us!

[Credits](#) · [Contributing guidelines](#)

## Table of Contents

---

### Purpose

- [Why an Open Guide?](#)
- [Scope](#)
- [Legend](#)

### AWS in General

- [General Information](#)
- [Learning and Career Development](#)
- [Managing AWS](#)

- Managing Servers and Applications

Specific AWS Services	Basics	Tips	Gotchas
<a href="#">ALB</a>			
<a href="#">AMIs</a>			
<a href="#">API Gateway</a>			
<a href="#">Auto Scaling</a>			
<a href="#">Batch</a>			
<a href="#">Certificate Manager</a>			
<a href="#">CLB (ELB)</a>			
<a href="#">CloudFront</a>			
<a href="#">CloudFormation</a>			
<a href="#">CloudWatch</a>			
<a href="#">Device Farm</a>			
<a href="#">DirectConnect</a>			
<a href="#">DynamoDB</a>			
<a href="#">EBS</a>			
<a href="#">EC2</a>			
<a href="#">ECS</a>			
<a href="#">EKS</a>			
<a href="#">EFS</a>			

Specific AWS Services	Basics	Tips	Gotchas
<a href="#">Elastic Beanstalk</a>			
<a href="#">Elastic IPs</a>			
<a href="#">ElastiCache</a>			
<a href="#">EMR</a>			
<a href="#">Fargate</a>			
<a href="#">Glacier</a>			
<a href="#">IoT</a>			
<a href="#">Kinesis Firehose</a>			
<a href="#">Kinesis Streams</a>			
<a href="#">KMS</a>			
<a href="#">Lambda</a>			
<a href="#">Load Balancers</a>			
<a href="#">Mobile Hub</a>			
<a href="#">OpsWorks</a>			
<a href="#">Quicksight</a>			
<a href="#">RDS</a>			
<a href="#">RDS Aurora</a>			
<a href="#">RDS Aurora MySQL</a>			
<a href="#">RDS Aurora PostgreSQL</a>			

Specific AWS Services	Basics	Tips	Gotchas
<a href="#">RDS MySQL and MariaDB</a>			
<a href="#">RDS PostgreSQL</a>			
<a href="#">RDS SQL Server</a>			
<a href="#">Redshift</a>			
<a href="#">Route 53</a>			
<a href="#">S3</a>			
<a href="#">Security and IAM</a>			
<a href="#">SES</a>			
<a href="#">SNS</a>			
<a href="#">SQS</a>			
<a href="#">Step Functions</a>			
<a href="#">WAF</a>			
<a href="#">VPCs, Network Security, and Security Groups</a>			

## Special Topics

- [High Availability](#)
- [Billing and Cost Management](#)
- [Further Reading](#)

## Legal

- [Disclaimer](#)
- [License](#)



- [Figure: Tools and Services Market Landscape](#): A selection of third-party companies/products
- [Figure: AWS Data Transfer Costs](#): Visual overview of data transfer costs
- [Table: Service Matrix](#): How AWS services compare to alternatives
- [Table: AWS Product Maturity and Releases](#): AWS product releases
- [Table: Storage Durability, Availability, and Price](#): A quantitative comparison

## Why an Open Guide?

A lot of information on AWS is already written. Most people learn AWS by reading a blog or a "[getting started guide](#)" or [standard AWS references](#). Nonetheless, trustworthy and practical information and recommendations aren't easy to come by. Documentation is a great but sprawling resource few have time to read fully, and it doesn't include anything but official experiences of engineers. The information in blogs or [Stack Overflow](#) is also not consistently up to date.

This guide is by and for engineers who use AWS. It aims to be a useful, living reference that consolidates links, tips, guides, and more. It arose from discussion and editing over beers by [several engineers](#) who have used AWS extensively.

Before using the guide, please read the [license](#) and [disclaimer](#).

[Back to top ↑](#)

Please help!

omissions or errors.

Chat 

Please help by [joining the Slack channel](#) (we like to talk about AWS in general, even if you only have questions — discussion, community and guides improvements) and [contributing to the guide](#). This guide is *open to contributions*, so unlike a traditional book, we are actively improving. Like any open source effort, we combine efforts but also review to ensure high quality.

## Scope

- Currently, this guide covers selected “core” services, such as EC2, S3, Load Balancers, EBS, and IAM, and partially covers other services. We expect it to expand.
- It is not a tutorial, but rather a collection of information you can read and return to. It is for both beginners and experts.
- The goal of this guide is to be:
  - **Brief:** Keep it dense and use links
  - **Practical:** Basic facts, concrete details, advice, gotchas, and other “folk knowledge”
  - **Current:** We can keep updating it, and anyone can contribute improvements
  - **Thoughtful:** The goal is to be helpful rather than present dry facts. Thoughtful opinion with rationale is welcome. Opinions based on real experience can be extremely valuable. (We believe this is both possible with a good writing style and in some [other venues](#).)
- This guide is not sponsored by AWS or AWS-affiliated vendors. It is written by and for engineers who use AWS.

## Legend

-  Marks standard/official AWS pages and docs
-  Important or often overlooked tip
-  “Serious” gotcha (used where risks or time or resource costs are significant: critical security risks, mistakes with security, or poor architectural choices that are fundamentally difficult to correct)
-  “Regular” gotcha, limitation, or quirk (used where consequences are things not working, breaking, or not scaling)

- Undocumented feature (folklore)
- Relatively new (and perhaps immature) services or features
- Performance discussions
- Lock-in: Products or decisions that are likely to tie you to AWS in a new or significant way — that is, later moving to an alternative would be costly in terms of engineering effort
- Alternative non-AWS options
- Cost issues, discussion, and gotchas
- A mild warning attached to “full solution” or opinionated frameworks that may take significant time to understand your needs exactly; the opposite of a point solution (the cathedral is a nod to [Raymond’s metaphor](#))
- Colors indicate basics, tips, and gotchas, respectively.
- Areas where correction or improvement are needed (possibly with link to an issue — do help!)

## General Information

---

### When to Use AWS

- [AWS](#) is the dominant public cloud computing provider.
  - In general, “[cloud computing](#)” can refer to one of three types of cloud: “public,” “private,” and “hybrid.” AWS is public since anyone can use it. Private clouds are within a single (usually large) organization. Many companies use hybrid clouds.
  - The core features of AWS are [infrastructure-as-a-service](#) (IaaS) — that is, virtual machines and supporting infrastructure. Service models include [platform-as-a-service](#) (PaaS), which typically are more fully managed services that don’t require much configuration, or [software-as-a-service](#) (SaaS), which are cloud-based applications. AWS does offer a few products in other models, too.
  - In business terms, with infrastructure-as-a-service you have a variable cost model — it is [OpEx, not CapEx](#) (though [purchased contracts](#) are still CapEx).
- AWS’s TTM revenue was [\\$37.549 billion](#) as of Q1 2020 according to their earnings results (slide 14 in the linked Amazon.com’s total revenue (slide 11 in the same deck) for the same TTM period).
- **Main reasons to use AWS:**

- If your company is building systems or products that may need to scale
  - and you have technical know-how
  - and you want the most flexible tools
  - and you're not significantly tied into different infrastructure already
  - and you don't have internal, regulatory, or compliance reasons you can't use a public cloud-based solution
  - and you're not on a Microsoft-first tech stack
  - and you don't have a specific reason to use Google Cloud
  - and you can afford, manage, or negotiate its somewhat higher costs
  - ... then AWS is likely a good option for your company.
- Each of those reasons above might point to situations where other services are preferable. In practice, many, if not all, of these reasons apply to AWS. AWS has become the dominant IaaS provider, and it's well known and used by a number of modern large companies can or already do benefit from using AWS. Many large enterprises have moved their internal infrastructure to Azure, Google Cloud, and AWS.
- **Costs:** Billing and cost management are such big topics that we have [an entire section on this](#).
- **◆ EC2 vs. other services:** Most users of AWS are most familiar with [EC2](#), AWS' flagship virtual server product, and often times they are the first AWS service that comes to mind. But AWS products now extend far beyond basic IaaS, and often companies do not properly understand the breadth of the AWS services and how they can be applied, due to the [sharply growing](#) number of services, their novel branding, and fear of  lock-in to proprietary AWS technology. Although a bit daunting, it's important for decision makers in companies to understand the breadth of the AWS services and make informed decisions. (We hope this post helps!)
- **■ AWS vs. other cloud providers:** While AWS is the dominant IaaS provider (31% market share in [this 2016 estimate](#)), there are other cloud providers that are better suited to some companies. [This Gartner report](#) has a good overview of the market:
  - **Google Cloud Platform.** GCP arrived later to market than AWS, but has vast resources and is now used widely, including a few large ones. It is gaining market share. Not all AWS services have similar or analogous services on GCP. In particular, GCP offers some more advanced machine learning-based services like the [Vision](#), [Speech](#), and [Natural Language](#) services, which are not common to switch once you're up and running, but it does happen: [Spotify migrated](#) from AWS to Google Cloud. There is a great discussion [on Quora](#) about relative benefits. Of particular note is that VPCs in GCP are [global by default](#) with a single global IP range, while AWS' VPCs have to live within a particular region. This gives GCP an edge if you're designing applications from the beginning. It's also possible to [share one GCP VPC](#) between multiple projects (roughly analogous to AWS Lambda functions). AWS you'd have to peer them. It's also possible to [peer GCP VPCs](#) in a similar manner to how it's done in AWS.

- [Microsoft Azure](#) is the de facto choice for companies and teams that are focused on a Microsoft stack, and significant emphasis on Linux as well
  - In China, AWS' footprint is relatively small. The market is dominated by Alibaba's [Alibaba Cloud](#), formerly called Aliyun.
  - Companies at (very) large scale may want to reduce costs by managing their own infrastructure. For example, Google and Facebook manage their own infrastructure.
  - Other cloud providers such as [Digital Ocean](#) offer similar services, sometimes with greater ease of use, more automation, and lower cost. However, none of these match the breadth of products, mind-share, and market domination AWS has.
  - Traditional managed hosting providers such as [Rackspace](#) offer cloud solutions as well.
- **AWS vs. PaaS:** If your goal is just to put up a single service that does something relatively simple, and you're not interested in managing operations engineering, consider a [platform-as-a-service](#) such as [Heroku](#). The AWS approach to PaaS is arguably more complex, especially for simple use cases.
  - **AWS vs. web hosting:** If your main goal is to host a website or blog, and you don't expect to be building an application on top of it, you may wish consider one of the myriad [web hosting services](#).
  - **AWS vs. managed hosting:** Traditionally, many companies pay [managed hosting](#) providers to maintain physical hardware and build and deploy their software on top of the rented hardware. This makes sense for businesses who want direct control over their infrastructure due to legacy, performance, or special compliance constraints, but is usually considered old fashioned or unnecessary for most startup and younger tech companies.
  - **Complexity:** AWS will let you build and scale systems to the size of the largest companies, but the complexity of doing so at large scale requires significant depth of knowledge and experience. Even very simple use cases often require more knowledge about AWS than in a simpler environment like Heroku or Digital Ocean. (This guide may help!)
  - **Geographic locations:** AWS has data centers in [over a dozen geographic locations](#), known as **regions**, in Europe, North America, South America, and now Australia and India. It also has many more **edge locations** globally for reduced latency
    - See the [current list](#) of regions and edge locations, including upcoming ones.
    - If your infrastructure needs to be in close physical proximity to another service for latency or throughput reasons (such as an ad exchange), viability of AWS may depend on the location.
  - **Lock-in:** As you use AWS, it's important to be aware when you are depending on AWS services that do not have equivalents in other clouds or on-premises
    - Lock-in may be completely fine for your company, or a significant risk. It's important from a business perspective to understand the dependencies explicitly, and consider the cost, operational, business continuity, and competitive risks of being tied to AWS.

and reliable vendor, many companies are comfortable with using AWS to its full extent. Others can tell stories about being trapped in the “[cloud jail](#)” when costs spiral.

- Generally, the more AWS services you use, the more lock-in you have to AWS — that is, the more engineering time and money it will take to change to other providers in the future.
- Basic services like virtual servers and standard databases are usually easy to migrate to other providers or open source equivalents. Load balancers and IAM are specific to AWS but have close equivalents from other providers. The key thing to remember is that engineers are architecting systems around specific AWS services that are not open source or relatively interoperable. Lambda, API Gateway, Kinesis, Redshift, and DynamoDB do not have substantially equivalent open source or free software equivalents, while EC2, RDS (MySQL or Postgres), EMR, and ElastiCache more or less do. (See more [below](#), we’ll discuss this in more detail in the next section.)
- **Combining AWS and other cloud providers:** Many customers combine AWS with other non-AWS services. For example, a company might store secure data in a managed hosting provider, while other systems are AWS. Or a company might only use AWS for certain parts of their system, while doing everything else. However small startups or projects starting fresh will typically stick to AWS or Google Cloud Platform.
- **Hybrid cloud:** In larger enterprises, it is common to have [hybrid deployments](#) encompassing private cloud or on-premises infrastructure, as well as AWS or Google Cloud Platform — or other enterprise cloud providers like [IBM/Bluemix](#), [Microsoft/Azure](#), [NetApp](#), or [EMC](#).
- **Major customers:** Who uses AWS and Google Cloud?
  - AWS’s [list of customers](#) includes large numbers of mainstream online properties and major brands, such as Expedia (moving to Google Cloud), Airbnb, Zynga, Comcast, Nokia, and Bristol-Myers Squibb.
  - Azure’s [list of customers](#) includes companies such as NBC Universal, 3M and Honeywell Inc.
  - Google Cloud’s [list of customers](#) is large as well, and includes a few mainstream sites, such as [Snapchat](#), [Bespoke](#), and [Spotify](#).

[Back to top](#) 

## Which Services to Use

- AWS offers a *lot* of different services — [about a hundred](#) at last count.
- Most customers use a few services heavily, a few services lightly, and the rest not at all. What services you’ll use depend on your company’s needs and goals. Choices differ substantially from company to company.

- **Immature and unpopular services:** Just because AWS has a service that sounds promising, it doesn't mean you services are very narrow in use case, not mature, are overly opinionated, or have limitations, so building your own. We try to give a sense for this by breaking products into categories.
- **Must-know infrastructure:** Most typical small to medium-size users will focus on the following services first. If you systems, you likely need to know at least a little about all of these. (Even if you don't use them, you should learn choice intelligently.)
  - [IAM](#): User accounts and identities (you need to think about accounts early on!)
  - [EC2](#): Virtual servers and associated components, including:
    - [AMIs](#): Machine Images
    - [Load Balancers](#): CLBs and ALBs
    - [Autoscaling](#): Capacity scaling (adding and removing servers based on load)
    - [EBS](#): Network-attached disks
    - [Elastic IPs](#): Assigned IP addresses
  - [S3](#): Storage of files
  - [Route 53](#): DNS and domain registration
  - [VPC](#): Virtual networking, network security, and co-location; you automatically use
  - [CloudFront](#): CDN for hosting content
  - [CloudWatch](#): Alerts, paging, monitoring
- **Managed services:** Existing software solutions you could run on your own, but with managed deployment:
  - [RDS](#): Managed relational databases (managed MySQL, Postgres, and Amazon's own Aurora database)
  - [EMR](#): Managed Hadoop
  - [Elasticsearch](#): Managed Elasticsearch
  - [ElastiCache](#): Managed Redis and Memcached
- **Optional but important infrastructure:** These are key and useful infrastructure components that are less widely have legitimate reasons to prefer alternatives, so evaluate with care to be sure they fit your needs:
  -  [Lambda](#): Running small, fully managed tasks "serverless"
  - [CloudTrail](#): AWS API logging and audit (often neglected but important)
  -  [CloudFormation](#): Templated configuration of collections of AWS resources

-  [Elastic Beanstalk](#): Fully managed (PaaS) deployment of packaged Java, .NET, PHP, Node.js, Python, Ruby, applications
  -  [EFS](#): Network filesystem compatible with NFSv4.1
  -  [ECS](#): Docker container/cluster management (note Docker can also be used directly, without ECS)
  -  [EKS](#): Kubernetes (K8) Docker Container/Cluster management
  -  [ECR](#): Hosted private Docker registry
  -  [Config](#): AWS configuration inventory, history, change notifications
  -  [X-Ray](#): Trace analysis and debugging for distributed applications such as microservices.
- **Special-purpose infrastructure:** These services are focused on specific use cases and should be evaluated if they Many also are proprietary architectures, so tend to tie you to AWS.
    -  [DynamoDB](#): Low-latency NoSQL key-value store
    -  [Glacier](#): Slow and cheap alternative to S3
    -  [Kinesis](#): Streaming (distributed log) service
    -  [SQS](#): Message queueing service
    -  [Redshift](#): Data warehouse
    -  [QuickSight](#): Business intelligence service
    - [SES](#): Send and receive e-mail for marketing or transactions
    -  [API Gateway](#): Proxy, manage, and secure API calls
    -  [IoT](#): Manage bidirectional communication over HTTP, WebSockets, and MQTT between AWS and clients "things" like appliances or sensors)
    -  [WAF](#): Web firewall for CloudFront to deflect attacks
    -  [KMS](#): Store and manage encryption keys securely
    - [Inspector](#): Security audit
    - [Trusted Advisor](#): Automated tips on reducing cost or making improvements
    -  [Certificate Manager](#): Manage SSL/TLS certificates for AWS services
    -  [Fargate](#): Docker containers management, backend for ECS and EKS
  - **Compound services:** These are similarly specific, but are full-blown services that tackle complex problems and m depends on your requirements. If you have large or significant need, you may have these already managed by i

engineering teams.

- [Machine Learning](#): Machine learning model training and classification
- [Lex](#): Automatic speech recognition (ASR) and natural language understanding (NLU)
- [Polly](#): Text-to-speech engine in the cloud
- [Rekognition](#): Service for image recognition
-  [Data Pipeline](#): Managed ETL service
-  [SWF](#): Managed state tracker for distributed polyglot job workflow
-  [Lumberyard](#): 3D game engine
- **Mobile/app development:**
  - [SNS](#): Manage app push notifications and other end-user notifications
  -  [Cognito](#): User authentication via Facebook, Twitter, etc.
  - [Device Farm](#): Cloud-based device testing
  - [Mobile Analytics](#): Analytics solution for app usage
  -  [Mobile Hub](#): Comprehensive, managed mobile app framework
- **Enterprise services:** These are relevant if you have significant corporate cloud-based or hybrid needs. Many small startups use other solutions, like Google Apps or Box. Larger companies may also have their own non-AWS IT systems.
  - [AppStream](#): Windows apps in the cloud, with access from many devices
  - [Workspaces](#): Windows desktop in the cloud, with access from many devices
  - [WorkDocs](#) (formerly Zocalo): Enterprise document sharing
  - [WorkMail](#): Enterprise managed e-mail and calendaring service
  - [Directory Service](#): Microsoft Active Directory in the cloud
  - [Direct Connect](#): Dedicated network connection between office or data center and AWS
  - [Storage Gateway](#): Bridge between on-premises IT and cloud storage
  - [Service Catalog](#): IT service approval and compliance
- **Probably-don't-need-to-know services:** Bottom line, our informal polling indicates these services are just not being used for good reasons:
  - [Snowball](#): If you want to ship petabytes of data into or out of Amazon using a physical appliance, read on.
  - [Snowmobile](#): Appliances are great, but if you've got exabyte scale data to get into Amazon, nothing beats a truck.

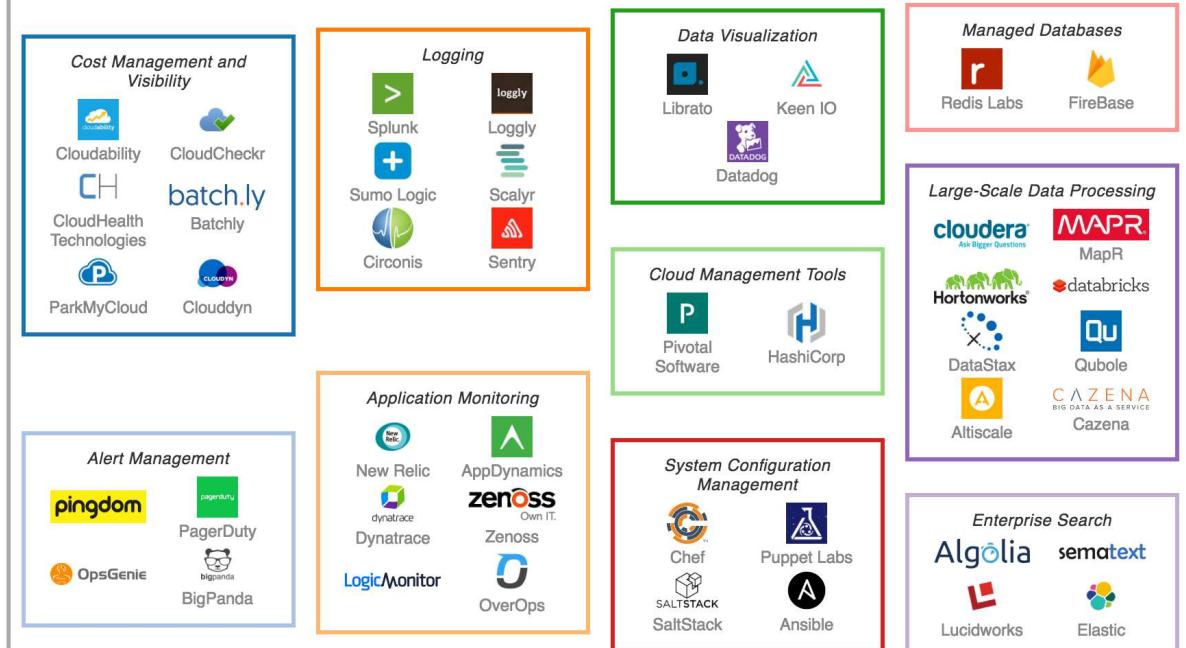
- [CodeCommit](#): Git service. You're probably already using GitHub or your own solution ([Stackshare](#) has information).
  -  [CodePipeline](#): Continuous integration. You likely have another solution already.
  -  [CodeDeploy](#): Deployment of code to EC2 servers. Again, you likely have another solution.
  -  [OpsWorks](#): Management of your deployments using Chef or (as of November 2017) Puppet Enterprise.
- [AWS in Plain English](#) offers more friendly explanation of what all the other different services are.

[Back to top](#) 

## Tools and Services Market Landscape

There are now enough cloud and "big data" enterprise companies and products that few can keep up with the market.

We've assembled a landscape of a few of the services. This is far from complete, but tries to emphasize services that practitioners — services that specifically help with AWS, or a complementary, or tools almost anyone using AWS must



Credits and latest version: <https://github.com/open-guides/og-aws>

🚧 Suggestions to improve this figure? Please [file an issue](#).

[Back to top ↑](#)

## Common Concepts

- 💡 The AWS [General Reference](#) covers a bunch of common concepts that are relevant for multiple services.
- AWS allows deployments in [regions](#), which are isolated geographic locations that help you reduce latency or offload. Regions contain availability zones(AZs), which are typically the first tool of choice for [high availability](#). AZs are [providing redundancy](#) across different physical locations.

[another](#) even within the same region, and [may span multiple physical data centers](#). While they are connected via disasters afflicting one should not affect others.

- Each service has API [endpoints](#) for each region. Endpoints differ from service to service and not all services are listed in [these tables](#).
- [Amazon Resource Names \(ARNs\)](#) are specially formatted identifiers for identifying resources. They start with 'arn:' services, and in particular for IAM policies.

[Back to top](#) 

## Service Matrix

Many services within AWS can at least be compared with Google Cloud offerings or with internal Google services. And assemble the same thing yourself with open source software. This table is an effort at listing these rough correspondences. This table is imperfect as in almost every case there are subtle differences of features!)

Service	AWS	Google Cloud	Google Internal	Microsoft Azure	Other providers	Open "but..."
Virtual server	EC2	Compute Engine (GCE)		Virtual Machine	DigitalOcean	OpenShift
PaaS	Elastic Beanstalk	App Engine	App Engine	Web Apps	Heroku, AppFog, OpenShift	Metaphysic, AppScale, CloudBees, FourQuint, Convoy
Serverless, microservices	Lambda, API Gateway	Functions		Function Apps	PubNub Blocks, Auth0 Webtask	Kong
Container, cluster	ECS, EKS, Fargate	Container Engine,	Borg or Omega	Container Service		Kubernetes, Mesos

Service	AWS	Google Cloud	Google Internal	Microsoft Azure	Other providers	Open "but"
manager		Kubernetes				
Object storage	S3	Cloud Storage	GFS	Storage Account	DigitalOcean Spaces	Swift, Minio
Block storage	EBS	Persistent Disk		Storage Account	DigitalOcean Volumes	NFS
SQL datastore	RDS	Cloud SQL		SQL Database		MySQL, PostgreSQL
Sharded RDBMS		Cloud Spanner	F1, Spanner	Azure Database for PostgreSQL - Hyperscale (Citus)		Cassandra, CockroachDB
Bigtable		Cloud Bigtable	Bigtable			HBase
Key-value store, column store	DynamoDB	Cloud Datastore	Megastore	Tables, DocumentDB		Cassandra, Couchbase, Redis, Redis
Memory cache	ElastiCache	App Engine Memcache		Redis Cache		Memcached, Redis
Search	CloudSearch, Elasticsearch			Search	Algolia, QBox, Elastic Cloud	Elasticsearch, Solr

Service	AWS	Google Cloud	Google Internal	Microsoft Azure	Other providers	Open "big data"
	(managed)					
Data warehouse	Redshift	BigQuery	Dremel	SQL Data Warehouse	Oracle, IBM, SAP, HP, many others	Gree
Business intelligence	QuickSight	Data Studio 360		Power BI	Tableau	
Lock manager	<a href="#">DynamoDB (weak)</a>		Chubby	Lease blobs in Storage Account		Zookeeper, Etcd,...
Message broker	SQS, SNS, IoT	Pub/Sub	PubSub2	Service Bus		RabbitMQ, Kafka
Streaming, distributed log	Kinesis	Dataflow	PubSub2	Event Hubs		Kafka, Apache Apex, Spark Streaming, Storm
MapReduce	EMR	Dataproc	MapReduce	HDInsight, DataLake Analytics	Qubole	Hadoop
Monitoring	CloudWatch	Stackdriver Monitoring	Borgmon	Monitor		Prometheus
Tracing	X-Ray	Stackdriver Trace		Monitor (Application)	DataDog, New Relic, Epsagon	Zipkin, AppNeta

Service	AWS	Google Cloud	Google Internal	Microsoft Azure	Other providers	Open "b..."
				Insights)		
Metric management			Borgmon, TSDB	Application Insights		Graphite InfluxDB OpenMetrics Grafana Riemann Pronosticator
CDN	CloudFront	Cloud CDN		CDN	Akamai, Fastly, Cloudflare, Limelight Networks	Apache Server
Load balancer	CLB/ALB	Load Balancing	GFE	Load Balancer, Application Gateway		Nginx HAProxy Apache Server
DNS	Route53	DNS		DNS		bind
Email	SES				Sendgrid, Mandrill, Postmark	
Git hosting	CodeCommit	Cloud Source Repositories		Visual Studio Team Services	GitHub, BitBucket	GitLab
User authentication	Cognito	Firebase Authentication		Azure Active Directory		oauth2

Service	AWS	Google Cloud	Google Internal	Microsoft Azure	Other providers	Op “bu
Mobile app analytics	Mobile Analytics	Firebase Analytics		HockeyApp	Mixpanel	
Mobile app testing	Device Farm	Firebase Test Lab		Xamarin Test Cloud	BrowserStack, Sauce Labs, Testdroid	
Managing SSL/TLS certificates	Certificate Manager				Let's Encrypt, Comodo, Symantec, GlobalSign	
Automatic speech recognition and natural language understanding	Transcribe (ASR), Lex (NLU)	Cloud Speech API, Natural Language API		Cognitive services	AYLIEN Text Analysis API, Ambiverse Natural Language Understanding API	Stand Core Suite Open Apa spaC
Text-to-speech engine in the cloud	Polly				Nuance, Vocalware, IBM	Mim Mary
Image recognition	Rekognition	Vision API		Cognitive services	IBM Watson, Clarifai	Tens Oper
OCR (Text recognition)	Texttract (documents),	Cloud Vision API		Computer Vision API		Tesse

Service	AWS	Google Cloud	Google Internal	Microsoft Azure	Other providers	Op “bu
	Rekognition (photographs)					
Language Translation	Translate	Translate		Translator Text API		Aper
File Share and Sync	WorkDocs	Google Docs		OneDrive	Dropbox, Box, Citrix File Share	own
Machine Learning	SageMaker, DeepLens, ML	ML Engine, Auto ML		ML Studio	Watson ML	
Data Loss Prevention	Macie	Cloud Data Loss Prevention		Azure Information Protection		

 [Please help fill this table in.](#)

Selected resources with more detail on this chart:

- Google internal: [MapReduce](#), [Bigtable](#), [Spanner](#), [F1 vs Spanner](#), [Bigtable vs Megastore](#)

[Back to top](#) 

## AWS Product Maturity and Releases

It's important to know the maturity of each AWS product. Here is a mostly complete list of first release date, with links. Most recently released services are first. Not all services are available in all regions; see [this table](#).

Service	Original release	Availability	CLI Support	HIPAA Compliant
 <a href="#">X-Ray</a>	2016-12	General	✓	✓

Service	Original release	Availability	CLI Support	HIPAA Compliant
 <a href="#">Lex</a>	2016-11	Preview		
 <a href="#">Polly</a>	2016-11	General	✓	✓
 <a href="#">Rekognition</a>	2016-11	General	✓	✓
 <a href="#">Athena</a>	2016-11	General	✓	✓
 <a href="#">Batch</a>	2016-11	General	✓	✓
 <a href="#">Database Migration Service</a>	2016-03	General		✓
 <a href="#">Certificate Manager</a>	2016-01	General	✓	✓
 <a href="#">IoT</a>	2015-08	General	✓	✓
 <a href="#">WAF</a>	2015-10	General	✓	✓
 <a href="#">Data Pipeline</a>	2015-10	General	✓	
 <a href="#">Elasticsearch</a>	2015-10	General	✓	✓
 <a href="#">Aurora</a>	2015-07	General	✓	✓ <sup>3</sup>
 <a href="#">Service Catalog</a>	2015-07	General	✓	✓
 <a href="#">Device Farm</a>	2015-07	General	✓	
 <a href="#">CodePipeline</a>	2015-07	General	✓	✓
 <a href="#">CodeCommit</a>	2015-07	General	✓	✓
 <a href="#">API Gateway</a>	2015-07	General	✓	✓ <sup>1</sup>
 <a href="#">Config</a>	2015-06	General	✓	✓
 <a href="#">EFS</a>	2015-05	General	✓	✓

Service	Original release	Availability	CLI Support	HIPAA Compliant
 <a href="#">Machine Learning</a>	2015-04	General	✓	
<a href="#">Lambda</a>	2014-11	General	✓	✓
<a href="#">ECS</a>	2014-11	General	✓	✓
<a href="#">EKS</a>	2018-06	General	✓ <sup>12</sup>	✓
<a href="#">KMS</a>	2014-11	General	✓	✓
<a href="#">CodeDeploy</a>	2014-11	General	✓	✓
<a href="#">Kinesis</a>	2013-12	General	✓	✓
<a href="#">CloudTrail</a>	2013-11	General	✓	✓
<a href="#">AppStream</a>	2013-11	Preview		✓
<a href="#">CloudHSM</a>	2013-03	General	✓	✓
<a href="#">Silk</a>	2013-03	Obsolete?		
<a href="#">OpsWorks</a>	2013-02	General	✓	✓
<a href="#">Redshift</a>	2013-02	General	✓	✓
<a href="#">Elastic Transcoder</a>	2013-01	General	✓	
<a href="#">Glacier</a>	2012-08	General	✓	✓
<a href="#">CloudSearch</a>	2012-04	General	✓	
<a href="#">SWF</a>	2012-02	General	✓	✓
<a href="#">Storage Gateway</a>	2012-01	General	✓	✓
<a href="#">DynamoDB</a>	2012-01	General	✓	✓

Service	Original release	Availability	CLI Support	HIPAA Compliant
<a href="#">DirectConnect</a>	2011-08	General	✓	✓
<a href="#">ElastiCache</a>	2011-08	General	✓	✓ <sup>14</sup>
<a href="#">CloudFormation</a>	2011-04	General	✓	✓
<a href="#">SES</a>	2011-01	General	✓	✓
<a href="#">Elastic Beanstalk</a>	2010-12	General	✓	✓
<a href="#">Route 53</a>	2010-10	General	✓	✓
<a href="#">IAM</a>	2010-09	General	✓	
<a href="#">SNS</a>	2010-04	General	✓	✓
<a href="#">EMR</a>	2010-04	General	✓	✓
<a href="#">RDS</a>	2009-12	General	✓	✓ <sup>2</sup>
<a href="#">VPC</a>	2009-08	General	✓	✓
<a href="#">Snowball</a>	2015-10	General	✓	✓
<a href="#">Snowmobile</a>	2016-11	General		✓
<a href="#">CloudWatch</a>	2009-05	General	✓	✓
<a href="#">CloudFront</a>	2008-11	General	✓	✓ <sup>4</sup>
<a href="#">Fulfillment Web Service</a>	2008-03	Obsolete?		
<a href="#">SimpleDB</a>	2007-12	! <a href="#">Nearly obsolete</a>	✓	
<a href="#">DevPay</a>	2007-12	General		
<a href="#">Flexible Payments Service</a>	2007-08	Retired		

Service	Original release	Availability	CLI Support	HIPAA Compliant
<a href="#">EC2</a>	2006-08	General	✓	✓ <sup>5,6,7</sup>
<a href="#">SQS</a>	2006-07	General	✓	✓
<a href="#">S3</a>	2006-03	General	✓	✓ <sup>8</sup>
<a href="#">Alexa Top Sites</a>	2006-01	General  HTTP-only		
<a href="#">Alexa Web Information Service</a>	2005-10	General  HTTP-only		

[Back to top](#)

#### Footnotes

1: Excludes use of Amazon API Gateway caching

2: RDS MySQL, Oracle, and PostgreSQL engines only

3: MySQL-compatible Aurora edition only

4: Excludes Lambda@Edge

5: Includes EC2 Systems Manager

6: Includes Elastic Block Storage (EBS)

7: Includes Elastic Load Balancing

8: Includes S3 Transfer Acceleration

9: Includes RDS MySQL, Oracle, PostgreSQL, SQL Server, and MariaDB

10: Includes Auto-Scaling

11: Data Analytics, Streams, Video Streams and Firehose

12: Kubernetes uses a custom CLI for Pod/Service management called kubectl. AWS CLI only handles Kubernetes Ma

13: IoT Core (includes Device Management) and Greengrass

14: ElastiCache for Redis only

15: Snowball and Snowball Edge

## Compliance

- Many applications have strict requirements around reliability, security, or data privacy. The [AWS Compliance](#) page lists several certifications, which include PCI DSS Level 1, SOC 1, 2, and 3, HIPAA, and ISO 9001.
- Security in the cloud is a complex topic, based on a [shared responsibility model](#), where some elements of compliance are the responsibility of AWS, and some are provided by your company.
- Several third-party vendors offer assistance with compliance, security, and auditing on AWS. If you have substantial needs, assistance is a good idea.
- From inside China, AWS services outside China [are generally accessible](#), though there are at times breakages in connectivity to services [inside China](#).

## Getting Help and Support

- **Forums:** For many problems, it's worth searching or asking for help in the [discussion forums](#) to see if it's a known issue.
- **Premium support:** AWS offers several levels of [premium support](#).
  - The first tier, called "Developer support" lets you file support tickets with 12 to 24 hour turnaround time, it starts at \$1000 per month. As monthly spend reaches around \$1000 it changes to a 3% surcharge on your bill.
  - The higher-level support services are quite expensive — and increase your bill by up to 10%. Many large companies never pay for this level of support. They are usually more helpful for midsize or larger companies needing rapid resolution of more perplexing problems.
  - Keep in mind, a flexible architecture can reduce need for support. You shouldn't be relying on AWS to solve all your problems. For example, if you can easily re-provision a new server, it may not be urgent to solve a rare kernel-level issue unless it's impacting critical services. If your EBS volumes have recent snapshots, you may be able to restore a volume before support can rectify the issue. If your services have an issue in one availability zone, you should in any case be able to rely on a redundant service in another zone.
  - Larger customers also get access to AWS Enterprise support, with dedicated technical account managers (TAMs) and 99.99% time SLAs.
  - There is definitely some controversy about how useful the paid support is. The support staff don't always seem to have the information and authority to solve the problems that are brought to their attention. Often your ability to have your issue resolved depends on your relationship with your account rep.
- **Account manager:** If you are at significant levels of spend (thousands of US dollars plus per month), you may be asked to request a dedicated account manager.

- These are a great resource, even if you're not paying for premium support. Build a good relationship with them for questions, problems, and guidance.
- Assign a single point of contact on your company's side, to avoid confusing or overwhelming them.
- **Contact:** The main web contact point for AWS is [here](#). Many technical requests can be made via these channels.
- **Consulting and managed services:** For more hands-on assistance, AWS has established relationships with many [managed service partners](#). The big consultants won't be cheap, but depending on your needs, may save you costs by helping you set up your architecture more effectively, or offering specific expertise, e.g. security. Managed service providers offer full-service management of cloud resources.
- **AWS Professional Services:** AWS provides [consulting services](#) alone or in combination with partners.

## Restrictions and Other Notes

- ◆ Lots of resources in Amazon have [limits](#) on them. This is actually helpful, so you don't incur large costs accidentally. It's important to know what limits apply to your usage so that quotas can be increased by opening support tickets. Some limits are easy to raise, and some are not. (Some of the most common are listed below.) Additionally, not all service limits are published.
  - **Obtaining Current Limits and Usage:** Limit information for a service may be available from the service API, documentation, or Support pages. [This page](#) from the awslimitchecker tool's documentation provides a summary of available retrieval options for each limit. The [tool](#) itself is also valuable for automating limit checks.
- ◆ [AWS terms of service](#) are extensive. Much is expected boilerplate, but it does contain important notes and restrictions. In particular, there are restrictions against using many AWS services in **safety-critical systems**. (Those appreciate it's important to review [clause 42.10](#).)

## Related Topics

- [OpenStack](#) is a private cloud alternative to AWS used by large companies that wish to avoid public cloud offerings.

## Learning and Career Development

---

### Certifications

- **Certifications:** AWS offers [certifications](#) for IT professionals who want to demonstrate their knowledge.

- [Certified Cloud Practitioner](#)
- [Certified Solutions Architect Associate](#)
- [Certified Developer Associate](#)
- [Certified SysOps Administrator Associate](#)
- [Certified Solutions Architect Professional](#)
- [Certified DevOps Engineer Professional](#)
- [Certified Security – Specialty](#)
- [Certified Advanced Networking – Specialty](#)
- [Certified Machine Learning – Specialty](#)
- [Certified Data Analytics – Specialty](#)
- [Certified Database – Specialty](#)

Associate level certifications were once required as pre-requisites to taking the Professional examinations - this is no longer the case.

- **Getting certified:** If you're interested in studying for and getting certifications, [this practical overview](#) tells you all you need to know. The official page is [here](#) and there is an [FAQ](#).
- **Training for certifications:** Training is offered by AWS themselves (mainly instructor-led and on-site) and various third parties (usually as video-based training) such as [A Cloud Guru](#), [CloudAcademy](#) and [Linux Academy](#).
- **Do you need a certification?** Especially in consulting companies or when working in key tech roles in large non-tech companies, certifications are important credentials. In others, including in many tech companies and startups, certifications are not always considered necessary. (In fact, fairly or not, some Silicon Valley hiring managers and engineers see them as a "nice-to-have" on a resume.)

Certifications are required to access certificate lounges at official AWS events such as [Summits](#) and [re:Invent](#). Lounges provide free food, charging points, seats and relatively better coffee.

## Managing AWS

---

### Managing Infrastructure State and Change

A great challenge in using AWS to build complex systems (and with DevOps in general) is to manage infrastructure state. In general, this boils down to three broad goals for the state of your infrastructure:

- **Visibility:** Do you know the state of your infrastructure (what services you are using, and exactly how)? Do you and anyone on your team — make changes? Can you detect misconfigurations, problems, and incidents with your service?
- **Automation:** Can you reconfigure your infrastructure to reproduce past configurations or scale up existing ones without manual work, or requiring knowledge that's only in someone's head? Can you respond to incidents easily or automatically?
- **Flexibility:** Can you improve your configurations and scale up in new ways without significant effort? Can you add new tools to your workflow? Can you reuse the same tools? Do you share, review, and improve your configurations within your team?

Much of what we discuss below is really about how to improve the answers to these questions.

There are several approaches to deploying infrastructure with AWS, from the console to complex automation tools, that attempt to help achieve visibility, automation, and flexibility.

## AWS Configuration Management

The first way most people experiment with AWS is via its web interface, the AWS Console. But using the Console is a good way to learn about AWS, but often works against automation or flexibility.

So if you're not going to manage your AWS configurations manually, what should you do? Sadly, there are no simple answers. Every approach has pros and cons, and the approaches taken by different companies vary widely, and include directly using the AWS API (without tooling on top yourself), using command-line tools, and using third-party tools and services.

## AWS Console

- The [AWS Console](#) lets you control much (but not all) functionality of AWS via a web interface.
- Ideally, you should only use the AWS Console in a few specific situations:
  - It's great for read-only usage. If you're trying to understand the state of your system, logging in and browsing the AWS console is a good way to do that.
  - It is also reasonably workable for very small systems and teams (for example, one engineer setting up one service).

- It can be useful for operations you're only going to do rarely, like less than once a month (for example, a one-time migration). In this case using the console can be the simplest approach.
- **! Think before you use the console:** The AWS Console is convenient, but also the enemy of automation, reproducibility, and communication. If you're likely to be making the same change multiple times, avoid the console. Favor some sort of automation over the console. There's almost always a path toward automation, as discussed next. Not only does using the console preclude automation, which makes it harder to document and standardize your processes, but it prevents documentation, clarity, and standardization around processes for yourself and your team.

## Command-Line tools

- The [aws command-line interface](#) (CLI), used via the `aws` command, is the most basic way to save and automate AWS operations.
- Don't underestimate its power. It also has the advantage of being well-maintained — it covers a large proportion of AWS services and is up to date.
- In general, whenever you can, prefer the command line to the AWS Console for performing operations.
- ◆ Even in the absence of fancier tools, you can [write simple Bash scripts](#) that invoke `aws` with specific arguments to perform common tasks. This is a primitive but effective way to document operations you've performed. It improves automation, allows others to reuse them on a team, and gives others a starting point for future work.
- ◆ For use that is primarily interactive (not scripted), consider instead using the [aws-shell](#) tool from AWS. It is easier to use than the CLI because it provides tab completion and a colorful UI, but still works on the command line. If you're using [SAWS](#), a previous version of the AWS CLI, [migrate to aws-shell](#).

## APIs and SDKs

- SDKs for using AWS APIs are available in most major languages, with [Go](#), [iOS](#), [Java](#), [JavaScript](#), [Python](#), [Ruby](#), and [.NET](#) among others. AWS maintains [a short list](#), but the [awesome-aws list](#) is the most comprehensive and current. Note [support](#) for each language and the [AWS Regions](#) supported by each.
- **Retry logic:** An important aspect to consider whenever using SDKs is error handling; under heavy use, a wide variety of programming errors to throttling to AWS-related outages or failures, can be expected to occur. SDKs typically implement [backoff](#) to address this, but this may need to be understood and adjusted over time for some applications. For example, [Amazon S3](#) has a [retry logic](#) that attempts to retransmit failed requests.
- **! Don't use APIs directly.** Although AWS documentation includes lots of API details, it's better to use the SDKs for your language to access APIs. SDKs are more mature, robust, and well-maintained than something you'd write yourself.

## Boto

- A good way to automate operations in a custom way is [Boto3](#), also known as the [Amazon SDK for Python](#). [Boto](#), this library, has been in wide use for years, but now there is a newer version with official support from Amazon, so projects.
- Boto3 contains a variety of APIs that operate at either a high level or a low level, here some explanation of both:
  - The low level APIs (Client APIs) are mapped to AWS Cloud service-specific APIs, and all service operations are Clients are generated from a JSON service definition file.
  - The high level option, Resource APIs, allows you to avoid calling the network at the low level and instead provide a way to interact with AWS Cloud services.
- Boto3 has a lot of helpful [features](#) like *waiters*, which provide a structure that allows for code to wait for changes. For example, when you are creating an EC2 instance and need wait until the instance is running in order to perform other operations.
- If you find yourself writing a Bash script with more than one or two CLI commands, you're probably doing it wrong. Instead, consider writing a Boto script instead. This has the advantages that you can:
  - Check return codes easily so success of each step depends on success of past steps.
  - Grab interesting bits of data from responses, like instance ids or DNS names.
  - Add useful environment information (for example, tag your instances with git revisions, or inject the latest build artifacts into an initialization script).

[Back to top](#) 

## General Visibility

- ◆ [Tagging resources](#) is an essential practice, especially as organizations grow, to better understand your resources. Through automation or convention, you can add tags:
  - For the org or developer that "owns" that resource
  - For the product that resource supports
  - To label lifecycles, such as temporary resources or one that should be deprovisioned in the future
  - To distinguish production-critical infrastructure (e.g. serving systems vs backend pipelines)
  - To distinguish resources with special security or compliance requirements

- To (once enabled) [allocate cost](#). Note that cost allocation tags only apply on a forward-looking basis; you can't apply them to items already billed.
- For many years, there was a notorious 10 tag limit per resource, which could not be raised and caused many headaches. As of 2016, this was [raised](#) to 50 tags per resource.
  - ◆ In 2017, AWS introduced the ability to [enforce tagging](#) on instance and volume creation, deprecating policies such as [Cloud Custodian](#).
  - ◆ Tags are case sensitive; 'environment' and 'Environment' are two different tags. Automation in setting tags is a sensible option at significant scale.
  - ◆ There is a bug in the ASG console where spaces after tag names are preserved. So if you type "Name" with a space, it will not get the expected behavior. This is probably true in other locations and SDKs also. Be sure you do not use spaces in tag names unless you really mean it. (As of Jul 2018)
  - ◆ When resources are shared across the org, tags are not shared with it. For example, sharing Transit Gateways between accounts creates correct tags in the account that created these resources but not in the accounts where these resources were shared.

## Managing Servers and Applications

---

[Back to top](#) 

### AWS vs Server Configuration

This guide is about AWS, not DevOps or server configuration management in general. But before getting into AWS infrastructure, it's important to understand that in addition to the configuration management for your AWS resources, there is the long-standing problem of configuration management for servers themselves.

[Back to top](#) 

### Philosophy

- Heroku's [Twelve-Factor App](#) principles list some established general best practices for deploying applications.
- **Pets vs cattle:** Treat servers [like cattle, not pets](#). That is, design systems so infrastructure is disposable. It should be easy to recreate a server if it is unexpectedly destroyed.

- The concept of [immutable infrastructure](#) is an extension of this idea.
- Minimize application state on EC2 instances. In general, instances should be able to be killed or die unexpectedly. State that is in your application should quickly move to RDS, S3, DynamoDB, EFS, or other data stores not on the instance itself. If you must store state on the instance, consider using ephemeral storage options, though it generally should not be the bootable volume, and EBS will require manual or automated re-mounting.

[Back to top](#) 

## Server Configuration Management

- There is a [large set](#) of open source tools for managing configuration of server instances.
- These are generally not dependent on any particular cloud infrastructure, and work with any variety of Linux (or non-Linux) operating systems.
- Leading configuration management tools are [Puppet](#), [Chef](#), [Ansible](#), and [Saltstack](#). These aren't the focus of this section, but we'll mention them as they relate to AWS.

[Back to top](#) 

## Containers and AWS

- [Docker](#) and the containerization trend are changing the way many servers and services are deployed in general.
- Containers are designed as a way to package up your application(s) and all of their dependencies in a known way. When you run a container, you are including every library or binary your application needs, outside of the kernel. A big advantage is that it's easy to test and validate a container locally without worrying about some difference between your computer and where it's deployed.
- A consequence of this is that you need fewer AMIs and boot scripts; for most deployments, the only boot script you need is one that fetches an exported docker image and runs it.
- Companies that are embracing [microservice architectures](#) will often turn to container-based deployments.
- AWS launched [ECS](#) as a service to manage clusters via Docker in late 2014, though many people still deploy Docker containers themselves. See the [ECS section](#) for more details.
- AWS launched [EKS](#) as a service to manage Kubernetes Clusters mid 2018, though many people still deploy ECS containers themselves. See the [EKS section](#) for more details.

## Visibility

- Store and track instance metadata (such as instance id, availability zone, etc.) and deployment info (application build number, commit hash, etc.) in your logs or reports. The [instance metadata service](#) can help collect some of the AWS data you'll need.
- **Use log management services:** Be sure to set up a way to view and manage logs externally from servers.
  - Cloud-based services such as [Sumo Logic](#), [Splunk Cloud](#), [Scalyr](#), [LogDNA](#), and [Loggly](#) are the easiest to set up and are also the most expensive, which may be a factor depending on how much log data you have).
  - Major open source alternatives include [Elasticsearch](#), [Logstash](#), and [Kibana](#) (the "[Elastic Stack](#)") and [Graylog](#).
  - If you can afford it (you have little data or lots of money) and don't have special needs, it makes sense to use a cloud-based service like Logstash whenever possible, since setting up your own scalable log processing systems is notoriously time consuming and error-prone.
- **Track and graph metrics:** The AWS Console can show you simple graphs from CloudWatch, but you typically will want to collect metrics from multiple sources and visualize them together. Collect and export helpful metrics everywhere you can (make sure they are manageable enough you can afford it).
  - Services like [Librato](#), [KeenIO](#), and [Datadog](#) have fancier features or better user interfaces that can save a lot of time (a good comparison is [here](#)).
  - Use [Prometheus](#) or [Graphite](#) as timeseries databases for your metrics (both are open source).
  - [Grafana](#) can visualize with dashboards the stored metrics of both timeseries databases (also open source).

## Tips for Managing Servers

- **! Timezone settings on servers:** unless *absolutely necessary*, always set the timezone on servers to [UTC](#) (see in-depth discussion on [AWS Lambda](#)). Most Linux distributions, such as [Ubuntu](#), [CentOS](#) or [Amazon Linux](#). Numerous distributed systems rely on time for synchronization, and UTC [provides](#) the universal reference plane: it is not subject to daylight savings changes and adjustments in different regions. This makes UTC a great choice for servers, as it saves you a lot of headache debugging [elusive timezone issues](#) and provide coherent timeline of events in your logs.
- **NTP and accurate time:** If you are not using Amazon Linux (which comes preconfigured), you should confirm you are using NTP correctly, to avoid insidious time drift (which can then cause all sorts of issues, from breaking API calls to misleading log timestamps).

part of your automatic configuration for every server. If time has already drifted substantially (generally >1000s), it won't shift it back, so you may need to remediate manually (for example, [like this](#) on Ubuntu).

- **Testing immutable infrastructure:** If you want to be proactive about testing your service's ability to cope with infrastructure failure, it can be helpful to introduce random instance termination during business hours, which will expose any bugs that are present. Engineers are available to identify and fix them. Netflix's [Simian Army](#) (specifically, [Chaos Monkey](#)) is a popular tool for this. BBC's [chaos-lambda](#) is a lightweight option which runs on AWS [Lambda](#).

## Security and IAM

---

We cover security basics first, since configuring user accounts is something you usually have to do early on when setting up a new AWS account.

### Security and IAM Basics

-  IAM [Homepage](#) · [User guide](#) · [FAQ](#)
- The [AWS Security Blog](#) is one of the best sources of news and information on AWS security.
- IAM is the service you use to manage accounts and permissioning for AWS.
- Managing security and access control with AWS is critical, so every AWS administrator needs to use and understand the service at a deep level.
- IAM identities include users (people or services that are using AWS), groups (containers for sets of users) and roles (containers for permissions assigned to AWS service instances). Permissions for these identities are governed by pre-defined policies or custom policies that you create.
- IAM manages various kinds of authentication, for both users and for software services that may need to authenticate to AWS.
  - [Passwords](#) to log into the console. These are a username and password for real users.
  - [Access keys](#), which you may use with command-line tools. These are two strings, one the "id", which is an unique identifier of the form 'XXXXXXXXXXXXXXXXXXXXXX', and the other is the secret, which is a 40-character mixed-case base64 string often set up for services, not just users.
    -  Access keys that start with AKIA are normal keys. Access keys that start with ASIA are session/temporary keys. They will require an additional "SessionToken" parameter to be sent along with the id and secret. See the [complete list of access key prefixes](#).

- [Multi-factor authentication \(MFA\)](#), which is the highly recommended practice of using a keychain fob or smart card as an additional layer of protection for user authentication.
- IAM allows complex and fine-grained control of permissions, dividing users into groups, assigning permissions to them via a [policy language](#) that can be used to customize security policies in a fine-grained way.
  - An excellent high level overview of IAM policy concepts lives at [IAM Policies In A Nutshell](#).
  - ◆ The policy language has a complex and error-prone JSON syntax that's quite confusing, so unless you are an expert, it's best to base yours off trusted examples or AWS' own pre-defined [managed policies](#).
- At the beginning, IAM policy may be very simple, but for large systems, it will grow in complexity, and need to be managed
  - ◆ Make sure one person (perhaps with a backup) in your organization is formally assigned ownership of managing IAM policies. Make sure every administrator works with that person to have changes reviewed. This goes a long way to avoiding serious misconfigurations.
- It is best to give each user or service the minimum privileges needed to perform their duties. This is the [principle of least privilege](#). It forms the foundations of good security. Organize all IAM users and groups according to levels of access they need.
- IAM has the [permission hierarchy](#) of:
  - i. Explicit deny: The most restrictive policy wins.
  - ii. Explicit allow: Access permissions to any resource has to be explicitly given.
  - iii. Implicit deny: All permissions are implicitly denied by default.
- You can test policy permissions via the AWS IAM [policy simulator tool](#). This is particularly useful if you write custom policies.

[Back to top](#) 

## Security and IAM Tips

- ◆ Use IAM to create individual user accounts and **use IAM accounts for all users from the beginning**. This is strongly recommended.
  - That way, you define different users, and groups with different levels of privilege (if you want, choose from a range of built-in roles, such as administrator, power user, etc.).
  - This allows credential revocation, which is critical in some situations. If an employee leaves, or a key is compromised, you can easily revoke their credentials with little effort.
  - You can set up [Active Directory federation](#) to use organizational accounts in AWS.

- **!** [Enable MFA](#) on your account.
  - You should always use MFA, and the sooner the better — enabling it when you already have many users is even better.
  - Unfortunately it can't be enforced in software, so an administrative policy has to be established.
  - Most users can use the Google Authenticator app (on [iOS](#) or [Android](#)) to support two-factor authentication. Consider a hardware fob.
- **!** Restrict use of significant IAM credentials as much as possible. Remember that in the cloud, loss of a highly credentialized user could essentially mean "game over," for your deployment, your users, or your whole company.
  - **Do NOT** use the [Root User account](#) other than when you initially create your account. Create custom IAM users and roles for your applications instead.
    - Lock up access and use of the root credentials as much as possible. Ideally they should be effectively "cold" — never used in deployments, this means attached to an actual MFA device, physically secured and rarely used.
- **!** Turn on CloudTrail: One of the first things you should do is [enable CloudTrail](#). Even if you are not a security hacker, it's important to do this from the beginning, so you have data on what has been happening in your AWS account should you ever need to investigate. You'll likely also want to set up a [log management service](#) to search and access these logs.
- **◆ Use IAM roles for EC2:** Rather than assign IAM users to applications like services and then sharing the sensitive [access keys](#), [assign roles to EC2 instances](#) and have applications retrieve credentials from the [instance metadata](#).
- Assign IAM roles by realm — for example, to development, staging, and production. If you're setting up a role, it's good to think about the specific realm so you have clean separation. This prevents, for example, a development instance from connecting to production resources.
- **Best practices:** AWS' [list of best practices](#) is worth reading in full up front.
- **IAM Reference:** [This interactive reference for all IAM actions, effects, and resources](#) is great to have open while you're learning how to understand existing IAM policies.
- **Multiple accounts:** Decide on whether you want to use multiple AWS accounts and [research](#) how to organize accounts
  - Number of users
  - Importance of isolation
    - Resource Limits
    - Permission granularity
    - Security
    - API Limits

- Regulatory issues
- Workload
- Size of infrastructure
- Cost of multi-account “overhead”: Internal AWS service management tools may need to be custom built or
  - ◆ It can help to use separate AWS accounts for independent parts of your infrastructure if you expect a high volume of calls since AWS [throttles calls](#) at the AWS account level.
- [Inspector](#) is an automated security assessment service from AWS that helps identify common security risks. This service adheres to certain security practices and may help with compliance.
- [Trusted Advisor](#) addresses a variety of best practices, but also offers some basic security checks around IAM usage, AWS Lambda configurations, and MFA. At paid support tiers, Trusted Advisor exposes additional checks around other areas, such as cost optimization.
- **Use KMS for managing keys:** AWS offers [KMS](#) for securely managing encryption keys, which is usually a far better approach than managing key security yourself. See [below](#).
- [AWS WAF](#) is a web application firewall to help you protect your applications from common attack patterns.
- **Security auditing:**
  - [Security Monkey](#) is an open source tool that is designed to assist with security audits.
  - [Scout2](#) is an open source tool that uses AWS APIs to assess an environment’s security posture. Scout2 is stable and actively maintained.
  - ◆ **Export and audit security settings:** You can audit security policies simply by exporting settings using AWS Lambda script like [SecConfig.py](#) (from [this 2013 talk](#)) and then reviewing and monitoring changes manually or automatically.

[Back to top ↑](#)

## Security and IAM Gotchas and Limitations

- **! Don’t share user credentials:** It’s remarkably common for first-time AWS users to create one account and one set of user credentials (username, key or password), and then use them for a while, sharing among engineers and others within a company. This is a bad idea for many reasons, but in particular, if you do, you will have reduced ability to revoke credentials on a per-service basis (for example, if an employee leaves or a key is compromised), which can lead to serious compliance issues.

- **! Instance metadata throttling:** The [instance metadata service](#) has rate limiting on API calls. If you deploy IAM should!) and have lots of services, you may hit global account limits easily.
  - One solution is to have code or scripts cache and reuse the credentials locally for a short period (say 2 minutes) be put into the `~/.aws/credentials` file but must also be refreshed automatically.
  - But be careful not to cache credentials for too long, as [they expire](#). (Note the other [dynamic metadata](#) also should not be cached a long time, either.)
- **◆ Some IAM operations are slower than other API calls** (many seconds), since AWS needs to propagate these globally.
- **! The uptime of IAM's API has historically been lower than that of the instance metadata API.** Be wary of incorporating IAM's API into critical paths or subsystems — for example, if you validate a user's IAM group membership when they log in and aren't careful about precaching group membership or maintaining a back door, you might end up locking users out if the API isn't available.
- **! Don't check in AWS credentials or secrets to a git repository.** There are bots that scan GitHub looking for credentials in code repositories. Use [AWS CloudWatch Metrics Insights](#) and other security tools, such as [git-secrets](#) to prevent anyone on your team from checking in sensitive information to your git repository.

## S3

---

### S3 Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- S3 (Simple Storage Service) is AWS' standard cloud storage service, offering file (opaque "blob") storage of arbitrary size and type, almost any size, from 0 to **5TB**. (Prior to [2011](#) the maximum size was 5 GB; larger sizes are now well supported via [Amazon S3 Large Object](#).)
- Items, or **objects**, are placed into named **buckets** stored with names which are usually called **keys**. The main components of an S3 object are the key, the bucket name, and the version ID.
- Objects are created, deleted, or updated. Large objects can be streamed, but you cannot modify parts of a value. Instead, you must upload a new object with a different key. You can use [Amazon S3 Select](#) to query the contents of a whole object.
- Every object also has [metadata](#), which includes arbitrary key-value pairs, and is used in a way similar to HTTP headers. Some are system-defined, some are significant when serving HTTP content from buckets or CloudFront, and you can also define your own use.
- **S3 URIs:** Although often bucket and key names are provided in APIs individually, it's also common practice to write URLs in the form `'s3://bucket-name/path/to/key'` (where the key here is `'path/to/key'`). (You'll also see `'s3n://'` and `'s3a://'` prefixes for S3 objects.)

- **S3 vs Glacier, EBS, and EFS:** AWS offers many storage services, and several besides S3 offer file-type abstractions infrequently accessed archival storage. [EBS](#), unlike S3, allows random access to file contents via a traditional filesystem attached to one EC2 instance at a time. [EFS](#) is a network filesystem many instances can connect to, but at higher cost than S3. [Amazon Glacier](#) is for infrequently accessed archival storage.

[Back to top](#) 

## S3 Tips

- For most practical purposes, you can consider S3 capacity unlimited, both in total size of files and number of objects in a bucket. The number of objects in a bucket is essentially also unlimited. Customers routinely have millions of objects.
- **Permissions:**
  - If you're storing business data on Amazon S3, it's important to manage permissions sensibly. In 2017 companies like [Facebook](#) and [Verizon](#) saw data breaches due to poorly-chosen S3 configuration for sensitive data. Fixing this later can be difficult because S3 has a lot of assets and internal users.
  - There are 3 different ways to grant permissions to access Amazon S3 content in your buckets.
    - **IAM policies** use the familiar [Identity and Access Management](#) permission scheme to control access to individual users or groups.
    - **Bucket policies** grant or deny permissions to an entire bucket. You might use this when hosting a website or making a bucket publicly readable, or to restrict access to a bucket by IP address. Amazon's [sample bucket policies](#) show some cases where these policies come in handy.
    - **Access Control Lists** (ACLs) can also be applied to every bucket and object stored in S3. ACLs grant additional permissions beyond those specified in IAM or bucket policies. ACLs can be used to grant access to another AWS user or even the general public. This is powerful but can be dangerous, because you need to inspect every object before granting access.
  - AWS' [predefined access control groups](#) allow access that may not be what you'd expect from their names.
    - "All Users", or "Everyone", grants permission to the general public, not only to users defined in your IAM policy. If an object is available to All Users, then it can be retrieved with a simple HTTP request of the form `http://\[bucket\].s3.amazonaws.com/\[object\]`. No authorization or signature is required to access data in this category.
    - "Authenticated Users" grants permissions to anyone with an AWS account, again not limited to your own account. Anyone can sign up for AWS, for all intents and purposes this is also open to the general public.
    - "Log Delivery" group is used by AWS to write logs to buckets and should be safe to enable on the bucket.

- A typical use case of this ACL is used in conjunction with the [requester pays](#) functionality of S3.
- ! Bucket permissions and object permissions are two different things and independent of each other. A private bucket can be seen when listing the bucket, but not downloaded. At the same time, a public object in a private bucket because the bucket contents can't be listed, but can still be downloaded by anyone who knows its exact key. Access to set bucket permissions can still make objects public if they have `s3:PutObjectAcl` or `s3:PutObject`.
- 🐚 In August 2017, AWS added [AWS Config rules to ensure your S3 buckets are secure](#).
  - ! These AWS Config rules only check the security of your bucket policy and bucket-level ACLs. You can't use them to detect or prevent rules that grant additional permissions, including opening files to the whole world.
- ♦ Do create new buckets if you have different types of data with different sensitivity levels. This is much less complex than creating a single bucket with many complex permissions rules. For example, if data is for administrators only, like log data, put it in a new bucket that only administrators can access.
- For more guidance, see:
  - [How to Secure an Amazon S3 Bucket](#)
  - [Deep dive into S3 access controls](#).
  - [How do S3 permissions work?](#)
- **Bucket naming:** Buckets are chosen from a global namespace (across all regions, even though S3 itself stores data in the region you select), so you'll find many bucket names are already taken. Creating a bucket means taking ownership of the bucket. Bucket names have [a few restrictions](#) on them.
  - Bucket names can be used as part of the hostname when accessing the bucket or its contents, like `<bucket>.1.amazonaws.com`, as long as the name is [DNS compliant](#).
  - A common practice is to use the company name acronym or abbreviation to prefix (or suffix, if you prefer) the bucket names (but please, don't use a check on this as a security measure — this is highly insecure and easily guessable).
  - ♦ Bucket names with '.' (periods) in them [can cause certificate mismatches](#) when used with SSL. Use '-' instead, which conforms with both SSL expectations and is DNS compliant.
- **Versioning:** S3 has [optional versioning support](#), so that all versions of objects are preserved on a bucket. This is useful for creating an archive of changes or the ability to back out mistakes (caution: it lacks the featureset of full version control systems).
- **Durability:** Durability of S3 is extremely high, since internally it keeps several replicas. If you don't delete it by accident, S3 won't lose your data. (AWS offers the seemingly improbable durability rate of [99.999999999%](#), but this is a marketing claim.)

based on independent failure rates and levels of replication — not a true probability estimate. Either way, S3 has durability.) Note this is *much* higher durability than EBS!

-  **S3 pricing** depends on [storage, requests, and transfer](#).
  - For transfer, putting data into AWS is free, but you'll pay on the way out. Transfer from S3 to EC2 in the same region is free; transfer to other regions or the Internet in general is not free.
  - Deletes are free.
- **S3 Reduced Redundancy and Infrequent Access:** Most people use the Standard storage class in S3, but there are other options with lower cost:
  -  [Reduced Redundancy Storage \(RRS\)](#) has been [effectively deprecated](#), and has lower durability (99.99%, standard S3). Note that it no longer participates in S3 price reductions, so it offers worse redundancy for most use cases. As a result, there's no reason to use it.
  - [Infrequent Access \(IA\)](#) lets you get cheaper storage in exchange for more expensive access. This is great for data that is already processed, but might want to look at later. To get an idea of the cost savings when using Infrequent Access, see [this S3 Infrequent Access Calculator](#).
  - [S3 - Intelligent Tiering](#) storage class is designed to optimize costs by automatically moving data to the most appropriate storage class without performance impact or operational overhead.
  - [S3 - One Zone - IA](#) is for data that is accessed less frequently, but requires rapid access when needed. Unlike Standard-IA, which stores data in a minimum of three Availability Zones (AZs), S3 One Zone-IA stores data in a single AZ and is called Standard-IA.
  - [Glacier](#) is a third alternative discussed as a separate product.
  - See [the comparison table](#).
-  **Performance:** Maximizing S3 performance means improving overall throughput in terms of bandwidth and number of objects per second.
  - S3 is highly scalable, so in principle you can get arbitrarily high throughput. (A good example of this is [S3Direct](#).)
  - But usually you are constrained by the pipe between the source and S3 and/or the level of concurrency of concurrent operations.
  - Throughput is of course highest from within AWS to S3, and between EC2 instances and S3 buckets that are in the same region.
  - Bandwidth from EC2 depends on instance type. See the "Network Performance" column at [ec2instances.info](#).
  - Throughput of many objects is extremely high when data is accessed in a distributed way, from many EC2 instances, to read or write objects from S3 from hundreds or thousands of instances at once.

- However, throughput is very limited when objects accessed sequentially from a single instance. Individual operations take milliseconds, and bandwidth to and from instances is limited.
  - Therefore, to perform large numbers of operations, it's necessary to use multiple worker threads and connect to multiple instances, and for larger jobs, multiple EC2 instances as well.
  - **Multi-part uploads:** For large objects you want to take advantage of the multi-part uploading capabilities (splitting the object into smaller chunks and uploading them in parallel). The chunk sizes of 5 MB.
  - **Large downloads:** Also you can download chunks of a single large object in parallel by exploiting the HTTP range header capability.
  - ◆ **List pagination:** Listing contents happens at 1000 responses per request, so for buckets with many millions of objects, this can take time.
  - ! **Key prefixes:** Previously randomness in the beginning of key names was necessary in order to avoid hotspots. This is no longer necessary as of July, 2018.
  - For data outside AWS, [DirectConnect](#) and [S3 Transfer Acceleration](#) can help. For S3 Transfer Acceleration, you can upload data to S3 from anywhere in the world at speeds equivalent of 1-2 months of storage for the transfer in either direction for using nearer endpoints.
- **Command-line applications:** There are a few ways to use S3 from the command line:
    - Originally, [s3cmd](#) was the best tool for the job. It's still used heavily by many.
    - The regular [aws](#) command-line interface now supports S3 well, and is useful for most situations.
    - [s4cmd](#) is a replacement, with greater emphasis on performance via multi-threading, which is helpful for large files, and also offers Unix-like globbing support.
  - **GUI applications:** You may prefer a GUI, or wish to support GUI access for less technical users. Some options:
    - The [AWS Console](#) does offer a graphical way to use S3. Use caution telling non-technical people to use it, however, as it offers full administrative permissions, it offers access to many other AWS features.
    - [Transmit](#) is a good option on macOS for most use cases.
    - [Cyberduck](#) is a good option on macOS and Windows with support for multipart uploads, ACLs, versioning, lifecycle rules, storage classes and server side encryption (SSE-S3 and SSE-KMS).
  - **S3 and CloudFront:** S3 is tightly integrated with the CloudFront CDN. See the CloudFront section for more information on [CloudFront and S3 transfer acceleration](#).
  - **Static website hosting:**

- S3 has a [static website hosting option](#) that is simply a setting that enables configurable HTTP index and error [support](#) to [public content](#) in S3. It's a simple way to host static assets or a fully static website.
- Consider using CloudFront in front of most or all assets:
  - Like any CDN, CloudFront improves performance significantly.
  - ◆ SSL is only supported on the built-in amazonaws.com domain for S3. S3 supports serving these sites [domain](#), but [not over SSL on a custom domain](#). However, [CloudFront allows you to serve a custom domain](#) provides free SNI SSL/TLS certificates via Amazon Certificate Manager. [SNI does not work on very old systems](#). Alternatively, you can provide your own certificate to use on CloudFront to support all browser fee.
  - ◆ If you are including resources across domains, such as fonts inside CSS files, you may need to [configure](#) serving those resources.
  - Since pretty much everything is moving to SSL nowadays, and you likely want control over the domain, set up CloudFront with your own certificate in front of S3 (and to ignore the [AWS example on this](#) as it is not applicable).
  - That said, if you do, you'll need to think through invalidation or updates on CloudFront. You may wish to [hashes in filenames](#) so invalidation is not necessary.
- **Data lifecycles:**
  - When managing data, the understanding the lifecycle of the data is as important as understanding the data. When you upload objects into a bucket, think about its lifecycle — its end of life, not just its beginning.
  - ◆ In general, data with different expiration policies should be stored under separate prefixes at the top level. Voluminous logs might need to be deleted automatically monthly, while other data is critical and should never be deleted. Storing the former in a separate bucket or at least a separate folder is wise.
  - ◆ Thinking about this up front will save you pain. It's very hard to clean up large collections of files created with varying lifecycles and no coherent organization.
  - Alternatively you can set a lifecycle policy to archive old data to Glacier. [Be careful](#) with archiving large numbers of objects to Glacier, since it may actually cost more.
  - There is also a storage class called [Infrequent Access](#) that has the same durability as Standard S3, but is designed to be suitable for objects that are infrequently accessed.
- **Data consistency:** Understanding [data consistency](#) is critical for any use of S3 where there are multiple producers and consumers of data. It's important to understand how S3 handles consistency between multiple requests and how it handles eventual consistency.

- Creation and updates to individual objects in S3 are **atomic**, in that you'll never upload a new object or change another client see only part half the change.
- The uncertainty lies with *when* your clients and other clients see updates.
- **New objects:** If you create a new object, you'll be able to read it instantly, which is called **read-after-write consistency**.
  - Well, with the additional caveat that if you do a read on an object before it exists, then create it, [you get read-after-write](#)).
  - This does not apply to any list operations; newly created objects are [not guaranteed to appear in a list](#).
- **Updates to objects:** If you overwrite or delete an object, you're only guaranteed **eventual consistency**, i.e. that you have no guarantee of when.
  - ◆ For many use cases, treating S3 objects as **immutable** (i.e. deciding by convention they will be created once) can greatly simplify the code that uses them, avoiding complex state management.
  - ◆ Note that [until 2015](#), 'us-standard' region had had a weaker eventual consistency model, and the other (non-US) regions had stronger consistency guarantees. This was finally corrected — but watch for many old blogs mentioning this!
- **Slow updates:** In practice, "eventual consistency" usually means within seconds, but expect rare cases of minutes or hours.
- **S3 as a filesystem:**
  - In general S3's APIs have inherent limitations that make S3 hard to use directly as a POSIX-style filesystem without its own object format. For example, appending to a file requires rewriting, which cripples performance, and atomic locking is mutual exclusion on opening files, and hardlinks are impossible.
  - [s3fs](#) is a FUSE filesystem that goes ahead and tries anyway, but it has performance limitations and surprises.
  - [Riofs](#) (C) and [Goofys](#) (Go) are more recent efforts that attempt to adopt a different data storage format to address some of the limitations of s3fs.
  - [S3QL](#) ([discussion](#)) is a Python implementation that offers data de-duplication, snap-shottting, and encryption at rest.
  - [ObjectiveFS](#) ([discussion](#)) is a commercial solution that supports filesystem features and concurrent clients.
- If you are primarily using a VPC, consider setting up a [VPC Endpoint](#) for S3 in order to allow your VPC-hosted resources to access S3 without the need for extra network configuration or hops.
- **Cross-region replication:** S3 has [a feature](#) for replicating a bucket between one region and another. Note that S3 buckets are replicated within one region, so usually this isn't necessary for durability, but it could be useful for compliance (geographic redundancy).

data storage), lower latency, or as a strategy to reduce region-to-region bandwidth costs by mirroring heavily used region.

- **IPv4 vs IPv6:** For a long time S3 only supported IPv4 at the default endpoint <https://BUCKET.s3.amazonaws.com>. In [2016](#) it now supports both IPv4 & IPv6! To use both, you have to [enable dualstack](#) either in your preferred API client or this url scheme <https://BUCKET.s3.dualstack.REGION.amazonaws.com>. This extends to S3 Transfer Acceleration as well.
- **S3 event notifications:** S3 can be configured to send an [SNS notification](#), [SQS message](#), or [AWS Lambda function](#).
- Limit your individual users (or IAM roles) to the minimal required S3 locations, and catalog the “approved” locations. S3 tends to become the dumping ground where people put data to random locations that are not cleaned up for years.
- If a bucket is deleted in S3, it can take up to 10 hours before a bucket with the same name can be created again.

[Back to top ↑](#)

## S3 Gotchas and Limitations

- S3 buckets sit outside the VPC and can be accessed from anywhere in the world if bucket policies are not set correctly. See the permissions section above carefully, there are countless cases of buckets exposed to the public.
- For many years, there was a notorious [100-bucket limit](#) per account, which could not be raised and caused many headaches. As of 2015, you can [request increases](#). You can ask to increase the limit, but it will still be capped (generally at 100 buckets per account).
- Be careful not to make implicit assumptions about transactionality or sequencing of updates to objects. Never assume that multiple writes to different objects in the same sequence will be applied sequentially. If you upload multiple objects to S3, modify a sequence of objects, the clients will see the same modifications in the same sequence, or if you upload multiple objects to S3, they will all appear at once to all clients.
- S3 has an [SLA](#) with 99.9% uptime. If you use S3 heavily, you'll inevitably see occasional errors accessing or storing data. Infrastructure fail. Availability is usually restored in seconds or minutes. Although availability is not extremely high, durability is excellent.
- After uploading, any change that you make to the object causes a full rewrite of the object, so avoid appending to regular files.
- Eventual data consistency, as discussed above, can be surprising sometimes. If S3 suffers from internal replication issues, data may not be visible from a subset of the machines, depending on which S3 endpoint they hit. Those usually resolve within seconds.

seen isolated cases when the issue lingered for 20-30 hours.

- ◆ **MD5s and multi-part uploads:** In S3, the [ETag header in S3](#) is a hash on the object. And in many cases, it is [not the case in general](#) when you use multi-part uploads. One workaround is to compute MD5s yourself and provide header (such as is done by [s4cmd](#)).
- ◆ **Incomplete multi-part upload costs:** Incomplete multi-part uploads accrue [storage charges](#) even if the upload is created. [Amazon \(and others\)](#) recommend using a lifecycle policy to clean up incomplete uploads and save on storage costs. If you have many of these, it may be worth investigating whatever's failing regularly.
- ◆ **US Standard region:** Previously, the us-east-1 region (also known as the US Standard region) was replicated across the globe, leading to greater variability of latency. Effective Jun 19, 2015 this is [no longer the case](#). All Amazon S3 regions now support regional consistency. Amazon S3 also renamed the US Standard region to the US East (N. Virginia) region to be consistent with other AWS services' naming conventions.
- ◆ **S3 authentication versions and regions:** In newer regions, S3 [only supports the latest authentication](#). If an S3 [SDK](#) doesn't work in one region, but works correctly in another region, make sure you are using the latest [authentication](#).

[Back to top ↑](#)

## Storage Durability, Availability, and Price

As an illustration of comparative features and price, the table below gives S3 Standard, RRS, IA, in comparison with Glacier, and d2.xlarge instance store using [Virginia region](#) as of Sept 2017.

	Durability (per year)	Availability “designed”	Availability SLA	Storage (per TB per month)	GET retrieve millions
Glacier	Eleven 9s	Slow	—	\$4	\$50
S3 IA	Eleven 9s	99.9%	99%	\$12.50	\$1
S3 RRS	99.99%	99.99%	99.9%	\$24 (first TB) \$0.40 (subsequent)	\$0.40
S3 Standard	Eleven 9s	99.99%	99.9%	\$23	\$0.40

	Durability (per year)	Availability "designed"	Availability SLA	Storage (per TB per month)	GET retrieve millions
EBS	<b>99.8%</b>	Unstated	99.99%	\$25/\$45/\$100/\$125+ ( <a href="#">sc1/st1/gp2/io1</a> )	
EFS	"High"	"High"	–	\$300	
EC2 d2.xlarge instance store	Unstated	Unstated	–	\$25.44	\$0

Especially notable items are in **boldface**. Sources: [S3 pricing](#), [S3 SLA](#), [S3 FAQ](#), [RRS info](#) (note that this is considered [durable storage](#))  
[EBS availability and durability](#), [EBS pricing](#), [EFS pricing](#), [EC2 SLA](#)

## EC2

---

### EC2 Basics

-  [Homepage](#) · [Documentation](#) · [FAQ](#) · [Pricing](#) (see also [ec2instances.info](#))
- EC2 (Elastic Compute Cloud) is AWS' offering of the most fundamental piece of cloud computing: A [virtual private cloud](#) can run [most Linux, BSD, and Windows operating systems](#). Internally, they've used a heavily modified [Xen](#) virtual machine monitor. New instance classes are being introduced with a KVM derived hypervisor instead, called [Nitro](#). So far, this is limited to [t2](#) and [m5](#) types. Lastly, there's a "bare metal hypervisor" available for [i3.metal instances](#)
- The term "EC2" is sometimes used to refer to the servers themselves, but technically refers more broadly to a whole range of services. This includes the compute instances, as well as the supporting services, too, like load balancing (CLBs/ALBs/NLBs), IP addresses (EIPs), bootable images (AMIs), security groups, and storage (EBS) (which we discuss individually in this guide).
-  [EC2 pricing](#) and [cost management](#) is a complicated topic. It can range from free (on the [AWS free tier](#)) to a lot of money depending on usage. Pricing is by instance type, by second or hour, and changes depending on AWS region and whether you are using On-Demand instances, [On-Demand](#), on the [Spot market](#) or pre-purchasing ([Reserved Instances](#)).
- **Network Performance:** For some instance types, AWS uses general terms like Low, Medium, and High to refer to network performance. Users have done [benchmarking](#) to provide expectations for what these terms can mean.

## EC2 Alternatives and Lock-In

- Running EC2 is akin to running a set of physical servers, as long as you don't do automatic scaling or toolled clustering. Migrating to another VPS or dedicated server provider should not be too hard.
- **Alternatives to EC2:** The direct alternatives are Google Cloud, Microsoft Azure, Rackspace, DigitalOcean, AWS Lambda, and other VPS providers, some of which offer similar APIs for setting up and removing instances. (See the comparison table below.)
- **Should you use Amazon Linux?** AWS encourages use of their own [Amazon Linux](#), which is evolved from [Red Hat Enterprise Linux](#) and [CentOS](#). It's used by many, but [others are skeptical](#). Whatever you do, think this decision through carefully. While Amazon Linux is heavily tested and better supported in the unlikely event you have deeper issues with OS and virtualization on EC2, many companies do just fine using a standard, non-Amazon Linux distribution, such as Ubuntu or CentOS. Using a standard Linux distribution means you have an exactly replicable environment should you use another hosting provider instead of (or in addition to) AWS. It's also helpful if you wish to test deployments on local developer machines running the same standard Linux distribution. Docker support is more common with Docker, too. Amazon now supports an official [Amazon Linux Docker image](#), aimed at assisting users in getting started on a comparable environment, though this is new enough that it should be considered experimental). Note that the latest version of [Amazon Linux 2](#) supports on-premise deployments explicitly.
- **EC2 costs:** See the [section on this](#).

## EC2 Tips

- **Picking regions:** When you first set up, consider which [regions](#) you want to use first. Many people in North America tend to set up in the us-east-1 (N. Virginia) region, which is the default, but it's worth considering if this is best up front. Factors include service availability (some services [are not available in all regions](#)), costing (baseline costs also [vary by region](#) with the lowest in us-east-1 for comparison purposes)), and compliance (various countries have differing regulations with respect to data storage, for example).
- **Instance types:** EC2 instances come in many types, corresponding to the capabilities of the virtual machine in CPU, memory, RAM, disk sizes and types (SSD or magnetic), and network bandwidth.

- Selecting instance types is complex since there are so many types. Additionally there are different generations of instances.
  - ♦ Use the list at [ec2instances.info](#) to review costs and features. [Amazon's own list](#) of instance types is harder to parse, but it includes more details about features and price together, which makes it doubly difficult.
  - Prices vary a lot, so use [ec2instances.info](#) to determine the set of machines that meet your needs and [ec2pricing.com](#) to find the cheapest type in the region you're working in. Depending on the timing and region, it might be much cheaper to use a larger instance with [more memory](#) or CPU than the bare minimum.
- **Turn off** your instances when they aren't in use. For many situations such as testing or staging resources, you may leave them running 24/7, but for most production workloads, you can turn them off when they aren't needed. This is a simple mechanism for cost savings. This can be achieved using [Lambda and CloudWatch](#), deploying the [Instance Scheduler](#) open source option like [cloudcycler](#), or a SaaS provider like [GorillaStack](#). (Note: if you turn off instances with an [AWS Lambda function](#), any state will be lost when the instance is turned off. Therefore, for stateful applications it is safer to turn off EBS volumes.)
- **Dedicated instances** and **dedicated hosts** are assigned hardware, instead of usual virtual instances. They are more expensive than regular instances but [can be preferable](#) for performance, compliance, financial modeling, or licensing reasons.
- **32 bit vs 64 bit:** A few micro, small, and medium instances are still available to use as 32-bit architecture. You'll likely see mostly "amd64" instances nowadays, though smaller instances still support 32 bit ("i386"). Use 64 bit unless you have good reasons to use 32.
- **HVM vs PV:** There are two kinds of virtualization technology used by EC2, [hardware virtual machine \(HVM\)](#) and [paravirtualization \(PV\)](#). Historically, PV was the usual type, but [now HVM is becoming the standard](#). If you want to use the newest instance types, you must use HVM. See the [instance type matrix](#) for details.
- **Operating system:** To use EC2, you'll need to pick a base operating system. It can be Windows or Linux, such as Amazon Linux or Ubuntu Server. You do this with AMIs, which are covered in more detail in their own section below.
- **Limits:** You can't create arbitrary numbers of instances. Default limits on numbers of EC2 instances per account vary by region and are described in [this list](#).
- **! Use termination protection:** For any instances that are important and long-lived (in particular, aren't part of a [auto-scaling group](#)), enable [termination protection](#). This is an important line of defense against user mistakes, such as accidentally terminating an instance of just one due to human error.

- **SSH key management:**
  - When you start an instance, you need to have at least one [ssh key pair](#) set up, to bootstrap, i.e., allow you to log in.
  - Aside from bootstrapping, you should manage keys yourself on the instances, assigning individual keys to individual users as appropriate.
  - Avoid reusing the original boot keys except by administrators when creating new instances.
  - Avoid sharing keys and [add individual ssh keys](#) for individual users.
- **GPU support:** You can rent GPU-enabled instances on EC2 for use in machine learning or graphics rendering workloads.
  - There are [three types](#) of GPU-enabled instances currently available:
    - The P3 series offers NVIDIA Tesla V100 GPUs in 1, 4 and 8 GPU configurations targeting machine learning and other high performance computing applications.
    - The P2 series offers NVIDIA Tesla K80 GPUs in 1, 8 and 16 GPU configurations targeting machine learning and other high performance computing applications.
    - The G3 series offers NVIDIA Tesla M60 GPUs in 1, 2, or 4 GPU configurations targeting graphics and video processing.
  - AWS offers two different AMIs that are targeted to GPU applications. In particular, they target deep learning and provide access to more stripped-down driver-only base images.
    - AWS offers both an Amazon Linux [Deep Learning AMI](#) (based on Amazon Linux) as well as an Ubuntu [Deep Learning AMI](#). These AMIs come with most NVIDIA drivers and ancillary software (CUDA, CUBLAS, CuDNN, TensorFlow, PyTorch, etc.) pre-installed, which removes a significant barrier to usage.
    - Note that using these AMIs can lead to lock in due to the fact that you have no direct access to software updates or versioning.
    - The compendium of frameworks included can lead to long instance startup times and difficult-to-remove dependencies.
  - ♦ As with any expensive EC2 instance types, [Spot instances can offer significant savings](#) with GPU workloads if you are willing to tolerate some downtime.
- All current EC2 instance types can take advantage of IPv6 addressing, so long as they are launched in a subnet with an IPv6 range in an IPv6-enabled VPC.

## EC2 Gotchas and Limitations

- **!** Never use ssh passwords. Just don't do it; they are too insecure, and consequences of compromise too severe. [up on this](#) and fully disable ssh password access to your ssh server by making sure 'PasswordAuthentication no' is set in the /etc/ssh/sshd\_config file. If you're careful about managing ssh private keys everywhere they are stored, it is a major security over password-based authentication.
- **◆** For all [newer instance types](#), when selecting the AMI to use, be sure you select the HVM AMI, or it just won't work.
- **!** When creating an instance and using a new ssh key pair, [make sure the ssh key permissions are correct](#).
- **◆** Sometimes certain EC2 instances can get scheduled for retirement by AWS due to "detected degradation of the instance". In which case you are given a couple of weeks to migrate to a new instance.
  - If your instance root device is an EBS volume, you can typically stop and then start the instance which moves the hardware, giving you control over timing of this event. Note however that you will lose any instance store volumes ([drives](#)) if your instance type has instance store volumes.
  - The instance public IP (if it has one) will likely change unless you're using Elastic IPs. This could be a problem on the IP address.
- **◆** Periodically you may find that your server or load balancer is receiving traffic for (presumably) a previous EC2 instance that has been terminated and replaced with a new one with the same IP address that you are handed out now (this may not matter, or it can be fixed by migrating to another IP).
- **!** If the EC2 API itself is a critical dependency of your infrastructure (e.g. for automated server replacement, custom scripts etc.) and you are running at a large scale or making many EC2 API calls, make sure that you understand when the API is [rate limited](#) and the limits are not published and subject to change) and code and test against that possibility.
- **!** Many newer EC2 instance types are either EBS-only, or backed by local NVMe disks assigned to the instance. Consider performance and costs when planning to use them.
- **!** If you're operating at significant scale, you may wish to break apart API calls that enumerate all of your resources either on individual resources, or a subset of the entire list. EC2 APIs will time out! Consider using [filters](#) to restrict the results.
- **!**  Instances come in two types: **Fixed Performance Instances** (e.g. M3, C3, and R3) and **Burstable Performance Instances** (T2). T2 instances receive CPU credits continuously, the rate of which depends on the instance size. T2 instances accrue credits when they are active, and use CPU credits when they are active. However, once an instance runs out of credits, you'll notice a performance drop. If you need consistently high CPU performance for applications such as video encoding, high volume databases, or machine learning inference, it is recommended to use Fixed Performance Instances.

- Instance user-data is [limited to 16 KB](#). (This limit applies to the data in raw form, not base64-encoded form.) If more than 16KB of user-data is provided, it will be truncated.
- Very new accounts may not be able to launch some instance types, such as GPU instances, because of an initially zero. This limit can be raised by making a support request. See [AWS Service Limits](#) for the method to make the service aware that this limit of zero is [not currently documented](#).
- Since multiple AWS instances all run on the same physical hardware, early cloud adopters encountered what became known as the [Neighbor problem](#). This feeling of not getting what you are paying for led to [user frustration](#), however "steal" mode was introduced to describe what's actually happening based on a [detailed explanation of how the kernel determine steal time](#). Avoiding this behavior and affecting your application in the cloud may be best handled by [properly designing your cloud architecture](#).
- AWS [introduced Dedicated Tenancy](#) in 2011. This allows customers to have all resources from a single server. So instead of sharing a server with other customers, a customer gets their own dedicated server. This solves the [noisy neighbor problem](#) since only that customer uses the CPU. This approach comes with a significant cost increase, but no shared maintenance. If a customer had 20 instances running using shared tenancy and one underlying server needed any type of maintenance. If a customer had 20 instances running using dedicated tenancy, only the instance on that server would go offline. If that customer had 20 instances running using dedicated tenancy, all 20 instances would go offline.
- ◆ Only **i3.metal** type instances providing an ability to run Android x86 emulators on AWS at the moment.

## CloudWatch

---

### CloudWatch Basics

-  [Homepage](#) · [Documentation](#) · [FAQ](#) · [Pricing](#)
- CloudWatch monitors resources and applications, captures logs, and sends events.
- CloudWatch monitoring is the standard mechanism for keeping tabs on AWS resources. A wide range of [metrics](#) are available via CloudWatch, allowing you to create time based graphs, [alarms](#), and [dashboards](#).
  - Alarms are the most practical use of CloudWatch, allowing you to trigger notifications from any given metric.
  - Alarms can trigger [SNS notifications](#), [Auto Scaling actions](#), or [EC2 actions](#).
  - Alarms also support [alerting when any M out of N datapoints cross the alarm threshold](#).
  - Publish and share graphs of metrics by creating [customizable dashboard views](#).
    - Monitor and report on EC2 [instance system check failure alarms](#).

- **Using CloudWatch Events:**
  - Events create a mechanism to automate actions in various services on AWS. You can create [event rules](#) from Auto Scaling, Run commands, deployments or time-based schedules (think Cron).
  - [Triggered events](#) can invoke Lambda functions, send SNS/SQS/Kinesis messages, or perform instance actions (like snapshot volumes).
  - Custom payloads can be sent to targets in JSON format, this is especially useful when triggering Lambdas.
- **Using CloudWatch Logs:**
  - [CloudWatch Logs](#) is a streaming log storage system. By storing logs within AWS you have access to unlimited storage and can also have the option of streaming logs directly to ElasticSearch or custom Lambdas.
  - A [log agent installed](#) on your servers will process logs over time and send them to CloudWatch Logs.
  - You can [export logged data to S3](#) or stream results to other AWS services.
  - CloudWatch Logs can be [encrypted using keys managed through KMS](#).
- **Detailed monitoring:** [Detailed monitoring](#) for EC2 instances must be enabled to get granular metrics, and is [billable](#).

[Back to top](#) 

## CloudWatch Alternatives and Lock-In

- CloudWatch offers fairly basic functionality that doesn't create significant (additional) AWS lock-in. Most of the required functionality can be obtained through APIs that can be imported into other aggregation or visualization tools or services that provide CloudWatch data import services).
- ■ Alternatives to CloudWatch monitoring services include [NewRelic](#), [Datadog](#), [Sumo Logic](#), [Zabbix](#), [Nagios](#), [RabbitMQ](#), and many others. These are open source options such as [StatsD](#) or [collectd](#) with [Graphite](#), and many others.
- ■ CloudWatch Log alternatives include [Splunk](#), [Sumo Logic](#), [Loggly](#), [LogDNA](#), [Logstash](#), [Papertrail](#), [Elastic Stack](#), and many others.

[Back to top](#) 

## CloudWatch Tips

- Some very common use cases for CloudWatch are [billing alarms](#), [instance or load balancer up/down alarms](#), and [AWS CloudTrail](#).

- You can use [EC2Config](#) to monitor watch memory and disk metrics on Windows platform instances. For Linux, there are other tools that do the same thing.
- You can [publish your own metrics](#) using the AWS API. [Incur additional cost.](#)
- You can stream directly from CloudWatch Logs to a Lambda or ElasticSearch cluster by creating [subscriptions](#) or [metrics filters](#).
- Don't forget to take advantage of the [CloudWatch non-expiring free tier](#).

[Back to top](#) 

## CloudWatch Gotchas and Limitations

- ♦ Metrics in CloudWatch originate [on the hypervisor](#). The hypervisor doesn't have access to OS information, so metrics (such as memory utilization) are not available unless pushed to CloudWatch from inside the instance.
- ♦ You can not use [more than one metric for an alarm](#).
- ♦ Notifications you receive from alarms will not have any contextual detail; they have only the specifics of the trigger.
- ♦ By default, CloudWatch metric resolution is 1 minute. If you send multiple values of a metric within the same minute, CloudWatch aggregates into minimum, maximum, average and total (sum) per minute.
- 🎉 In July 2017, a new [high-resolution option](#) was added for CloudWatch metrics and alarms. This feature allows CloudWatch to publish metrics at 1-second resolution, and to evaluate CloudWatch alarms every 10 seconds.
  - The [blog post introducing this feature](#) describes how to publish a high-resolution metric to CloudWatch. Note that the `PutMetricData` API's `StorageResolution` parameter is an attribute of each item you send in the `MetricData` array, not a parameter of the `PutMetricData` API call.
- ♦ Data about metrics is kept in CloudWatch [for 15 months](#), starting November 2016 (used to be 14 days). Metrics older than 15 months are deleted.

## AMIs

### AMI Basics

- 📖 [User guide](#)

- AMIs (Amazon Machine Images) are immutable images that are used to launch preconfigured EC2 instances. They can have private or public flavors. Access to public AMIs is either freely available (shared/community AMIs) or bought and sold in the AWS Marketplace.
- Many operating system vendors publish ready-to-use base AMIs. For Ubuntu, see the [Ubuntu AMI Finder](#). Amazon also provides its own [Amazon Linux](#).

[Back to top ↑](#)

## AMI Tips

- AMIs are built independently based on how they will be deployed. You must select AMIs that match your deployment strategy:
  - EBS or instance store
  - PV or HVM [virtualization types](#)
  - 32 bit ("i386") vs 64 bit ("amd64") architecture
- As discussed above, modern deployments will usually be with **64-bit EBS-backed HVM**.
- You can create your own custom AMI by [snapshotting the state](#) of an EC2 instance that you have modified.
- [AMIs backed by EBS storage](#) have the necessary image data loaded into the EBS volume itself and don't require it to be copied to the instance storage, which results in EBS-backed instances coming up much faster than instance storage-backed ones.
- **AMIs are per region**, so you must look up AMIs in your region, or copy your AMIs between regions with the [AWS Lambda](#) function `CopyImage`.
- As with other AWS resources, it's wise to [use tags](#) to version AMIs and manage their lifecycle.
- If you create your own AMIs, there is always some tension in choosing how much installation and configuration to include:
  - Baking less into your AMIs (for example, just a configuration management client that downloads, installs, and configures software on new EC2 instances when they are launched) allows you to minimize time spent automating AMI creation and management. This approach results in a longer instance launch time but reduces the complexity of the AMI creation pipeline.
  - Baking more into your AMIs (for example, pre-installing but not fully configuring common software along with a configuration management client that loads configuration settings at launch time) results in a faster launch time and fewer errors during the instance launch process but increases the need for you to manage the complexity of the AMI creation pipeline.

- Baking even more into your AMIs (for example, installing all required software as well and potentially also configuration information) results in fast launch times and a much lower chance of instance launch-time failure (additional re-deployment and re-configuration considerations) can require time consuming AMI updates instead of configuration as well as more complex AMI creation automation processes.
- Which option you favor depends on how quickly you need to scale up capacity, and size and maturity of your system.
  - When instances boot fast, auto-scaled services require less spare capacity built in and can more quickly scale to sudden increases in load. When setting up a service with autoscaling, consider baking more into your AMIs than EBS storage option.
  - As systems become larger, it's common to have more complex AMI management, such as a multi-stage AMI where a few (ideally one) common base AMIs are infrequently regenerated when components that are common to all instances are updated and then a more frequently run "service-level" AMI generation process that includes installation and configuration of application-specific software.
- More thinking on AMI creation strategies [here](#).
- Use tools like [Packer](#) to simplify and automate AMI creation.
- If you use RHEL instances and happen to have existing RHEL on-premise Red Hat subscriptions, then you could use the [Access program](#) to migrate a portion of your subscriptions to AWS, and thereby not having AWS charge you for a second time. You can either use your own self-created RHEL AMI's or Red Hat provided [Gold Images](#) that will be used for your AMI's once you sign up for Red Hat Cloud Access.

[Back to top ↑](#)

## AMI Gotchas and Limitations

- **Amazon Linux package versions:** [By default](#), instances based on Amazon Linux AMIs are configured point to packages in Amazon's package repository. This means that the package versions that get installed are not locked to changes, including breaking ones, to appear when applying updates in the future. If you bake your AMIs with updated packages, it is unlikely to cause problems in running services whose instances are based on those AMIs – breaks will appear at the stage of your build process, and will need to be fixed or worked around before new AMIs can be generated. The feature that allows you to configure Amazon Linux instances to target the repository of a particular major version of an AMI, reducing the likelihood that breaks caused by Amazon-initiated package version changes will occur at package update time. The cost of not having updated packages get automatically installed by future update runs. Pairing use of the "lock" feature with the build process to advance the Amazon Linux AMI at your discretion can give you tighter control over update behaviors.
- **Cloud-Init Defaults:** Oftentimes users create AMIs after performing customizations (albeit manually or via something like Ansible). If you're not careful to alter cloud-init settings that correspond to the system service (e.g. sshd, etc.) you may find that your changes are no longer in effect after booting your new AMI for the first time, as cloud-init has overwritten them.

Some distros have different files than others, but all are generally located in `/etc/cloud/`, regardless of distro. You should review these files carefully for your chosen distro before rolling your own AMIs. A [complete reference to cloud-init](#) is available on the Cloud-Init site. This is an advanced configuration mechanism, so test any changes made to these files in a sandbox prior to deployment.

## Auto Scaling

---

### Auto Scaling Basics

- [Homepage](#) · [User guide](#) · [FAQ](#) · [Pricing](#) at no additional charge
- [Auto Scaling Groups \(ASGs\)](#) are used to control the number of instances in a service, reducing manual effort to manage EC2 instances.
- They can be configured through [Scaling Policies](#) to automatically increase or decrease instance counts based on instance utilization, or based on a schedule.
- There are three common ways of using ASGs - dynamic (automatically adjust instance count based on metrics for instance utilization), static (maintain a specific instance count at all times), scheduled (maintain different instance counts across days of the week).
- ASGs [have no additional charge](#) themselves; you pay for underlying EC2 and CloudWatch services.

## Auto Scaling Tips

-  Better matching your cluster size to your current resource requirements through use of ASGs can result in significant cost savings and improved performance.
- Pairing ASGs with CLBs is a common pattern used to deal with changes in the amount of traffic a service receives.
- Dynamic Auto Scaling is easiest to use with stateless, horizontally scalable services.
- Even if you are not using ASGs to dynamically increase or decrease instance counts, you should seriously consider using them inside of ASGs – given a target instance count, the ASG will work to ensure that number of instances running is equal to the target. It will automatically replace instances for you if they die or are marked as being unhealthy. This results in consistent capacity and better service.
- Autoscalers can be [configured to terminate](#) instances that a CLB or ALB has marked as being unhealthy.

## Auto Scaling Gotchas and Limitations

-  **ReplaceUnhealthy setting:** By default, ASGs will kill instances that the EC2 instance manager considers to be unhealthy for instances whose CPU is completely saturated for minutes at a time to appear to be unresponsive, causing an alarm to trigger. You can turn off the [ReplaceUnhealthy setting](#) turned on to replace them. When instances that are managed by ASGs are expected to have high CPU, consider deactivating this setting. If you do so, however, detecting and killing unhealthy nodes will be slower.

## EBS

---

### EBS Basics

-  [Homepage](#) · [User guide](#) · [FAQ](#) · [Pricing](#)
- **EBS** (Elastic Block Store) provides block level storage. That is, it offers storage volumes that can be attached as fixed network drives.

- EBS volumes can only be attached to one EC2 instance at a time. In contrast, EFS can be shared but has a much [comparison](#).

[Back to top ↑](#)

## EBS Tips

- 🌐 RAID: Use [RAID drives](#) for [increased performance](#).
- 🌐 A worthy read is AWS' [post on EBS IO characteristics](#) as well as their [performance tips](#).
- 🌐 One can [provision IOPS](#) (that is, pay for a specific level of I/O operations per second) to ensure a particular level of performance.
- 🌐 A single gp2 EBS volume allows 16k IOPS max To get the maximum performance out of a gp2 EBS volume, it needs to be large and attached to an EBS-optimized EC2 instance.
- 🌐 Standard and gp2 EBS volumes improve IOPS with size. It may make sense for you to simply enlarge a volume to get better performance explicitly. This can in many cases reduce costs by 2/3.
- A standard block size for an EBS volume is 16kb.

[Back to top ↑](#)

## EBS Gotchas and Limitations

- ! EBS durability is reasonably good for a regular hardware drive (annual failure rate of [between 0.1% - 0.2%](#)). Consider it very poor if you don't have backups! By contrast, S3 durability is extremely high. *If you care about your data, back it up and take snapshots.*
- ♦ EBS has an [SLA](#) with 99.99% uptime. See notes on high availability below.
- ! EBS volumes have a [volume type](#) indicating the physical storage type. The types called "standard" (st1 or sc1) are slow platter disks, which deliver only hundreds of IOPS — not what you want unless you're really trying to cut costs. Most likely io1 are typically the options you want.
- ! When restoring a snapshot to create an EBS volume, blocks are lazily read from S3 the first time they're referenced. If you experience a period of high latency, you may wish to use `dd` or `fio` as per the [official documentation](#).

## EFS Basics

-  [Homepage](#) · [User guide](#) · [FAQ](#) · [Pricing](#)
-  EFS is Amazon's network filesystem. It's presented as an [NFSv4.1](#) server. Any compatible NFSv4 client can mount it.
- It is designed to be highly available and durable and each EFS file system object is redundantly stored across multiple regions.
- EFS is designed to be used as a shared network drive and it can automatically scale up to petabytes of stored data across instances attached to it.
- EFS can offer [higher throughput](#) (multiple gigabytes per second) and better durability and availability than EBS (Amazon's block storage service), but with higher latency.
- EFS is priced based on the volume of data stored, and costs [much more than EBS](#); it's in the ballpark of three times the cost of general purpose gp2 EBS volumes.
-  [Performance](#) is dependent on the volume of data stored, as is the price:
  - Like EBS, EFS uses a credit based system. Credits are earned at a rate of 50 KiB/s per GiB of storage and consumed during reading/writing files or metadata. Unlike EBS, operations on metadata (file size, owner, date, etc.) also consume credits. The [BurstCreditBalance metric](#) in CloudWatch should be monitored to make sure the file system doesn't run out of credits.
  - Throughput capacity during bursts is also dependent on size. Under 1 TiB, throughput can go up to 100 MiB/s. For each additional TiB, an additional 50 MiB/s is added for each stored TiB. For instance, a file system storing 5 TiB would be able to burst at a rate of 500 MiB/s. The [PercentIOLimit metric](#) in CloudWatch hovers around 100%, Max I/O is recommended. Changing performance modes is not recommended for new EFS and migrating data.
  - EFS has two performance modes that can only be set when a file system is created. One is "General Purpose", which provides higher throughput and lower latency. The other is "Throughput Optimized for Amazon Lambda", which provides higher throughput but at the cost of higher latency. When in doubt, use General Purpose, which is also the default.
- High availability is achieved by having [mount targets in different subnets / availability zones](#).

[Back to top ↑](#)

## EFS Tips

- With EFS being based on NFSv4.1, any directory on the EFS can be mounted directly, it doesn't have to be the root directory.
- [User and group level permissions](#) can be used to control access to certain directories on the EFS file system.
-  **Sharing EFS filesystems:** One EFS filesystem can be used for multiple applications or services, but it should be aware of the following:

Pros:

- Because performance is based on total size of stored files, having everything on one drive will increase performance.
- application consuming credits faster than it can accumulate might be offset by another application that just rarely accesses them.

Cons:

- Since credits are shared, if one application over-consumes them, it will affect the others.
- A compromise is made with regards to [security](#): all clients will have to have network access to the drive. Some clients may be able to mount any directory on the EFS and they have read-write access to all files on the drive, including access to the applications hosted on other clients. There isn't a no-root-squash equivalent for EFS.

[Back to top](#) 

## EFS Gotchas and Limitations

- ◆ A number of NFSv4.1 features are [not supported](#) and there are some [limits](#) to the service.
- ◆ As of 2017-08, EFS offers disk level encryption for new drives. For file systems created before that date, encryption is achieved by moving the data to a new EFS volume.
- ◆ An EFS file system [can be mounted on premises](#) over Direct Connect.
- ◆ An EFS file system can NOT be mounted over VPC peering or VPN, even if the VPN is running on top of Direct Connect.
- ◆ Using an EFS volume on Windows is not supported.
-  When a file is uploaded to EFS, it can take hours for EFS to update the details for billing and burst credit purposes.
-  Metadata operations can be costly in terms of burst credit consumption. Recursively traversing a tree can easily ramp up to tens or even hundreds of megabytes of burst credits being consumed, even if no file is being modified.

## Load Balancers

---

### Load Balancer Basics

- AWS has 3 load balancing products - "Classic Load Balancers" (CLBs), "Application Load Balancers" (ALBs), and "Network Load Balancers" (NLBs).
- Before the introduction of ALBs, "Classic Load Balancers" were known as "Elastic Load Balancers" (ELBs), so older documentation and blog posts may still reference "ELBs".
- CLBs have been around since 2009, ALBs in 2016, NLBs were added in 2017 to AWS.
- CLBs support TCP and HTTP load balancing. ALBs support HTTP load balancing only. NLBs support TCP layer 4 load balancing.
- CLBs and ALBs can optionally handle termination for a single SSL certificate.
- All can optionally perform active health checks of instances and remove them from the destination pool if they fail.
- CLBs don't support complex / rule-based routing. ALBs support a (currently small) set of rule-based routing features, while NLBs support extensive routing options.
- CLBs can only forward traffic to a single globally configured port on destination instances, while ALBs can forward traffic to multiple ports on destination instances (configured on a per-instance basis, better supporting routing to services on shared clusters with dynamic port assignments like Docker Swarm or Mesos). NLBs support multiple ports on same IP; registering targets by IP address, including targets outside the VPC. ECS can select unused port for scheduling a task then register a target group using this port.
- CLBs are supported in EC2 Classic as well as in VPCs while ALBs are supported in VPCs only.
- ALBs can target groups of instances and IP based targets in the RFC1918 ranges allowing you to use on-premise servers via Direct Connect.

[Back to top ↑](#)

### Load Balancer Tips

- If you don't have opinions on your load balancing up front, and don't have complex load balancing needs like a large number of targets or specific routing requirements, it's reasonable just to use a CLB or ALB for load balancing instead.

- Even if you don't want to think about load balancing at all, because your architecture is so simple (say, just one server in front of it anyway). This gives you more flexibility when upgrading, since you won't have to change any DNS settings to propagate, and also it lets you do a few things like terminate SSL more easily.
- **CLBs and ALBs have many IPs:** Internally, an AWS load balancer is simply a collection of individual software load balancers running on EC2, with DNS load balancing traffic among them. The pool can contain many IPs, at least one per availability zone. They also support SSL termination, which is very convenient.
- **Scaling:** CLBs and ALBs can scale to very high throughput, but scaling up is not instantaneous. If you're expecting a spike in traffic suddenly, it can make sense to load test them so they scale up in advance. You can also [contact Amazon](#) about the load balancer.
- **Client IPs:** In general, if servers want to know true client IP addresses, load balancers must forward this information via the standard [X-Forwarded-For](#) header. When using a CLB as an HTTP load balancer, it's possible to get the client's IP address directly.
- **Using load balancers when deploying:** One common pattern is to swap instances in the load balancer after spinning up your latest version, keep old stack running for one or two hours, and either flip back to old stack in case of problems or proceed to the next step.
- **Rotating Certificates while retaining ARN:** Rotating IAM Server Certificates can be difficult as the standard practice is to delete the old certificate and then update all resources with the new ARN. You can however retain the same ARN using the `update-certificate` process:
  1. Upload a new IAM Server Certificate with a unique name (e.g fuzzy.com.new)
  2. Rename the existing IAM Server Certificate (e.g fuzzy.com to fuzzy.com.expired)
  3. Rename the new IAM Server Certificate to the name of the previously existing certificate (e.g fuzzy.com.new to fuzzy.com)
  4. Jiggle the CLB/ALB Listener to pick up the change:
    - ALB: Invoke `modify-listener` with the existing details for the ALB Listener
    - CLB: Invoke `create-load-balancer-listeners` with the existing details for the CLB listener

[Back to top](#) 

## Load Balancer Gotchas and Limitations

- ! CLBs and ALBs have **no fixed external IP** that all clients see. For most consumer apps this doesn't matter, but for some yours may want this. IPs will be different for each user, and will vary unpredictably for a single client over time (within [IP ranges](#)). And similarly, never resolve a CLB name to an IP and put it as the value of an A record — it will work for a short time but then fail.

- ! Some web clients or reverse proxies cache DNS lookups for a long time, which is problematic for CLBs and ALBs. This means after a few minutes, hours, or days, your client will stop working, unless you disable DNS caching in [settings](#) and be sure to [adjust them properly](#). Another example is nginx as a reverse proxy, which [normally resolves IP addresses](#) (although there is [a way to get around this](#)).
- ! It's not unheard of for IPs to be recycled between customers without a long cool-off period. So as a client, if you are not using SSL (to verify the server), you might get not just errors, but responses from completely different service providers.
- ♦ As an operator of a service behind a CLB or ALB, the latter phenomenon means you can also see puzzling or unexpected behavior from clients of other companies. This is most common with clients using back-end APIs (since web browsers typically do not cache API endpoints).
- ! CLBs and ALBs take time to scale up, it does not handle sudden spikes in traffic well. Therefore, if you anticipate a spike in traffic, "pre-warm" the load balancer by gradually sending an increasing amount of traffic.
- ! Tune your healthchecks carefully — if you are too aggressive about deciding when to remove an instance and then adding it back into the pool, the service that your load balancer is fronting may become inaccessible for seconds. Be extra careful about this when an autoscaler is configured to terminate instances that are marked as being unhealthy by the load balancer.
- ! CLB HTTPS listeners don't support Server Name Indication (SNI). If you need SNI, you can work around this limitation by providing a certificate with Subject Alternative Names (SANs) or by using TCP listeners and terminating SSL at your application layer.
- ♦ There is a limit on the number of ALBs, CLBs and NLBs per region (separately). As of late 2017, the default limit is 100 ALBs, 100 CLBs and 100 NLBs per region. These limits can be easily raised for ALB and CLB, but AWS is quite reluctant to raise the limit on NLBs.
- ♦ If using a Network Load Balancer (NLB) then EC2 clients cannot connect to an NLB that resides in another VPC or a managed VPN unless the EC2 client is a C5, i3.metal or M5 instance type. For VPC peering, both VPCs must be in the same region. See the [NLB Troubleshooting](#).

## CLB

---

### CLB Basics

-  [Homepage](#) · [User guide](#) · [FAQ](#) · [Pricing](#)
- Classic Load Balancers, formerly known as Elastic Load Balancers, are HTTP and TCP load balancers that are managed by Amazon.

## CLB Tips

- Best practices: [This article](#) is a must-read if you use CLBs heavily, and has a lot more detail.

## CLB Gotchas and Limitations

- In general, CLBs are not as “smart” as some load balancers, and don’t have fancy features or fine-grained control a load balancer would offer. For most common cases involving sessionless apps or cookie-based sessions over HTTP, they work well.
- ◆ By default, CLBs will refuse to route traffic from a load balancer in one Availability Zone (AZ) to a backend instance [cause 503s](#) if the last instance in an AZ becomes unavailable, even if there are healthy instances in other zones. If you’re using two backend instances per AZ, you almost certainly want to [enable cross-zone load balancing](#).
- ◆ Complex rules for directing traffic are not supported. For example, you can’t direct traffic based on a regular expression like HAProxy offers.
- Apex DNS names: Once upon a time, you couldn’t assign a CLB to an apex DNS record (i.e. example.com instead of www.example.com), because it needed to be an A record instead of a CNAME. This is now possible with a Route 53 alias record directed at a balancer.
- ◆ CLBs use [HTTP keep-alives](#) on the internal side. This can cause an unexpected side effect: Requests from different clients, which each have their own TCP connection on the external side, can end up on the same TCP connection on the internal side. Never assume requests on the same TCP connection are from the same client!
- ◆ Traffic between CLBs and back-end instances in the same subnet **will** have [Network ACL](#) rules evaluated (EC2-Classic subnet would not have Network ACL rules evaluated). If the default '0.0.0.0/0 ALLOW' rule is removed from the subnet, a rule that allows traffic on both the health check port and any listener port must be added.
- As of December 2016, CLBs launched in VPCs do not support IPv6 addressing. CLBs launched in EC2-Classic support [with the "dualstack" DNS name](#).

## ALB Basics

-  [Homepage](#) · [User guide](#) · [FAQ](#) · [Pricing](#)
-  **Websockets and HTTP/2 are [now supported](#).**
-  **Internet Protocol Version 6 (IPv6) is [now supported](#).**
-  **Load Balancing via IP is [now supported](#).**
- Prior to the Application Load Balancer, you were advised to use TCP instead of HTTP as the protocol to make it work and use [the obscure but useful Proxy Protocol](#) ([more on this](#)) to pass client IPs over a TCP load balancer.

[Back to top](#) 

## ALB Tips

- Use ALBs to route to services that are hosted on shared clusters with dynamic port assignment (like ECS or Mesos).
- ALBs support [HTTP host-based routing](#) (send HTTP requests for "api.mydomain.com" -> {target-group-1}, "blog.mydomain.com" -> {target-group 2}) as well as [HTTP path-based routing](#) (send HTTP requests for "/api/\*" -> {target-group-1}, "/blog/\*" -> {target-group-2}).

[Back to top](#) 

## ALB Gotchas and Limitations

- ♦ ALBs only support HTTP/2 over HTTPS (no plain-text HTTP/2).
- ♦ ALBs only support HTTP/2 to external clients and not to internal resources (instances/containers).
- ALBs support HTTP routing but not port-based TCP routing.
- Instances in the ALB's target groups have to either have a single, fixed healthcheck port ("EC2 instance"-level healthcheck port) or no healthcheck port at all. If a target has a healthcheck port, the healthcheck port for a target has to be the same as its application port ("Application instance"-level healthcheck port). The ALB supports per-target healthcheck port that is different than the application port.
- ALBs are VPC-only (they are not available in EC2 Classic)

- In a target group, if there is no healthy target, all requests are routed to all targets. For example, if you point a list containing a single service that has a long initialization phase (during which the health checks would fail), requests while it is still starting up.
-  Although ALBs [now support SNI](#), they only support 25 HTTPS certificates per Load Balancer. This limitation is might be subject to change.

## Elastic Beanstalk

---

### Elastic Beanstalk Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- EB (Elastic Beanstalk) is a PaaS (Platform as a Service) that helps developers create, deploy and scale web applications.
- EB handles deployment, configuration, provisioning, load balancing, auto-scaling, monitoring, and logging.
- EB creates AWS resources on your behalf but you retain full access and control of the underlying resources.
-  There is no cost to use EB but you will still be charged the full cost of the underlying AWS resources created.

[Back to top](#) 

### Elastic Beanstalk Tips

- To speed up deployment before launch or in a dev stage, turn off health checks and set the [Deployment policy](#).
- If you have a configuration you want to re-use for multiple EB apps, you can save the current configuration using `myEBConfig`.
- By default, EB doesn't have any alarms. You'll need to add them yourself on metrics that you're monitoring.
- By default, EB doesn't enable [managed platform updates](#). Enable them in configuration to have EB automatically pre-specified maintenance window.

[Back to top](#) 

### Elastic Beanstalk Gotchas and Limitations

- ◆ Don't edit [apache|nginx] conf files manually on ec2 instances as they will be re-written on each deployment
- ◆ After creating an EB environment, it's no longer possible to change the `Name` tag
- ◆ EB will sometimes quarantine instances that cause multiple deployment issues. Despite being quarantined, EB will still attempt to use them on subsequent deployments. To prevent this behavior, said instances will need to be terminated (or the underlying instance type changed)
- File uploads are capped at 10MB for most default eb configurations - update [nginx config](#) to change
- If you edit `.elasticbeanstalk/saved_configs/`, be aware that this is not kept in sync with the EB environment configuration. You will need to manually fetch and save for changes to take effect

## Elastic IPs

---

### Elastic IP Basics

-  [Documentation](#) · [FAQ](#) · [Pricing](#)
- Elastic IPs are static IP addresses you can rent from AWS to assign to EC2 instances.

[Back to top](#) 

### Elastic IP Tips

- ◆ **Prefer load balancers to elastic IPs:** For single-instance deployments, you could just assign elastic IP to an instance and consider that your deployment. Most of the time, you should provision a [load balancer](#) instead:
  - It's easy to add and remove instances from load balancers. It's also quicker to add or remove instances from a load balancer than to reassigned an elastic IP.
  - It's more convenient to point DNS records to load balancers, instead of pointing them to specific IPs you manage. You can also use Route 53 aliases, which are easier to change and manage.
  - But in some situations, you do need to manage and fix IP addresses of EC2 instances, for example if a custom application requires specific IP addresses. In such situations require elastic IPs.
- Elastic IPs are limited to 5 per account. It's possible to [request more](#).
- If an Elastic IP is not attached to an active resource there is a small [hourly fee](#).

- Elastic IPs are [no extra charge](#) as long as you're using them. They have a (small) cost when not in use, which is a deterrent to people from squatting on excessive numbers of IP addresses.

[Back to top ↑](#)

## Elastic IP Gotchas and Limitations

- ♦ There is [officially no way](#) to allocate a contiguous block of IP addresses, something you may desire when giving your network a range of IP addresses. Though when allocating at once, you may get lucky and have some be part of the same CIDR block. If this is important to you, you may want to [bring your own IP](#), which is more involved than this guide will go into.
- Unofficially, if you have Enterprise support, you can ask your account rep to try to allocate a block of Elastic IPs via AWS Support. For example, some of [Duo's fixed ranges](#) are [blocks of AWS IP space reassigned to AWS-DUOSECURITYINC](#). This is a good option to consider if you expect a denial.

## Glacier

---

### Glacier Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- **Glacier** is a lower-cost alternative to S3 when data is infrequently accessed, such as for archival purposes.
- It's only useful for data that is rarely accessed. It generally takes [3-5 hours](#) to fulfill a retrieval request.
- AWS [has not officially revealed](#) the storage media used by Glacier; it may be low-spin hard drives or even tapes.
- AWS has released an even more cost effective storage tier called [Glacier Deep Archive](#) that offers ~12 hour retrieval times for roughly a thousand dollars per month per petabyte.

[Back to top ↑](#)

### Glacier Tips

- You can physically [ship](#) your data to Amazon to put on Glacier on a USB or eSATA HDD.

[Back to top ↑](#)

## Glacier Gotchas and Limitations

- ♦ Getting files off Glacier is glacially slow (typically 3-5 hours or more).
- ♦ Due to a fixed overhead per file (you pay per PUT or GET operation), uploading and downloading many small objects will be very expensive. There is also a 32k storage overhead per file. Hence it's a good idea is to archive files before upload.
- 🎨 Be aware of the per-object costs of archiving S3 data to Glacier. [It costs \\$0.05 per 1,000 requests](#). If you have many objects of relatively small size, [it will take time to reach a break-even point](#) (initial archiving cost versus lower storage costs over time).

## Quicksight

---

### Quicksight Basics

-  [Homepage](#) · [User guide](#) · [Pricing](#)

[Back to top](#) 

### Quicksight Gotchas and Limitations

- ! Out of the box Quicksight is not able to access tables that are linked to a Schema in the AWS Glue Schema Registry. This is because the auto-generated IAM role `aws-quicksight-service-role-v0` does not have the necessary permissions. You can't provide a different IAM role when creating a dataset, but you can add more permissions to the role. The error message you will receive is an `SQL_EXCEPTION` with the following stack trace:

```
line 2:8: Column 'columnname' cannot be resolved where columnname is the first column in your table.
```
- ♦ Only QuickSight accounts that were created in the US East (N. Virginia) region can access the QuickSight Forum. This means that you can only have a QuickSight account in one region per AWS account.

## RDS

---

### RDS Basics

-  [Homepage](#) · [User guide](#) · [FAQ](#) · [Pricing](#) (see also [ec2instances.info/rds/](#))

- RDS is a managed relational database service, allowing you to deploy and scale databases more easily. It supports [Microsoft SQL Server](#), [PostgreSQL](#), [MySQL](#), [MariaDB](#), and Amazon's own [Aurora](#).
- RDS offers out of the box support for [high availability and failover](#) for your databases.

[Back to top](#) 

## RDS Tips

- If you're looking for the managed convenience of RDS for other data stores such as MongoDB or Cassandra, you can use third-party services from providers such as [Compose](#), or [InstaClstr](#).
- ◆ Make sure to create a new [parameter group](#) and option group for your database since the default parameter groups are shared across all databases.
- RDS instances start with a default timezone of UTC. If necessary, this can be [changed to a different timezone](#).
- Save costs by stopping the RDS instance(s) when not in use, you can achieve this by configuring a stop and start using the [Instance Scheduler](#) solution.

[Back to top](#) 

## RDS Gotchas and Limitations

- 🚧 RDS instances run on EBS volumes (either general-purpose or provisioned IOPS), and hence are constrained by the performance of EBS.
- ◆ Verify what database features you need, as not everything you might want is available on RDS. For example, it's not possible to use all MySQL features on RDS. You can check the list of [supported features and extensions](#). If the features you need aren't supported by RDS, you'll have to implement them yourself.
- ◆ If you use the failover support offered by RDS, keep in mind that it is based on DNS changes, and make sure to propagate these changes appropriately. This is particularly important for Java, given how its DNS resolver's TTL is [configured](#).
- ◆ **DB migration to RDS:** While importing your database into RDS ensure you take into consideration the maintenance window. If a backup is running at the same time, your import can take a considerably longer time than you would have expected.
- [Database sizes are limited](#) to 6TB for all database engines except for SQL Server which has a 4TB limit and Aurora which has a 64TB limit.

# RDS MySQL and MariaDB

---

## RDS MySQL and MariaDB Basics

- RDS offers MySQL versions 5.5, 5.6, 5.7 and 5.8.
- RDS offers MariaDB versions 10.0, 10.1, 10.2 and 10.3.

[Back to top](#) 

## RDS MySQL and MariaDB Tips

- MySQL RDS allows access to [binary logs](#).
- Multi-AZ instances of MySQL transparently replicate data across AZs using DRBD. Automated backups of multi-[backup instance](#) to reduce latency spikes on the primary.
- ♦ **Performance Schema:** While [Performance Schema](#) is enabled by default in MySQL 5.6.6 and later, it is disabled on RDS. If you wish to enable Performance Schema, a reboot of the RDS instance will be required.
- ♦ **MySQL vs MariaDB vs Aurora:** If you prefer a MySQL-style database but are starting something new, you probably want Aurora and MariaDB as well. Aurora has increased availability and is the next-generation solution. That said, Aurora is faster than MySQL for certain workloads. MariaDB, the modern [community fork](#) of MySQL, [likely now has the enterprise features](#) you need for mission-critical purposes and is supported by RDS.

[Back to top](#) 

## RDS MySQL and MariaDB Gotchas and Limitations

- ♦ **No SUPER privileges.** RDS provides some [stored procedures](#) to perform some tasks that require SUPER privileges, such as stopping replication.
- ♦ You can replicate to non-RDS instances of MySQL, but [replication to these instances will break during AZ failover](#).

- ◆ There is no ability to manually CHANGE MASTER on replicas, so they must all be rebuilt after a failover of the master.
- ◆ Most global options are exposed only via [DB parameter groups](#). Some variables that were introduced in later versions as [avoid\\_temporal\\_upgrade](#) in MySQL 5.6.24 are not made available in RDS's 5.6.x parameter group and making an upgrade to MySQL 5.7.x.
- ◆ RDS features such as Point-In-Time restore and snapshot restore are not supported on MyISAM tables. Ensure your MyISAM table before executing a snapshot or backup operation to ensure consistency.

## RDS PostgreSQL

---

### RDS PostgreSQL Basics

- RDS offers PostgreSQL 9.3, 9.4, 9.5, 9.6, and 10.

[Back to top](#) 

### RDS PostgreSQL Tips

- Recently Logical Replication is being supported, [both as subscriber and publisher](#).
- Supports a relatively large range of native [extensions](#).
- RDS PostgreSQL 10 Supports native partitioning and most of the major features and tunables.
- Supports connections over SSL.
- Supports multi A-Z and Point-in-time recovery.

[Back to top](#) 

### RDS PostgreSQL Gotchas and Limitations

- No superuser privileges. RDS provides a role `rds_superuser` that can do most of the needed operations but the full superuser privilege is missing.
- Some major features are delayed compared to open source PostgreSQL.
- By default RDS is spec'd with general purpose SSD, if you need better performance you have to spec provisioned IOPS.
- You can't use RDS as a replica outside RDS without using logical replication.

- There are settings that cannot be changed and most of the settings that can change can only be changed using groups.
- It's harder to troubleshoot performance problems since you have no access to the host.
- Be sure to verify that all the [extensions](#) you need are available. If you are using an extension not listed there, you may need a work around, or deploy your own database in EC2.
- Many Postgres utilities and maintenance items expect command line access, that can usually be satisfied by using the AWS Lambda interface.

## RDS SQL Server

---

### RDS SQL Server Basics

- [RDS offers SQL Server 2008 R2, 2012, 2014, 2016 and 2017](#) including Express, Web, Standard and Enterprise.

[Back to top](#) 

### RDS SQL Server Tips

- Recently added support for [backup and restore to/from S3](#) which may make it an attractive [DR option](#) for on-premises environments.

[Back to top](#) 

### RDS SQL Server Gotchas and Limitations

- ♦ The user is granted only db\_owner privileges for each database on the instance.
- ♦ Storage cannot be expanded for existing databases. If you need more space, you must restore your database from a backup to larger storage.
- ♦ There is a **16TB** database size limit for non-Express editions. There is also a minimum storage size, 20GB for Web, Standard and Enterprise.
- ♦ Limited to [30 databases per instance](#)

## RDS Aurora Basics

Aurora is a cloud only database service designed to provide a distributed, fault-tolerant relational database with self-scaling up to 64TB per instance. It currently comes in two versions, a MySQL compatible system, and a PostgreSQL

## RDS Aurora MySQL

---

### RDS Aurora MySQL Basics

- Amazon's proprietary fork of MySQL intended to scale up for high concurrency workloads. Generally speaking, i performance under Aurora is not expected to improve significantly relative to MySQL or MariaDB, but Aurora is performance while executing many more queries concurrently than an equivalent MySQL or MariaDB server cou
- Notable new features include:
  - Log-structured storage instead of B-trees to improve write performance.
  - Out-of-process buffer pool so that databases instances can be restarted without clearing the buffer pool.
  - The underlying physical storage is a specialized SSD array that automatically maintains 6 copies of your data
  - Aurora read replicas share the storage layer with the write master which significantly reduces replica lag, eli master to write and distribute the binary log for replication, and allows for zero-data-loss failovers from the master and all the read replicas that share storage are known collectively as an **Aurora cluster**. Read replica

[Back to top](#) 

### RDS Aurora MySQL Tips

- In order to take advantage of Aurora's higher concurrency, applications should be configured with large databases should execute as many queries concurrently as possible. For example, Aurora servers have been tested to prod on some OLTP workloads with up to 5,000 connections.
- Aurora scales well with multiple CPUs and may require a large instance class for optimal performance.

- The easiest migration path to Aurora is restoring a database snapshot from MySQL 5.6 or 5.7. The next easiest migration path is creating a new Aurora instance and setting it up as a replica of your existing database. If none of those methods are options, Amazon offers a migration service.
- You can replicate [from an Aurora cluster to MySQL](#) or [to another Aurora cluster](#). This requires binary logging to be enabled and may not be as performant as native Aurora replication.
- Because Aurora read replicas are the [equivalent of a multi-AZ backup](#) and they can be configured as zero-data-loss, there are fewer scenarios in which the creation of a multi-AZ Aurora instance is required.

[Back to top](#) 

## RDS Aurora MySQL Gotchas and Limitations

-  [Aurora 1.x is based on MySQL 5.6.x](#) with some cherry-picking of later MySQL features. It is missing most 5.7 features and lacks online DDL features introduced in 5.6.17.
-  [Aurora 2.x is based on MySQL 5.7.x](#)
- Aurora does not support GTID transactions in either the 5.6/Aurora 1.x or the 5.7/Aurora 2.x release lines.
- Aurora maximum cluster size is 64 TB

## RDS Aurora PostgreSQL

---

### RDS Aurora PostgreSQL Basics

- Amazon's proprietary fork of PostgreSQL, intended to scale up for high concurrency workloads while maintaining compatibility with PostgreSQL 9.6.
- Higher throughput (up to 3x with similar hardware).
- Automatic storage scale in 10GB increments up to 64TB.
- Low latency read replicas that share the storage layer with the master which significantly reduces replica lag.
- Point in time recovery.
- Fast database snapshots.

[Back to top](#) ↑

## RDS Aurora PostgreSQL Tips

- Aurora Postgres by default is supposed to utilize high connection rates and for this reason connection pooling must be used accordingly.
- Because Aurora is based on PostgreSQL 9.6, it lacks features like declarative partitioning or logical replication.

[Back to top](#) ↑

## RDS Aurora PostgreSQL Gotchas and Limitations

- Aurora PostgreSQL falls behind normal RDS when it comes to available versions, so if you need features from the latest PostgreSQL version you might be better off with plain RDS.
- Patching and bug fixing is separate from open source PostgreSQL.

## ElastiCache

---

### ElastiCache Basics

-  [Homepage](#) · [User guide for Redis](#) · [User guide for Memcached](#) · [FAQ](#) · [Pricing](#)
- ElastiCache is a managed in-memory cache service, that can be used to store temporary data in a fast in-memory store to avoid repeating the same computation multiple times when it could be reused.
- It supports both the [Memcached](#) and [Redis](#) open source in-memory cache software and exposes them both using standard protocols.
- The main benefit is that AWS takes care of running, patching and optimizing the cache nodes for you, so you just need to connect to the service and configure its endpoint in your application, while AWS will take of most of the operational work of running the service.

[Back to top](#) ↑

## ElastiCache Tips

- Choose the [engine](#), clustering configuration and [instance type](#) carefully based on your application needs. The documentation will detail the pros, cons and limitations of each engine in order to help you choose the best fit for your application. Redis is preferable for storing more complex data structures, while Memcached is just a plain key/value store. The simplified API makes it to be slightly faster and allows it to scale out if needed, but Redis has more features which you may use in your application.
- For Memcached AWS provides enhanced SDKs for certain programming languages which implement [auto-discovery](#) and failover logic which is not available in the normal memcached client libraries.

[Back to top](#) 

## ElastiCache Gotchas and Limitations

- Since in some cases changing the cache clusters may have some restrictions, like for [scaling](#) purposes, it may be necessary to replace the entire cluster. If the cluster was created using CloudFormation, it is recommended to avoid launching new clusters in the same stack as the existing ones. Instead, it is recommended to launch ElastiCache clusters in separate stacks (with similar constraints) in dedicated stacks which can be replaced entirely with new stacks having the required configurations.

## DynamoDB

---

### DynamoDB Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- DynamoDB is a [NoSQL](#) database with focuses on speed, flexibility, and scalability.
- DynamoDB is priced on a combination of throughput and storage.

[Back to top](#) 

### DynamoDB Alternatives and Lock-in

-  Unlike the technologies behind many other Amazon products, DynamoDB is a proprietary AWS product without an open source alternative. While there are several NoSQL databases available, they may not offer the same level of performance or consistency guarantees as DynamoDB. If you tightly couple your application to its API and featureset, it may be difficult to switch to a different database without significant changes to your codebase.

- The most commonly used alternative to DynamoDB is [Cassandra](#).

[Back to top ↑](#)

## DynamoDB Tips

- There is a [local version of DynamoDB](#) provided for developer use.
- [DynamoDB Streams](#) provides an ordered stream of changes to a table. Use it to replicate, back up, or drive event processing.
- DynamoDB can be used [as a simple locking service](#).
- DynamoDB indexing can include **primary keys**, which can either be a single-attribute hash key or a composite key, and you can also query non-primary key attributes using [secondary indexes](#).
- **Data Types:** DynamoDB supports three [data types](#) – number, string, and binary – in both scalar and multi-value formats. It also supports [JSON](#).
- As of late 2017, DynamoDB supports both [global tables](#) and [backup / restore functionality](#).

[Back to top ↑](#)

## DynamoDB Gotchas and Limitations

- ◆ DynamoDB doesn't provide an easy way to bulk-load data (it is possible through [Data Pipeline](#)) and this has [consequences](#). Since you need to use the regular service APIs to update existing or create new rows, it is common to run into performance issues when trying to bulk-load data. One way to mitigate this is to increase the destination table's write throughput to speed import. But when the table's write capacity is increased, DynamoDB will automatically split the partitions underlying the table, spreading the total table capacity evenly across the new generation of partitions. This leaves the table in a state where it is much easier for hotspots to overwhelm individual partitions.
- ◆ It is important to make sure that DynamoDB [resource limits](#) are compatible with your dataset and workload. For example, the maximum size value that can be added to a DynamoDB table is 400 KB (larger items can be stored in S3 and a User-defined function can be used to handle them).
- ◆ Dealing with **time series data** in DynamoDB can be challenging. A global secondary index together with document-level access control can be a possible solution as explained [here](#).
- ◆ When setting up [fine grained policies](#) for access to DynamoDB tables, be sure to include their secondary indexes in the policy document as well.

## ECS Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- ECS (EC2 Container Service) is a relatively new service (launched end of 2014) that manages clusters of services on EC2.
- See the [Containers and AWS](#) section for more context on containers.
- ECS is growing in adoption, especially for companies that embrace microservices.
- Deploying Docker directly in EC2 yourself is another common approach to using Docker on AWS. Using ECS is not (yet) seem to be the predominant way many companies are using Docker on AWS.
- It's also possible to use [Elastic Beanstalk with Docker](#), which is reasonable if you're already using Elastic Beanstalk.
- Using Docker may change the way your services are deployed within EC2 or Elastic Beanstalk, but it does not radically change how other services are used.
- [ECR](#) (EC2 Container Registry) is Amazon's managed Docker registry service. While simpler than running your own registry, it lacks some features that might be desired by some users:
  - Doesn't support cross-region replication of images.
    - If you want fast fleet-wide pulls of large images, you'll need to push your image into a region-local registry.
  - Doesn't support custom domains / certificates.
- A container's health is monitored via [CLB](#) or [ALB](#). Those can also be used to address a containerized service. When using ALB, you need to handle port contention (i.e. services exposing the same port on the same host) since an ALB's target groups can't address ECS-based services directly.
- [The Hitchhiker's Guide to AWS ECS and Docker](#) by J. Cole Morrison is an excellent article for Introduction to AWS ECS.

[Back to top](#) 

## ECS Tips

- **Log drivers:** ECS supports multiple log drivers (awslogs, splunk, fluentd, syslog, json, ...). Use [awslogs](#) for CloudWatch Logs (CloudWatch Metrics is made for the logs first). [Drivers such as fluentd are not enabled by default](#). You can, install the agent and enable it by setting `ECS_AVAILABLE_LOGGING_DRIVERS='["awslogs","fluentd"]'` to `/etc/ecs/ecs.config`.

- [This blog from Convoy](#) (and [commentary](#)) lists a number of common challenges with ECS as of early 2016.
- It is possible to optimize disk clean up on ECS. By default, the unused containers are deleted after 3 hours and 30 minutes. These settings can be changed by adding `ECS_ENGINE_TASK_CLEANUP_WAIT_DURATION=10m` and `ECS_IMAGE_CLEANUP_WAIT_DURATION=10m` to `/etc/ecs/ecs.config`. [More information on optimizing ECS disk cleanup.](#)

[Back to top ↑](#)

## ECS Alternatives and Lock-in

- [Kubernetes](#): Extensive container platform. Available as a hosted solution on [Google Cloud](#), [AWS](#), [Azure](#), [DigitalOcean](#), etc.
- [Nomad](#): Orchestrator/Scheduler, tightly integrated in the HashiCorp stack (Consul, Vault, etc).

 [Please help expand this incomplete section.](#)

## EKS

---

### EKS Basics

-  [Homepage](#) · [User guide](#) · [FAQ](#) · [Pricing](#)
- EKS (Elastic Kubernetes Service) is a new service (launched June 2018) that provides managed Kubernetes Masters to deploy K8s Services and Pods on top of EC2 based Kubernetes nodes.
- See the [Containers and AWS](#) section for more context on containers.
- EKS is AWS's solution to hosting Kubernetes natively on AWS. It is not a replacement for ECS directly but is in response to the dominance of Kubernetes.
- EKS does not launch EC2 nodes and would have to be configured and setup either manually or via Cloudformation (CloudFormation solution)
- EKS management is done through a utility called kubectl, and with Kube configuration files. These files will need to be exchanged with the K8s Master with a certificate and URL. The AWS CLI can autogenerate the configuration file that kubectl uses for communicating with the cluster.<sup>1</sup>

- EKS authentication is integrated with IAM roles/permissions. The AWS CLI has an integrated sub-command for generating tokens.<sup>2</sup> This was formerly done via a custom plugin for kubectl called [aws-iam-authenticator](#) (formerly heptio-authenticator).
- EKS provides [Calico](#) from Tigera for securing workloads within a cluster using Kubernetes network policy.

[Back to top](#) 

## EKS Tips

- Multiple clusters can be supported by using different kubeconfig files.
- AWS has a [Kubernetes Quickstart](#) developed in collaboration with Heptio.

[Back to top](#) 

## EKS Alternatives and Lock-in

- [ECS](#): Amazon's native Container Scheduled platform released in 2014. If you don't utilise containers today and aren't planning to, ECS is an excellent product.
- [Kubernetes](#): Extensive container platform. Available as a hosted solution on [Google Cloud](#), [AWS](#), [Digital Ocean](#) and many others.
- [Nomad](#): Orchestrator/Scheduler, tightly integrated in the HashiCorp stack (Consul, Vault, etc).

[Back to top](#) 

## EKS Gotchas and Limitations

- Pods and Service configurations can rapidly consume IP addresses inside a VPC. Proper care and maintenance should be taken to ensure IP exhaustion does not occur.
- There is currently no integrated monitoring in CloudWatch for EKS pods or services, you will need to deploy a monitoring solution like CloudWatch Metrics or CloudWatch Metrics Insights. Prometheus also supports Kubernetes such as Prometheus.
- Autoscaling based off CPU/Memory of a node is limited as you will not be aware of pending Services/Pods that require resources. An [autoscaler](#) can be useful for scaling based on Node resource usage and unschedulable Pods.
- [Prometheus](#) is a very popular monitoring solution for K8s, metrics and alerts can be used to send events to Lambda functions to take autoscaling actions.

## Footnotes

1: <https://docs.aws.amazon.com/eks/latest/userguide/create-kubeconfig.html>

2: <https://aws.amazon.com/about-aws/whats-new/2019/05/amazon-eks-simplifies-kubernetes-cluster-authentication/>

## Fargate

---

### Fargate Basics

-  [Homepage](#) · [FAQ](#) · [Pricing](#)
- Fargate allows you to manage and deploy containers without having to worry about running the underlying compute infrastructure.
- Fargate serves as a new backend (in addition to the legacy EC2 backend) on which ECS and EKS tasks can be run.
- Fargate and EC2 backends are called "Launch Types"
- Fargate allows you to treat containers as fundamental building blocks of your infrastructure

### Fargate Tips

- Fargate follows a similar mindset to Lambda, which lets you focus on applications, instead of dealing with underlying infrastructure.
- Fargate is supported by CloudFormation, aws-cli and ecs-cli
- Fargate tasks can be launched alongside tasks that use EC2 Launch Type
-  Before creating a large Fargate deployment, make sure to estimate costs and compare them against alternatives. A traditional EC2 deployment - Fargate prices can be several times those of equivalently-sized EC2 instances. To evaluate the cost of a Fargate deployment based on potential costs, refer to pricing for [EC2](#) and [Fargate](#).

-  [Azure Container Instances](#): Available on Microsoft Azure in preview version, allows to run applications in containers without managing virtual machines

[Back to top](#) 

## Fargate Gotchas and Limitations

- The smallest resource values that can be configured for an ECS Task that uses Fargate is 0.25 vCPU and 0.5 GB of memory.
- [Task storage is ephemeral. After a Fargate task stops, the storage is deleted.](#)

## Lambda

---

### Lambda Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- Lambda is AWS' serverless compute offering, allowing users to define Lambda functions in a selection of runtimes and environments. Lambda supports a variety of triggers, including SNS notifications and API Gateway invocations. Lambda is the key service that enables the serverless [architecture on AWS](#), alongside AWS API Gateway, AWS Batch, and AWS DynamoDB.

[Back to top](#) 

### Lambda Tips

- The idea behind 'serverless' is that users don't manage provisioning, scaling, or maintenance of the physical machines that run their application code. With Lambda, the machine that actually executes the user-defined function is abstracted as a function. When creating a Lambda function, users are able to declare the amount of memory available to the function, which directly affects the performance and cost of the function.
- Changing the amount of memory available to your Lambda functions also affects the amount of [CPU power](#) available to the function.
- While AWS does not offer hard guarantees around container reuse, in general it can be expected that an unaltered Lambda function will reuse a warm (previously used) container if called shortly after another invocation. Users can use this as a way to reduce cold start times by smartly caching application data on initialization.

- A Lambda that hasn't been invoked in some time may not have any warm containers left. In this case, the Lambda and initialize the Lambda code in a 'cold start' scenario, which can add significant latency to Lambda invocations performance [has improved significantly over the 2018-2019 timeframe](#) and is now typically in the range of 200-300ms function depending on the language runtime.
- Lambda functions running insides of VPCs have also seen [recent improvements](#) to cold start times. Previously the Lambda would have cold starts of ~15 seconds; now those same functions cold start in < 1 second.
- There are a few strategies to avoiding or mitigating cold starts. [Provisioned concurrency](#) was announced at re:Invent 2018 as an effective means to eliminating cold starts. Other techniques include keeping containers warm by periodic triggering of lightweight runtimes such as Node as opposed to Java.
- Lambda is integrated with AWS CloudWatch and provides a logger at runtime that publishes CloudWatch events to CloudWatch Logs.
- Lambda offers out-of-the-box opt-in support for AWS X-Ray. X-Ray can help users diagnose Lambda issues by capturing metrics and traces of their Lambda's execution flow. This is especially useful when investigating issues calling other AWS services as it provides a detailed and easy-to-parse [visualization of the call graph](#).
- Using [timed CloudWatch events](#), users can use Lambda to run periodic jobs in a cron-like manner.
- Events sent to Lambda that fail processing can be managed using a [Dead Letter Queue \(DLQ\) in SQS](#).
- More on serverless:
  - [Mike Roberts's thoughts on martinfowler.com](#).
  - [AWS Serverless Application Model \(SAM\)](#), a simplification built on top of CloudFormation that can help to define serverless applications using Lambda.
  - [Serverless](#), one of the most popular frameworks for building serverless applications using AWS Lambda and AWS Step Functions.
  - [Other helpful frameworks](#).

[Back to top ↑](#)

## Lambda Alternatives and Lock-in

-  Other clouds offer similar services with different names, including [Google Cloud Functions](#), [Azure Functions](#), and [Amazon Lambda](#). If you are running Kubernetes another Lambda alternative is [OpenFaaS](#)

## Lambda Gotchas and Limitations

- ◆ Testing Lambdas, locally and remotely, can be a challenge. Several tools are available to make this easier, including supported [SAM Local](#).
- ◆ Managing lots of Lambda functions is a workflow challenge, and tooling to manage Lambda deployments is limited.
- ◆ AWS' official workflow around managing function [versioning and aliases](#) is painful. One option is to avoid Lambda and abstracting your deployment workflow outside of Lambda. One way this can be accomplished is by deploying your Lambda function to multiple AWS accounts across successive stages, with a distinct AWS account per stage, where each account only needs to be aware of the latest version of the Lambda function. This approach requires external tooling to handle the deployment and updates.
- ◆ While adding/removing S3 buckets as triggers for Lambda function, this error may occur: "There was an error: Configuration is ambiguously defined. Cannot have overlapping suffixes in two rules if the prefixes are overlapping type." In this case, you can manually remove the Lambda event in the "Events" tab in the "Properties" section of the S3 bucket.
- ◆ Managing the size of your deployment artifact can be a challenge, especially if using Java. Options to mitigate this include using smaller jars or using Java's built-in mechanism for loading dependencies at runtime into /tmp.
- When using DynamoDB as a trigger for your Lambda functions, this error may occur: "PROBLEM: internal Lambda error. Please contact Lambda customer support." This usually just means that Lambda can't detect anything in the DynamoDB stream. If the issue persists, deleting and recreating your trigger may help.
- ◆ If your lambda needs access to resources in a VPC (for example ElastiCache or RDS), it will need to be deployed in a VPC. This will increase cold-start times as an Elastic Network Interface (ENI) will have to be registered within the VPC for each Lambda function. AWS also has a relatively low initial limit (350) on the number ENI's that can be created within an VPC, however this can be increased up to 1000s if a good case is made to AWS support.
- ◆ If your lambda is in a VPC and needs access to resources outside VPC such as outbound internet access, it will need to be assigned a public IP. Lambdas won't be assigned a public IP irrespective of subnet. If it is accessing a service like S3, you may use VPC endpoints.
- ◆ Lambda has several [resource limits](#) as of 2017-06:
  - A **6MB** request or response payload size.
  - A **50 MB** limit on the compressed .zip/jar file deployment package size.
  - A **250 MB** limit on the code/dependencies in the package before compression.

- A 500 MB limit on local storage in /tmp.

[Back to top](#) 

## Lambda Code Samples

- [Fan-out](#) is an example of using Lambda to “fan-out” or copy data from one service, in this case Kinesis, to multiple Destinations for fan-out data in the sample include IoT, SQS and more.
- This [AWS limit monitor using Lambdas](#) shows use of multiple Lambdas for monitoring.
- This [Lambda ECS Worker Pattern](#) shows use of Lambda in a workflow where data from S3 is picked up by the Lambda function and then sent to ECS for more processing.
- The [Secure Pet Store](#) is a sample Java application which uses Lambda and API Gateway with Cognito (for user identity).
- [aws-lambda-list](#) is a list of “hopefully useful AWS lambdas and lambda-related resources”. Quite a few code samples are guaranteed tested. Caveat Emptor.

 [Please help expand this incomplete section.](#)

## API Gateway

---

### API Gateway Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- **API Gateway** provides a scalable, secured front-end for service APIs, and can work with Lambda, Elastic Beanstalk, and CloudFront.
- It allows “serverless” deployment of applications built with Lambda.
- ♦ Switching over deployments after upgrades can be tricky. There are no built-in mechanisms to have a single API gateway point to multiple Lambda functions. So it may be necessary to build an additional layer in front (even another API Gateway) to handle migration from one deployment to another.

[Back to top](#) 

### API Gateway Alternatives and Lock-In

- [Kong](#) is an open-source, on-premises API and microservices gateway built on nginx with Lua. Kong is extensible.
- [Tyk](#) is an open-source API gateway implemented in Go and available in the cloud, on-premises or hybrid.

[Back to top](#) 

## API Gateway Tips

- ♦ Prior to 2016-11, you could only send and receive plain text data (so people would base64-encode binary data).
- API Gateway supports the OpenAPI specification (aka [Swagger](#)). This allows you to describe your API in a language that various tools to generate code supporting your API.
- Generating clients is extremely easy, either through the AWS console or using the get-sdk API.
- API Gateway integrates with CloudWatch out-of-the-box, allowing for easy logging of requests and responses.
  - Note that if your request or response are too large, CloudWatch will truncate the log. For full request/reply logs, do so in your integration (e.g. Lambda).
  - A good practice when calling API Gateway APIs is to log the request ID on the client. You can later refer to the CloudWatch logs for easier tracing and debugging.
- There are multiple ways to secure your API, including built-in support for [AWS Cognito](#). For most use-cases, Cognito is the simplest way to authenticate users.
  - Although you can roll your own solution using a [custom authorizer](#), which is basically a Lambda you define that checks if the request is acceptable or not.
- While API Gateway lends itself well to REST-style development, it's perfectly reasonable to implement an RPC-style API as well. Depending on your use-case, this can often lead to a much simpler API structure and smoother client experience.
  - RPC-style APIs are particularly useful when designing services that sit deeper in the stack and don't serve client requests directly.

[Back to top](#) 

## API Gateway Gotchas and Limitations

- ♦ API Gateway only supports encrypted (https) endpoints, and does not support unencrypted HTTP. (This is probably a good thing.)

- ♦ API Gateway doesn't support multi-region deployments for high availability. It is a service that is deployed in one region with a global endpoint that is served from AWS edge locations (similar to a CloudFront distribution). You cannot have multiple API Gateways with the same hostname in different AWS regions and use Route 53 to distribute the traffic. More in [this blog post](#).
- ♦ Integration timeout: All of the various integration types (eg: Lambda, HTTP) for API Gateway have timeouts, and while there are some limits, these timeouts can't be increased.
- ♦ API Gateway returns a 504 status code for any network or low level transport related issue. When this happens, check the CloudWatch logs for the request that includes the message: Execution failed due to an internal error. The reason for this error is that even though your backend server is up and running, it may be doing something outside of the API Gateway's control (such as sending well-formed chunked messages). You can test by hitting your backend directly with the curl --raw -s command and seeing if it complains.
- ♦ AWS X-Ray support exists but cumbersome to use. If you have other AWS services calling API Gateway, your X-Ray trace will end at the API Gateway boundary. API Gateway will also not appear as a node in your service map. [More here](#).
- ♦ Be careful using the export feature. The resulting Swagger template is often incomplete and doesn't integrate well with other tools like AWS Lambda's API Gateway extension or AWS CloudFormation's API Gateway extension.
- ♦ Many changes to API Gateway resources need to be 'deployed' via console or API call. Unfortunately, API Gateway does not provide any notifications for staged changes. If you're changing a resource and it's not taking effect, there's a decent chance you just need to deploy it.
  - In particular, when deploying an API Gateway as part of a CloudFormation stack, changes will not automatically be applied until the deployment resource itself was changed. You can change work around this by always changing the deployment resource during a CloudFormation update, or running a custom resource that ensures the deployment is made.
  - Alternatively, by using the [Serverless Application Model](#) definition for an API Gateway resource, you can always have a new deployment every time the stack is updated.
- ♦ API Gateway does not support nested query parameters on method requests.
- ♦ API Gateway limits number of resources to 300, as described [here](#). This is something to be considered when using API Gateway as a platform where your team/organization deploys to the same API Gateway.

 [Please help expand this incomplete section.](#)

# Step Functions

---

## Step Functions Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- Step Functions is AWS' way to create state machines that manage a serverless workflow.

[Back to top](#) 

## Step Functions Tips

- A variety of structures are supported including branching, parallel operations and waits
- [Tasks](#) represent the real work nodes and are frequently Lambda functions, but can be [Activities](#) which are external implemented any way you like.
- State machines have [data](#) that "flows" through the steps and can be modified and added to as the state machine
- It's best if your tasks are idempotent, in part because you may want to re-run the state machine with the same in debugging
- The AWS Console facilitates your examining the execution state at various steps.
  - The console lets you do this with a few steps:
    - select the "input" tab from the failed execution
    - copy the input data (JSON)
    - select the state machine name in the breadcrumbs
    - start a new execution, pasting the input data you copied previously

[Back to top](#) 

## Step Functions Gotchas and Limitations

- Step Functions are free tier eligible up to an initial 4000 transitions per month. Thereafter, the charge is \$0.025 per-second per transition.
- You can have many, simultaneous, executions, but be aware of lambda throttling limits. This has been per-account became settable per-lambda.

- Step Function executions are limited to 25,000 events. Each step creates multiple events. This means that [iterating](#) is limited to an iteration count of around 3000 before needing to [continue as a new execution](#).

## Route 53

---

### Route 53 Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- Route 53 is AWS' DNS service.

[Back to top](#) 

### Route 53 Alternatives and Lock-In

- Historically, AWS was slow to penetrate the DNS market (as it is often driven by perceived reliability and long-term stability), but Route 53 has matured and [is becoming the standard option](#) for many companies. Route 53 is cheap by historical standards and has a fairly large global network with geographic DNS and other formerly "premium" features. It's convenient if you want to use AWS services like CloudFront or Lambda@Edge.
-  Generally you don't get locked into a DNS provider for simple use cases, but increasingly become tied in once you start using more advanced features like geographic routing or Route 53's alias records.
-  Many alternative DNS providers exist, ranging from long-standing premium brands like [UltraDNS](#) and [Dyn](#) to more modestly priced brands like [DNSMadeEasy](#). Most DNS experts will tell you that the market is opaque enough that performance don't really correlate well with price.
-  Route 53 is usually somewhere in the middle of the pack on performance tests, e.g. the [SolveDNS reports](#).

[Back to top](#) 

### Route 53 Tips

-  Know about Route 53's "alias" records:
  - Route 53 supports all the standard DNS record types, but note that [alias resource record sets](#) are not standard and specific to Route 53.

- Aliases are like an internal name (a bit like a CNAME) that is resolved internally on the server side. For example, you could have a CNAME to the DNS name of a CLB or ALB, but it's often better to make an alias to the same load balancer. In this case, same, but in the latter case, externally, all a client sees is the target the record points to.
  - It's often wise to use alias record as an alternative to CNAMEs, since they can be updated instantly with an API without waiting for DNS propagation.
  - You can use them for CLBs/ALBs or any other resource where AWS supports it.
  - Somewhat confusingly, you can have CNAME and A aliases, depending on the type of the target.
  - Because aliases are extensions to regular DNS records, if exported, the output [zone file](#) will have additional information in it.
- [Latency-based routing](#) allows users around the globe to be automatically directed to the nearest AWS region without having to specify a preference, so that latency is reduced.
  - Understand that domain registration and DNS management (hosted zones) are two separate Route 53 services. When you register a domain, Route 53 automatically assigns four name servers to it (e.g. ns-2.awsdns-00.com). Route 53 also offers the ability to host a private hosted zone for DNS management, but you are not required to do your DNS management in the same account or you can choose to point to the servers assigned to your domain in Route 53.
  - One use case would be to put your domain registration (very mission critical) in a [bastion account](#) while managing your application's DNS in another account which is accessible by your applications.

[Back to top](#) 

## Route 53 Gotchas and Limitations

- ♦ Private Hosted Zone will only respond to DNS queries that originate from within a VPC. As a result Route 53 will not respond to queries from outside the VPC unless they are made via a VPN or Direct connect. To get around this you will need to implement [Hybrid Cloud DNS Solutions](#) or use the provided IP addresses to query the hosted zone.

## CloudFormation

---

### CloudFormation Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#) at no additional charge
- **CloudFormation** allows you to manage sets of resources from other AWS services grouped into [stacks](#). CloudFormation defines these stacks in a template using [JSON](#) or [YAML](#). CloudFormation is one of the major services underpinning [code capabilities](#) and is crucial in enabling repeatable and consistent deployments of infrastructure.
-  CloudFormation itself has [no additional charge](#); you pay for the underlying resources.

[Back to top](#) 

## CloudFormation Alternatives and Lock-In

- HashiCorp's [Terraform](<https://www.hashicorp.com/terraform/intro/vs/cloudformation.html>) is a third-party alternative that can support multiple platforms/providers including [Azure](#) and [OpenStack](#).
- ◆ Some AWS features may not be available in Terraform (e.g. multi-AZ ElastiCache using Redis), and you may have to use CloudFormation templates.
- [Pulumi](#) enables teams to define and deliver Cloud Native Infrastructure as Code on any cloud, with any language from serverless to Kubernetes to infrastructure.

[Back to top](#) 

## CloudFormation Tips

- Validate your stack in a different AWS account! CloudFormation truly shines when making multiple deployments across different accounts and regions. A common practice is to deploy stacks in successive stages ending in a production environment.
- Avoid potentially time-consuming syntax errors from eating into your deployment time by running `validate-template`.
- CloudFormation is sometimes slow to update what resources (and new features on old services) a user is able to use. If you need to deploy a resource or feature that isn't supported by the template, CloudFormation allows running a [Lambda](#) function on a stack create or update via [custom resources](#).
- Custom resources make CloudFormation into a truly powerful tool, as you can do all sorts of neat things quite easily like initial configuration of Dynamo tables or S3 buckets, cleaning up old CloudWatch logs, etc.
  - For writing Custom Resources in Javascript, AWS provides a good reference in the [documentation](#).
- CloudFormation offers a visual [template designer](#) that can be useful when getting up to speed with the template language.

- By using [StackSets](#), users can define and deploy an entire production application consisting of multiple stacks (one single CloudFormation template).
- If you're developing a serverless application (i.e., using Lambda, API Gateway) CloudFormation offers a simplified [SAM](#).
- ! Use a restrictive [stack policy](#)! Without one, you can inadvertently delete live production resources, probably causing downtime.
- ! Turn on [termination protection](#) on all of your stacks to avoid costly accidents!
- The CloudFormation [template reference](#) is indispensable when discovering what is and isn't possible in a CloudFormation template.
- [Troposphere](#) is a Python library that makes it much easier to create CloudFormation templates.
  - Currently supports [AWS](#) and [OpenStack](#) resource types.
  - Troposphere attempts to support all resources types that can be described in CloudFormation templates.
  - Built in [error](#) checking.
  - A recommended soft dependency is [awacs](#), which allows you to generate AWS access policy in JSON by writing Python code.
- [stacker](#) is a Python application that makes it easy to define, configure, orchestrate and manage dependencies for multiple stacks across multiple user-defined environments.
- If you are building different stacks with similar layers, it may be useful to build separate templates for each layer using the [AWS::CloudFormation::Stack](#).
- ♦ Avoid hardcoding resource parameters that can potentially change. Use stack parameters as much as you can to define parameter values.
- ♦ Until [2016](#), CloudFormation used only an awkward JSON format that makes both reading and debugging difficult. This typically involved building additional tooling, including converting it to YAML, but now [this is supported directly](#).
- Wherever possible, export relevant [physical IDs](#) from your Stacks by defining [Outputs in your CloudFormation Template](#). These are the actual names assigned to the resources being created. Outputs can be returned from `DescribeStack` API calls, and can be used in other Stacks as part of the [recent addition](#) of [cross-stack references](#). Note that importing outputs in a stack from another stack creates a dependency that is tracked by CloudFormation. You will not be able to delete the stack with the outputs until the stack that depends on it is deleted.
- CloudFormation can be set up to [send SNS notifications](#) upon state changes, enabling programmatic handling of errors, such as when a stack fails to build, or simple email alerts so the appropriate people are informed.
- CloudFormation allows the use of [conditionals](#) when creating a stack.

- One common way to leverage this capability is in support of multi-environment CloudFormation templates – use ‘if-else’ statements on the value of a [parameter passed in](#) (e.g. “env”), environment-specific values for the SecurityGroup IDs, and AMI names can be passed into reusable generic templates.
- **Version control your CloudFormation templates!** In the Cloud, an application is the combination of the code we run and the infrastructure it runs on. By version controlling both, it is easy to roll back to known good states.
- Avoid naming your resources explicitly (e.g. DynamoDB tables). When deploying multiple stacks to the same AWS account, names can come into conflict, potentially slowing down your testing. Prefer using resource references instead.
- For things that shouldn’t ever be deleted, you can set an explicit `DeletionPolicy` on the resource that will prevent it from being deleted even if the CloudFormation stack itself is deleted. This is useful for anything that can maintain expensive state, such as DynamoDB tables, and things that are exposed to the outside world, such as API Gateway APIs.

[Back to top ↑](#)

## CloudFormation Gotchas and Limitations

- ♦ A given CloudFormation stack can end up in a wide variety of states. Error reporting is generally weak, and often observe-tweak-redeploy cycles are needed to get a working template. The internal state machine for [all the various states](#) is opaque.
- ♦ Some cross-region operations are not possible in CloudFormation without using a custom resource, such as [AWS Lambda subscriptions](#).
- ♦ While having hand-made resources live alongside CloudFormation-created resources is inadvisable, it’s sometimes necessary. If possible, leave ALL resource management up to a CloudFormation template and only provide read-only access to the hand-made resources.
- ! Modifications to stack resources made outside CloudFormation can potentially lead to stacks stuck in UPDATE\_IN\_PROGRESS mode. Stacks in this state can be recovered using the [continue-update-rollback command](#). This command can be run in the CLI. The [--resources-to-skip](#) parameter usable in the CLI can be useful if the continue-update-rollback command fails. [Drift Detection](#) can be used to detect outside changes made to stack.
- ♦ CloudFormation is useful but complex and with a variety of pain points. Many companies find alternate solutions to use it, but only with significant additional tooling.
- ♦ CloudFormation can be very slow, especially for items like CloudFront distributions and Route53 CNAME entries.
- ♦ It’s hard to assemble good CloudFormation configurations from existing state. AWS does [offer a trick to do this](#).

- CloudFormer also hasn't been updated in ages (as of Oct 2017), doesn't support templating many new services, and can't define even existing services that have since been updated. For example, Dynamo tables defined through CloudFormer don't support TTL definitions or auto-scaling configuration. There is a third-party version of the tool with more supported features.
- ◆ Many users don't use CloudFormation at all because of its limitations, or because they find other solutions provide better ways to accomplish the same goals, such as local scripts (Boto, Bash, Ansible, etc.) you manage yourself than Docker-based solutions ([Convoy](#), etc.).
- ◆ Deploying large stacks (i.e., many resources) can be problematic due to unintuitive API limits. For instance, the `CreateDeployment` API has a default limit of [3 requests per minute](#) as of 1/12/2018. This limit is readily exceeded when creating many CloudFormation stacks. Creating CW alarms is another commonly seen limit ( `PutMetricAlarm` , 3 tps as of 1/12/2018) when creating many autoscaling policies for DynamoDB. One way to work around this limit is to include CloudFormation artifacts that artificially chain resource creation.
- ◆ Creating/deleting stacks can be a little less clean than ideal. Some resources will leave behind traces in your AWS account after deletion. E.g., Lambda will leave behind CloudWatch log groups that never expire.

## VPCs, Network Security, and Security Groups

---

### VPC Basics

- [Homepage](#) · [User guide](#) · [FAQ](#) · [Security groups](#) · [Pricing](#)
- VPC (Virtual Private Cloud) is the virtualized networking layer of your AWS systems.
- Most AWS users should have a basic understanding of VPC concepts, but few need to get into all the details. VPCs can be trivial or extremely complex, depending on the extent of your network and security needs.
- All modern AWS accounts (those created [after 2013-12-04](#)) are "EC2-VPC" accounts that support VPCs, and all instances run in a VPC. Older accounts may still be using "EC2-Classic" mode. Some features don't work without VPCs, so you probably want to switch to VPC mode.

[Back to top ↑](#)

### VPC and Network Security Tips

- **! Security groups** are your first line of defense for your servers. Be extremely restrictive of what ports are open connections. In general, if you use CLBs, ALBs or other load balancing, the only ports that need to be open to incoming traffic are port 22 and whatever port your application uses. Security groups access policy is 'deny by default'.
- **Port hygiene:** A good habit is to pick unique ports within an unusual range for each different kind of production service. Your web frontend might use 3010, your backend services 3020 and 3021, and your Postgres instances the usual 5432. Using fine-grained security groups for each set of servers. This makes you disciplined about listing out your services, better documented, and easier to audit. For example, should you accidentally have an extra Apache server running on the default port 80 on a backend service, it will be easily detectable by security groups.
- **Migrating from Classic:** For migrating from older EC2-Classic deployments to modern EC2-VPC setup, [this article](#)
  - You can [migrate Elastic IPs between EC2-Classic and EC2-VPC](#).
- For basic AWS use, one default VPC may be sufficient. But as you scale up, you should consider mapping out network boundaries thoroughly. A good overview of best practices is [here](#).
- Consider controlling access to your private AWS resources through a [VPN](#).
  - You get better visibility into and control of connection and connection attempts.
  - You expose a smaller surface area for attack compared to exposing separate (potentially authenticated) services directly to the internet.
    - e.g. A bug in the YAML parser used by the Ruby on Rails admin site is much less serious when the admin interface is accessed through a VPN.
  - Another common pattern (especially as deployments get larger, security or regulatory requirements get more stringent) is to provide a [bastion host](#) behind a VPN through which all SSH connections need to transit.
  - For a cheap VPN to access private AWS resources, consider using a point-to-site software VPN such as [OpenConnect](#) installed using the [official AMI](#), though you are limited to 2 concurrent users on the free license, or it can be run using the [openvpn](#) package on Linux. The Linux package allows for unlimited concurrent users but the installation is less straightforward. [OpenVPN installer script](#) can help you install it and add client keys easily.

- ◆ Consider using other security groups as sources for security group rules instead of using CIDRs — that way, a security group and only hosts in that security group are allowed access. This is a much more dynamic and secure group rules.
- **VPC Flow Logs** allow you to monitor the network traffic to, from, and within your VPC. Logs are stored in CloudWatch Logs and can be used for security monitoring (with third party tools), performance evaluation, and forensic investigation.
  - See the [VPC Flow Logs User Guide](#) for basic information.
  - See the [flowlogs-reader](#) CLI tool and Python library to retrieve and work with VPC Flow Logs.
- **IPv6 is available in VPC**. Along with this announcement came the introduction of the [Egress-Only Internet Gateways](#). One can now choose to use NAT Gateways to enable egress-only traffic for their VPC in IPv4, one can use an Egress-Only Internet Gateway for the same purpose in IPv6.
- Amazon provides an IPv6 CIDR block for your VPC at your request - at present you cannot implement your own CIDR block to own one already.
- New and existing VPCs can both use IPv6. Existing VPCs will need to be configured to have an IPv6 CIDR block assigned as new VPCs do.

[Back to top](#) 

## PrivateLink

-  [Homepage](#) · [User Guide](#) · [Pricing](#)
- One of the uses for Private link is [Interface VPC Endpoints](#) which deploys an ENI into your VPC and subnets which allows your services to interact with AWS API's as if they were accessible locally in your VPC without having to go out to the internet.
- Another use case would be to expose a service of your own to other accounts in AWS through a [VPC Endpoint Service](#).

[Back to top](#) 

## VPC and Network Security Gotchas and Limitations

- ◆ VPCs are tied to one Region in one Account. Subnets are tied to one VPC and limited to one Availability Zone per Region.

- ♦ Security groups are tied to one VPC. If you are utilizing infrastructure in multiple VPCs you should make sure configuration/deployment tools take that into account.
- ♦ [VPC Endpoints](#) are currently only available for S3 and DynamoDB. If you have a security requirement to lock down your VPC you may want to use [DNS filtering](#) to control outbound traffic to other services.
- ! Be careful when choosing your VPC IP CIDR block: If you are going to need to make use of [ClassicLink](#), make sure range [doesn't overlap](#) with that of EC2 Classic.
- ! If you are going to peer VPCs, carefully consider the cost of [data transfer between VPCs](#), since for some workloads it can be prohibitively expensive.
- ! New RDS instances require a [subnet group](#) within your VPC. If you're using the [default VPC](#) this isn't a concern, as it will automatically create a subnet group for each availability zone in your region. However, if you're creating your own VPC and plan on using RDS, make sure to create subnets within the VPC to act as the subnet group.
- ! If you delete the default VPC, you can [recreate it via the CLI or the console](#).
- ! Be careful with VPC VPN credentials! If lost or compromised, the VPN endpoint must be deleted and recreated. See [Replacing Compromised Credentials](#).
- ! Security Groups and Route Tables apply entries separately for IPv4 and IPv6, so one must ensure they add entries accordingly.
- 📈 Managed NAT gateways are a convenient alternative to manually managing [NAT instances](#), but they do come with trade-offs. Consider [alternatives](#) if you're transferring many terabytes from private subnets to the internet. If you transfer traffic from EC2 instances in private subnets to S3, avoid the [NAT gateway data processing charge](#) by setting up a Gateway endpoint and route the traffic to/from S3 through the VPC endpoints instead of going through the NAT gateways.

## KMS

---

### KMS Basics

- 📱 [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- KMS (Key Management Service) is a secure service for creating, storing and auditing usage of cryptographic keys.
- **Service integration:** KMS [integrates with other AWS services](#): EBS, Elastic Transcoder, EMR, Redshift, RDS, SES, S3, Lambda, and Amazon Workspaces.

- **Encryption APIs:** The [Encrypt](#) and [Decrypt API](#) allow you to encrypt and decrypt data on the KMS service side, not key contents.
- **Data keys:** The [GenerateDataKey](#) API generates a new key off of a master key. The data key contents are exposed to encrypt and decrypt any size of data in your application layer. KMS does not store, manage or track data keys this in your application.
- ♦ **Auditing:** Turn on CloudTrail to audit all KMS API events.
- **Access:** Use [key policies](#) and [IAM policies](#) to grant different levels of KMS access. For example, you create an IAM user to encrypt and decrypt with a specific key.

[Back to top](#) 

## KMS Tips

- ♦ It's very common for companies to manage keys completely via home-grown mechanisms, but it's far preferable as KMS from the beginning, as it encourages more secure design and improves policies and processes around management.
- A good motivation and overview is in [this AWS presentation](#).
- The cryptographic details are in [this AWS whitepaper](#).
- [This blog from Convoy](#) demonstrates why and how to use KMS for encryption at rest.

[Back to top](#) 

## KMS Gotchas and Limitations

- ♦ The Encrypt API only works with < 4KB of data. Larger data requires generating and managing a [data key](#) in your application.
- ♦ KMS audit events are not available in the [CloudTrail Lookup Events API](#). You need to look find them in the raw CloudTrail saves in S3.
- ♦ In order to encrypt a multi-part upload to S3, the KMS Key Policy needs to allow "kms:Decrypt" and "kms:GenerateDataKey" to "kms:Encrypt", otherwise the upload will fail with an "AccessDenied" error.
- ♦ KMS keys are region specific — they are stored and can only be used in the region in which they are created. They cannot be moved to other regions.

- ♦ KMS keys have a key policy that must grant access to something to manage the key. If you don't grant anything creation, then you have to reach out to support to have the key policy reset [Reduce the Risk of the Key Becoming](#)
- ♦ If you use a key policy to grant access to IAM roles or users and then delete the user/role, recreating the user permission to the key again.

## CloudFront

---

### CloudFront Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- CloudFront is AWS' [content delivery network \(CDN\)](#).
- Its primary use is improving latency for end users through accessing cacheable content by hosting it at [over 60 geographies](#).

[Back to top](#) 

### CloudFront Alternatives and Lock-in

-  CDNs are [a highly fragmented market](#). CloudFront has grown to be a leader, but there are many alternatives for specific needs.

[Back to top](#) 

### CloudFront Tips

-  IPv6 is [supported](#). This is a configurable setting, and is enabled by default on new CloudFront distributions. It also allows for use of WAF with CloudFront.
-  HTTP/2 is [now supported!](#) Clients [must support TLS 1.2 and SNI](#).
- While the most common use is for users to browse and download content (GET or HEAD methods) requests, CloudFront also supports [uploading](#) data (POST, PUT, DELETE, OPTIONS, and PATCH).
  - You must enable this by specifying the [allowed HTTP methods](#) when you create the distribution.
  - Interestingly, the cost of accepting (uploaded) data is [usually less](#) than for sending (downloaded) data.

- In its basic version, CloudFront [supports SSL](#) via the [SNI extension to TLS](#), which is supported by all modern web browsers. If you need to support older browsers, you need to pay a few hundred dollars a month for dedicated IPs.
  - Consider invalidation needs carefully. CloudFront [does support invalidation](#) of objects from edge locations, but it takes many minutes to propagate to edge locations, and costs \$0.005 per request after the first 1000 requests (CloudFront does not support this better.)
- Everyone should use TLS nowadays if possible. [Ilya Grigorik's table](#) offers a good summary of features regarding the use of TLS with CloudFront.
- An alternative to invalidation that is often easier to manage, and instant, is to configure the distribution to [cache based on query strings](#), then append unique query strings with versions onto assets that are updated frequently.
- For good web performance, it is recommended to [enable compression](#) on CloudFront distributions if the origin server does not already compress its responses.

[Back to top](#)

## CloudFront Gotchas and Limitations

- If using S3 as a backing store, remember that the endpoints for website hosting and for general S3 are different. "bucketname.s3.amazonaws.com" is a standard S3 serving endpoint, but to have redirect and error page support, you must use the website hosting endpoint listed for that bucket, e.g. "bucketname.s3-website-us-east-1.amazonaws.com" (or the equivalent for other regions).
- By default, CloudFront will not forward HTTP Host: headers through to your origin servers. This can be problematic if you run multiple sites switched with host headers. You can [enable host header forwarding](#) in the default cache behavior.
- 4096-bit SSL certificates: CloudFront does not support 4096-bit SSL certificates as of late 2016. If you are using a 4096-bit certificate, you'll need to make sure it's 2048 bits. See [ongoing discussion](#).
- Although connections from clients to CloudFront edge servers can make use of IPv6, [connections to the origin servers are limited to IPv4](#).

## DirectConnect

---

### DirectConnect Basics

-  [Homepage](#) · [User guide](#) · [FAQ](#) · [Pricing](#)
- Direct Connect is a private, dedicated connection from your network(s) to AWS.

[Back to top](#) 

## DirectConnect Tips

- If your data center has a [partnering relationship](#) with AWS, setup is streamlined.
- Use for more consistent predictable network performance guarantees (1 Gbps or 10 Gbps per link).
- Use to peer your colocation, corporate, or physical datacenter network with your VPC(s).
  - Example: Extend corporate LDAP and/or Kerberos to EC2 instances running in a VPC.
  - Example: Make services that are hosted outside of AWS for financial, regulatory, or legacy reasons callable from your VPC.

## Redshift

---

### Redshift Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- Redshift is AWS' managed [data warehouse](#) solution, which is massively parallel, scalable, and columnar. It is very fast using [ParAccel](#) technology and exposes [Postgres](#)-compatible interfaces.

[Back to top](#) 

### Redshift Alternatives and Lock-in

-   Whatever data warehouse you select, your business will likely be locked in for a long time. Also (and not coincidentally), the data warehouse market is highly fragmented. Selecting a data warehouse is a choice to be made carefully, with research into the [market landscape](#) and what [business intelligence](#) tools you'll be using.

[Back to top](#) 

## Redshift Tips

- Although Redshift is mostly Postgres-compatible, its SQL dialect and performance profile are different.
- Redshift supports only [12 primitive data types](#). ([List of unsupported Postgres types](#))
- It has a leader node and computation nodes (the leader node distributes queries to the computation ones). Not [be executed only on the lead node](#).
- ◆ Make sure to create a new [cluster parameter group](#) and option group for your database since the default parameters allow dynamic configuration changes.
- Major third-party BI tools support Redshift integration (see [Quora](#)).
- [Top 10 Performance Tuning Techniques for Amazon Redshift](#) provides an excellent list of performance tuning techniques.
- [Amazon Redshift Utils](#) contains useful utilities, scripts and views to simplify Redshift ops.
- [VACUUM](#) regularly following a significant number of deletes or updates to reclaim space and improve query performance.
- Avoid performing blanket [VACUUM](#) or [ANALYZE](#) operations at a cluster level. The checks on each table to determine whether an ANALYZE action needs to be taken is wasteful. Only perform ANALYZE and VACUUM commands on the objects using [Analyze & Vacuum Schema Utility](#) to perform this work. The SQL to determine whether a table needs to be VACUUMed can be found in the [Schema Utility README](#) if you wish to create your own maintenance process.
- Redshift provides various [column compression](#) options to optimize the stored data size. AWS strongly encourages using [compression](#) at the [COPY](#) stage, when Redshift uses a sample of the data being ingested to analyze the column. However, automatic compression can only be applied to an empty table with no data. Therefore, make sure the sample size is large enough to provide Redshift with a representative sample of the data (the default sample size is 100,000 rows).
- Redshift uses columnar storage, hence it does not have indexing capabilities. You can, however, use [distribution](#) strategies to improve performance. Redshift has two types of sort keys: compounding sort key and interleaved sort key.
- A compound sort key is made up of all columns listed in the sort key definition. It is most useful when you have many columns and want to sort by a subset of them, using the prefix of the sortkey.
- An interleaved sort key on the other hand gives equal weight to each column or a subset of columns in the sort key. This allows you to specify ahead of time which column(s) you want to choose for sorting and filtering, this is a much better choice than the default random sort key.
- ◆ ⏱ **Distribution strategies:** Since data in Redshift is physically distributed among nodes, choosing the right distribution style is crucial for adequate query performance. There are three possible distribution style settings – [RANDOM](#), [KEY](#) or [ALL](#). Use KEY to colocate join key columns for tables which are joined in queries. Use ALL to place the data in all cluster nodes.

## Redshift Gotchas and Limitations

- ! ⚡ While Redshift can handle heavy queries well, it does not scale horizontally, i.e. does not handle multiple concurrent queries well. Therefore, if you expect a high parallel load, consider replicating or (if possible) sharding your data across multiple clusters.
- ♦ The leader node, which manages communications with client programs and all communication with compute nodes, is a single point of failure.
- ⚡ Although most Redshift queries parallelize well at the compute node level, certain stages are executed on the leader node, which can become the bottleneck.
- ♦ Redshift data commit transactions are very expensive and serialized at the cluster level. Therefore, consider grouping multiple small COPY/INSERT/UPDATE commands into a single transaction whenever possible.
- ♦ Redshift does not support multi-AZ deployments. Building multi-AZ clusters is not trivial. [Here](#) is an example.
- ♦ Beware of storing multiple small tables in Redshift. The way Redshift tables are laid out on disk makes it impractical to store many small tables. The total space required to store a table (in MB) is nodes \* slices/node \* columns. For example, on a 16 node cluster an empty table with 100 columns will occupy 640MB on disk.
- ⚡ Query performance degrades significantly during data ingestion. [WLM \(Workload Management\)](#) tweaks help mitigate this. If you need consistent read performance, consider having replica clusters (at the extra cost) and swap them during ingest.
- ! Never resize a live cluster. The resize operation can take hours depending on the dataset size. In rare cases, the resize fails and the cluster becomes unresponsive. You'll end up having a non-functional cluster. The safer approach is to create a new cluster from a snapshot of the old one and shut down the old one.
- ♦ Redshift has **reserved keywords** that are not present in Postgres (see full list [here](#)). Watch out for DELTA ([here](#)).
- ♦ Redshift does not support many Postgres functions, most notably several date/time-related and aggregation functions ([here](#)).
- ♦ Uniqueness, primary key, and foreign key constraints on Redshift tables are informational only and are not enforced by the query optimizer to generate query plans. NOT NULL column constraints are enforced. See [here](#) for more information on defining constraints.
- ♦ Compression on sort key [can result in significant performance impact](#). So if your Redshift queries involving sort keys might want to consider removing compression on a sort key.

- ♦ [Choosing a sort key](#) is very important since you can not change a table's sort key after it is created. If you need to change the distribution key of a table, you need to create a new table with the new key and move your data into it with a query like "create table new\_table select \* from old\_table".
- ! When moving data with a query that looks like "insert into x select from y", you need to have twice as much memory as the table "y" takes up on the cluster's disks. Redshift first copies the data to disk and then to the new table. [Here](#) is a good article for big tables.

## EMR

---

### EMR Basics

-  [Homepage](#) · [Release guide](#) · [FAQ](#) · [Pricing](#)
- EMR (which used to stand for Elastic Map Reduce, but not anymore, since it now extends beyond map-reduce) is a managed deployment of [Hadoop](#), [HBase](#) and [Spark](#). It reduces the management burden of setting up and maintaining yourself.

[Back to top](#) 

### EMR Alternatives and Lock-in

- 🚧 Most of EMR is based on open source technology that you can in principle deploy yourself. However, the job submission tooling is AWS-specific. Migrating from EMR to your own clusters is possible but not always trivial.

[Back to top](#) 

### EMR Tips

- EMR relies on many versions of Hadoop and other supporting software. Be sure to check [which versions are in use](#).
- 🎯 Off-the-shelf EMR and Hadoop can have significant overhead when compared with efficient processing on a cluster. If your data is small and performance matters, you may wish to consider alternatives, as [this post](#) illustrates.
- Python programmers may want to take a look at Yelp's [mrjob](#).

- It takes time to tune performance of EMR jobs, which is why third-party services such as [Qubole's data service](#) always to improve performance or reduce costs.

[Back to top ↑](#)

## EMR Gotchas and Limitations

-  **EMR costs** can pile up quickly since it involves lots of instances, efficiency can be poor depending on cluster choice of workload, and accidents like hung jobs are costly. See the [section on EC2 cost management](#), especially instances. [This blog post](#) has additional tips, but was written prior to the shift to per-second billing.
-  Beware of "double-dipping". With EMR, you pay for the EC2 capacity and the service fees. In addition, EMR storage means you pay for the storage and **PUT requests** at [S3 standard rates](#). While the log files tend to be relatively small depending on the size, generates thousands of log files that can quickly add up to thousands of dollars on the bill. [Aggregation](#) is not available on EMR.

## Kinesis Streams

---

### Kinesis Streams Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- Kinesis Streams (which used to be only called Kinesis, before Kinesis Firehose and Kinesis Analytics were launched) allows you to ingest high-throughput data streams for immediate or delayed processing by other AWS services.
- Kinesis Streams' subcomponents are called **shards**. Each shard provides 1MB/s of write capacity and 2MB/s of read capacity, with a maximum of 5 reads per second. A stream can have its shards programmatically increased or decreased based on a variety of factors.
- All records entered into a Kinesis Stream are assigned a unique sequence number as they are captured. The records are ordered by this number, so any time-ordering is preserved.
- [This page](#) summarizes key terms and concepts for Kinesis Streams.

[Back to top ↑](#)

## Kinesis Streams Alternatives and Lock-in

- 📡 Kinesis is most closely compared to [Apache Kafka](#), an open-source data ingestion solution. It is possible to host Kinesis on [EC2 instances](#) (or any other VPS), however you are responsible for managing and maintaining both Zookeeper brokers in a highly available configuration. Confluent has a good blog post with their recommendations on how to link to several other blogs they have written on the subject.
- 💡 Kinesis uses very AWS-specific APIs, so you should be aware of the potential future costs of migrating away from it.
- An application that efficiently uses Kinesis Streams will scale the number of shards up and down based on the rate of data. (Note there is no direct equivalent to this with Apache Kafka.)

[Back to top ↑](#)

## Kinesis Streams Tips

- The [KCL](#) (Kinesis Client Library) provides a skeleton interface for Java, Node, Python, Ruby and .NET programs to interact with a Kinesis Stream. In order to start consuming data from a Stream, you only need to provide a config file to point to the Stream, and functions for initialising the consumer, processing the records, and shutting down the consumer with grace.
  - The KCL uses a DynamoDB table to keep track of which records have been processed by the KCL. This ensures that each record is processed "at least once". It is up to the developer to ensure that the program can handle doubly-processed records.
  - The KCL also uses DynamoDB to keep track of other KCL "workers". It automatically shares the available Kinesis workers as equally as possible.

[Back to top ↑](#)

## Kinesis Streams Gotchas and Limitations

- 🔍 Kinesis Streams' shards each only permit [5 reads per second](#). If you are evenly distributing data across many shards, the total aggregate throughput for the Stream will remain at 5 reads per second on aggregate, as each consuming application will need to check for new records. This puts a hard limit on the number of different consuming applications possible per Stream for a given latency.
  - For example, if you have 5 consuming applications reading data from one Stream with any number of shards, the latency will be at least one second, as each of the 5 consumers will need to poll *each shard* every second, reaching a maximum of 5 reads per second per shard.

- [This blog post](#) further discusses the performance and limitations of Kinesis in production.
-  **Kinesis Streams are not included in the free tier.** Make sure if you do any experimentation with it on a personal account to turn off the stream or it may run up unexpected costs (~\$11 per shard-month.)

## Kinesis Firehose

---

[Back to top](#) 

### Kinesis Firehose Gotchas and Limitations

-  When delivering from Firehose to Elasticsearch, the JSON document cannot contain an “\_id” property. Firehose will ignore those documents and won’t log any error.

## Device Farm

---

### Device Farm Basics

-  [Homepage](#) · [Developer guide](#) · [FAQ](#) · [Pricing](#)
- **Device Farm** is an AWS service that enables mobile app testing on real devices.
- Supports iOS and Android (including Kindle Fire) devices, as well as the mobile web.
- Supports remote device access in order to allow for interactive testing/debugging.

[Back to top](#) 

### Device Farm Tips

- [AWS Mobile blog](#) contains several examples of Device Farm usage for testing.
- Device Farm offers a free trial for users who want to evaluate their service.
- Device Farm offers two pricing models: Paying **per device minute** is useful for small usage levels or for situations where usage is low. **Unmetered plans** are useful in situations where active usage is expected from the beginning.

- To minimize waiting time for device availability, one approach is to create several device pools with different devices and choose one of the unused device pools on every run.

[Back to top](#) 

## Device Farm Gotchas and Limitations

- ! Devices don't have a SIM card and therefore can't be used for testing SIM card-related features.
- ♦ Device Farm supports testing for most popular languages/frameworks, but not for all. An actual list of supported languages is presented on [this page](#).
- ♦ The API and CLI for Device Farm is quite a low level and may require developing additional tools or scripts on top of it.
- ♦ AWS provide several tools and plugins for Device Farm, however, it doesn't cover all cases or platforms. It may be necessary to use specific tools or plugins to support specific requirements.
- ! In general, Device Farm doesn't have Android devices from Chinese companies like Huawei, Meizu, Lenovo, etc. A list of supported devices located [here](#).
- ♦ Device availability is uneven. It depends on several factors including device popularity. Usually, more modern devices are in higher demand, thus the waiting time for them will be higher compared to relatively old devices.

## Mobile Hub

---

### Mobile Hub Basics

-  [Homepage](#) · [User guide](#) · [FAQ](#) · [Pricing](#)
- **Mobile Hub** orchestrates multiple services to create an AWS backend for mobile and web applications.
- Each *project* in Mobile Hub has one *backend* made up of configurable features, plus one or more *applications*.
- Features include Analytics, Cloud Logic, Conversational Bots, Hosting and Streaming, NoSQL Database, User Data Store, and Device Farm. Each feature uses one or two services to deliver a chunk of functionality.
- Services used include [API Gateway](#), [CloudFront](#), Cognito, [Device Farm](#), [DynamoDB](#), [Lambda](#), Lex, Pinpoint and [S3](#).
- Application SDKs exist for Android (Java), iOS (Swift), Web (JS) and React Native (JS). There is also a CLI for JavaS

## Mobile Hub Tips

- The Mobile Hub [console](#) has starter kits and tutorials for various app platforms.
- The CLI allows local development of Lambda code (JS by default) with `awsmobile {pull|push}` commands, to sync and back again.
- Mobile Hub itself is free, but each of the services has its own pricing model.

## Mobile Hub Gotchas and Limitations

- ♦ The Cloud API feature allows importing an existing Lambda function instead of defining a new one, but there are some gotchas with the CLI. Check the GitHub [issues](#).
- ! Mobile Hub uses CloudFormation under the covers, and gets confused when a service is changed outside of the console.

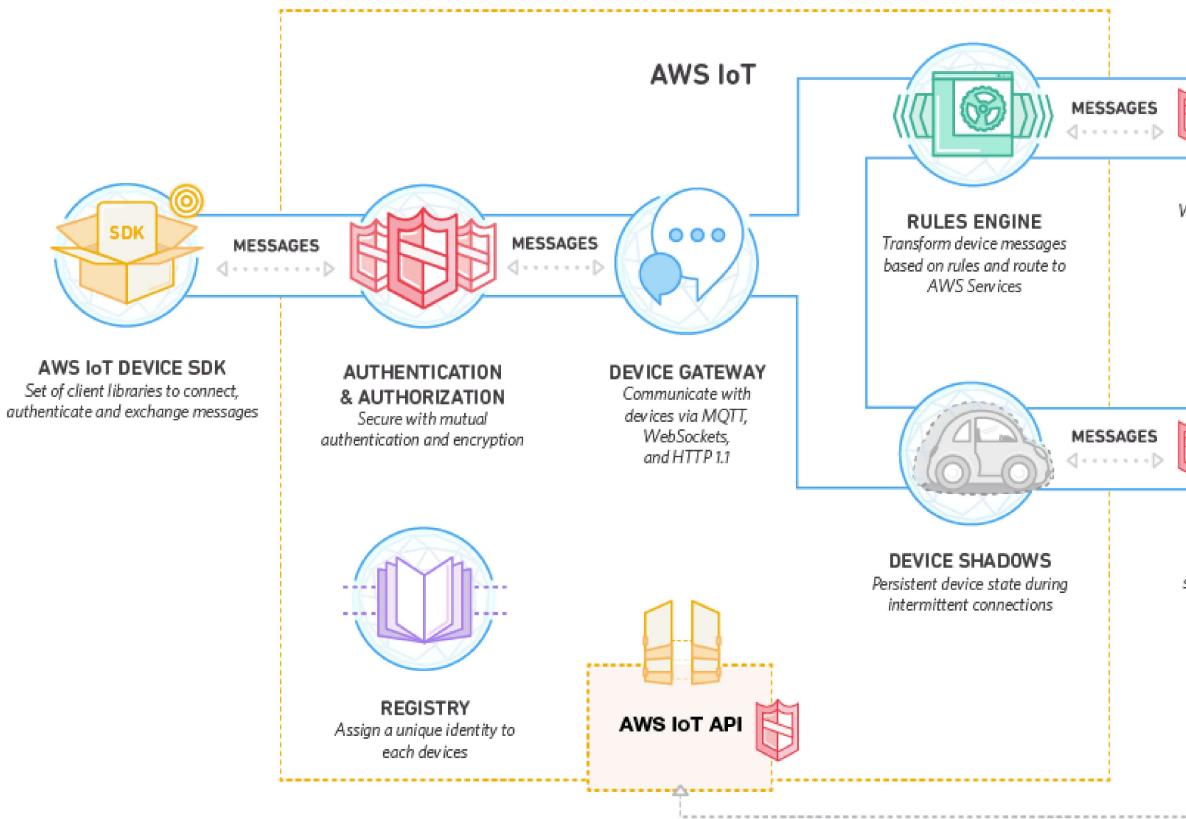
# IoT

---

## IoT Basics

-  [Homepage](#) · [User guide](#) · [FAQ](#) · [Pricing](#)
- IoT is a platform for allowing clients such as IoT devices or software applications ([examples](#)) to communicate with each other.
- Clients are also called **devices** (or **things**) and include a wide variety of device types. Roughly there are three categories:
  - **Device Gateway:**
    - Send messages only: For example, the [AWS IoT Button](#) on an [eddystone beacon](#).
    - Send, receive, and process messages: For example, a simple processing board, such as a [Raspberry Pi](#) ([quick-start](#)) device, such as the [Echo or Echo Dot](#). These are designed to work with the [Alexa skills kit](#), a programmable interface for building Alexa skills.
  - **Cloud Services:** AWS provides several services for managing and processing IoT data, such as [AWS IoT Core](#), [AWS IoT Device农](#), and [AWS IoT Analytics](#).
  - **Edge Devices:** These are physical devices that run locally and handle data processing and storage without needing to connect to the cloud. Examples include [AWS Greengrass](#) and [AWS IoT Greengrass](#).
- AWS has a useful [quick-start](#) (using the Console) and a [slide presentation](#) on core topics.

- IoT terms:
  - AWS [IoT Things](#) (metadata for devices in a [registry](#)) and can store device state in a JSON document, which is immutable. Device metadata can also be stored in [IoT Thing Types](#). This aids in device metadata management by allowing description and configuration for more than one device. Note that IoT Thing Types can be deprecated, but are immutable.
  - AWS [IoT Certificates](#) (device authentication) are the logical association of a unique certificate to the logical device. This association can be done in the Console. In addition, the public key of the certificate must be copied to the device to cover the authentication of devices to a particular AWS Device Gateway (or message broker). You can associate those certificates with an IoT device or you can [register your own CA \(Certificate Authority\) with AWS](#), generate your own certificates, and associate those certificates with your devices via the AWS Console or cli.
  - AWS [IoT Policies](#) (device/topic authorization) are JSON files that are associated to one or more AWS IoT core things and allow associated devices to publish and/or subscribe to messages from one or more MQTT topics.
  - AWS [IoT Rules](#) are SQL-like queries which allows for reuse of some or all device message data, as described in [summarizes design patterns with for IoT Rules](#).
  - Shown below is a [diagram](#) which summarizes the flow of messages between the AWS IoT services:



[Back to top ↑](#)

## IoT Greengrass

- [Homepage](#)

-  **Greengrass** is a software platform that extends AWS IoT capabilities allowing Lambda functions to be run directly on the device. It also enables IoT devices to be able to securely communicate on a local network without having to connect to the cloud.
  - Greengrass includes a local pub/sub message manager that can buffer messages if connectivity is lost so that when connectivity is restored, messages to the cloud are preserved. Locally deployed Lambda functions can be triggered by local events, such as sensor data or other sources.
  - Greengrass includes secure authentication and authorization of devices within the local network and also between the local network and the AWS cloud. It also provides secure, over-the-air software updates of Lambda functions.
- Greengrass core software includes a message manager object, Lambda runtime, local copy service for IoT Thing deployment agent to manage Greengrass group configuration.
- **Greengrass groups** are containers for selected IoT devices settings, subscriptions and associated Lambda functions. A device is either a Greengrass core or an IoT device which will be connected to that particular Greengrass core.
- The Greengrass core SDK enables Lambda functions to interact with the AWS Greengrass core on which they run, send MQTT messages, interact with the local Thing Shadows service, or invoke other deployed Lambda functions.
- The AWS Greengrass Core SDK only supports sending MQTT messages with QoS = 0.
- Shown below is a [diagram](#) which shows the architecture of AWS IoT Greengrass services:



[Back to top ↑](#)

## IoT Alternatives and Lock-in

- AWS, Microsoft and Google have all introduced IoT-specific sets of cloud services since late 2015. AWS was first, releasing their set of IoT services to [general availability](#) in Dec 2015. Microsoft released their set of IoT services for Azure in [Feb 2016](#). Google has released their IoT services [Android Things](#) and [Weave](#).
- Issues of lock-in center around your devices — [protocols](#) (for example MQTT, AMQP), message formats (such as JSON), security (certificates).

[Back to top ↑](#)

## IoT Tips

- **Getting started with Buttons:** One way to start is to use an [AWS IoT Button](#). AWS provides a number of code samples for the AWS IoT Button, you can use the AWS IoT console, click the “connect AWS IoT button” link and you’ll be taken to the AWS Lambda function configuration page. There you fill out your button’s serial number to associate it with a Lambda. (As of this writing, AWS IoT buttons are only available in the US.)
- **Connections and protocols:** It is important to understand the details of about the devices you wish to connect to, including how you will secure the device connections, the device protocols, and more. Cloud vendors differ significantly in the common IoT protocols, such as MQTT, AMQP, XMPP. AWS IoT supports **secure MQTT**, **WebSockets** and **HTTPS**.
- **Support for device security** via certificate processing is a key differentiator in this space. In August 2016, AWS announced [registrations](#) for IoT devices to their services.
- **Combining with other services:** It’s common to use other AWS services, such as AWS Lambda, Kinesis and DynamoDB, with no means required. Sample IoT application reference architectures are in this [screencast](#).
- **Testing tools:**
  - To get started, AWS includes a lightweight MQTT client in the AWS IoT console. Here you can create and test device messages to and from various MQTT topics.
  - When testing locally, if using MQTT, it may be helpful to download and use the open source [Mosquitto broker](#) to interact with devices and/or device simulators
  - Use this [MQTT load simulator](#) to test device message load throughout your IoT solution.

[Back to top ↑](#)

## IoT Gotchas and Limitations

- **◆ IoT protocols:** It is important to verify the exact type of support for your particular IoT device message protocol. The most commonly used IoT protocol is [MQTT](#). Within MQTT there are [three possible levels of QoS in MQTT](#). AWS IoT supports QoS 0 (forget and forget, or at most once) and QoS 1(at least once, or includes confirmation), but *not* QoS 2 (exactly once, requires acknowledgement). This is important in understanding how much code you’ll need to write for your particular application message requirements. See the [presentation about the nuances of connecting](#).
- **◆ The ecosystems to match IAM users or roles to IoT policies** and their associated authorized AWS IoT devices. In many cases, coding to enforce your security requirements is common.

- ! A common mistake is to misunderstand the importance of IoT device security. It is imperative to associate each device with its own certificate (public key). You can generate your own certificates and upload them to AWS, or you can use AWS generated certificates. It's best to read and understand AWS's own guidance on this [topic](#).
- ♦ There is only one AWS IoT Gateway (endpoint) per AWS account. For production scenarios, you'll probably need multiple accounts in order to separate device traffic for development, test and production. It's interesting to note that the AWS IoT service supports configuration of multiple endpoints, so that a single Azure account can be used with separate pub/sub endpoints for development, testing and production
- ♦ **Limits:** Be aware of [limits](#), including device message size, type, frequency, and number of AWS IoT rules.

[Back to top ↑](#)

## IoT Code Samples

- [Simple Beer Service](#) is a surprisingly useful code example using AWS IoT, Lambda, etc.
- [IoT-elf](#) offers clean Python sample using the AWS IoT SDK.
- [IoT Button projects](#) on Hackster include many different code samples for projects.
- [5 IoT code examples](#): a device simulator, MQTT sample, just in time registration, truck simulator, prediction data processing.
- [AWS Alexa trivia voice example](#) is a quick-start using Alexa voice capability and Lambda.
- Some Raspberry Pi examples include the [Beacon project](#), [Danbo](#), and [GoPiGo](#).

## SES

---

### SES Basics

-  [Homepage](#) · [Documentation](#) · [FAQ](#) · [Pricing](#)
- SES (or Simple Email Service) is a service that exposes SMTP endpoints for your application to directly integrate with email providers.

[Back to top ↑](#)

### SES Tips

- **Bounce Handling:** Make sure you handle this early enough. Your ability to send emails can be removed if SES bounces too many emails.
- **Credentials:** Many developers get confused between [SES credentials](#) and AWS API keys. Make sure to enter [AWS API keys](#) using the SMTP APIs.

[Back to top](#) 

## SES Gotchas and Limitations

- **Internet Access:** SES SMTP endpoints are on the Internet and will not be accessible from a location without Internet access (such as a private subnet without NAT gateway route in the routing table). In such a case, set up an SMTP relay instance in your VPC and access and configure your application to send emails to this SMTP relay instance rather than SES. The relay should support [TLS](#) and [TLSv1.2](#). **!** If you are using a proxy instead of a NAT, confirm that your proxy service supports [TLS](#) and [TLSv1.2](#).

## Certificate Manager

### Certificate Manager Basics

-  [Homepage](#) · [User guide](#) · [FAQ](#) · [Pricing](#)
- Use the **Certificate Manager** to manage SSL/TLS certificates in other AWS services.
- Supports importing existing certificates as well as issuing new ones.
- Provides Domain Validated (DV) certificates. [Validation](#) can be done in two ways. The first (and recommended) way is to validate the domain which lives within Route 53 and the user has access, the necessary record can be added in the console via a single click. The second way is to validate the domain which requires more user interaction, and is done by sending an email to 3 contact addresses in WHOIS and 5 contacts in the domain, for each domain name present in the request.
- ACM will attempt to automatically [renew](#) a certificate issued by Amazon. It will first attempt to connect to the domain to verify that the certificate used by the domain is the same with the certificate that it intends to renew. Failing that, it will attempt to renew the certificate using previously for validation. Failing that, ACM will attempt manual validation by sending emails to all domains in the certificate.

[Back to top](#) 

## Certificate Manager Alternatives and Lock-in

- Certificates issued by the Certificate Manager can't be used outside of the services that support it. Imported certificates can still be used elsewhere.

[Back to top](#)

## Certificate Manager Tips

- Supported services: Managed [Load Balancers](#), [CloudFront](#), [API Gateway](#) and [Elastic Beanstalk](#).
- During the domain validation process, if DNS validation is unsuccessful Certificate Manager will send an email to the administrative contact specified in the domain's WHOIS record and up to five common administrative addresses. Some anti-spam filters might block these emails because of this. You should check the spam folder of your email if you don't receive a confirmation email.
- Setting up a certificate for a test domain you don't have email set up on? You can now use DNS validation instead.
- Remember when requesting a wildcard domain that the request will not be valid for the level just below the wildcard. For example, if you request a certificate for \*.bar.example.com, it would not be valid for foo.bar.example.com but not bar.example.com. Likewise it would also not be valid for www.bar.foo.example.com. This applies to each of these domains to the certificate request.

[Back to top](#)

## Certificate Manager Gotchas and Limitations

- In order to use **Certificate Manager** for CloudFront distributions, the certificate must be issued or imported from the same AWS Region.
- Certificates used with Elastic Load Balancers must be issued in the same region as the load balancer. Certificates cannot be copied between regions, as of July 2017. If a domain uses load balancers present in multiple regions, a different certificate must be requested for each region.
- IoT has its [own way](#) of setting up certificates.
- By default the maximum number of domains per certificate is 10. You can get this limit increased to a maximum of 42 with AWS support. Note for every different domain you have on the requested cert, you'll need to press accept on an individual basis. For example if you request a cert with 42 different domains or sub domains, you'll need to press accept on 42 different domains.

- ◆ If you request a limit increase to AWS support for this, they will respond to you asking to confirm this. By body of your initial request: "I acknowledge at the moment, there is no method to add or remove a name Instead, you must request a new certificate with the revised namelist and you must then re-approve all certificate, even if they'd been previously approved."
- ◆ There is no way at the moment to add or remove a domain to/from an existing certificate. You must request approve it from each of the domains requested.

## WAF

---

### WAF Basics

- [Homepage](#) · [Documentation](#) · [FAQ](#) · [Pricing](#)
- WAF (Web Application Firewall) is used in conjunction with the CloudFront and ALB services to inspect and block on user-configurable conditions.
- HTTPS and HTTP requests are supported with this service.
- WAF's strength is in detecting malicious activity based on pattern-matching inputs for attacks such as SQL injection.
- WAF supports inspection of requests [received through both IPv6 and IPv4](#).

[Back to top](#) 

### WAF Tips

- Getting a WAF API call history can be done through CloudTrail. This is enabled through the CloudTrail console.
- It's also possible to get [full logs of all the web requests inspected](#)

[Back to top](#) 

### WAF Gotchas and Limitations

- As of May 2019, AWS WAF is available on Amazon CloudFront and in 12 commercial AWS regions: US East (N. Virginia), US West (Oregon), US West (N. California), EU (Ireland), EU (Frankfurt), EU (London), EU (Stockholm), Asia Pacific (Tokyo), Asia Pacific (Singapore), Asia Pacific (Sydney), Asia Pacific (Seoul), and South America (Sao Paulo).

# OpsWorks

---

## OpsWorks Basics

-  [Homepage](#) · [Documentation](#) · [FAQ](#) · Pricing: [Stacks](#), [Chef Automate](#), [Puppet Enterprise](#)
- OpsWorks is a configuration management service that uses [Chef](#) or [Puppet](#) configuration management. It is broad services:
  - [OpsWorks Stacks](#): The service lets you configure and launch stacks specific to your application's needs, and application deployments. Chef runs can be performed manually via the Execute Cookbooks command, other part of lifecycle events.
    - OpsWorks Stacks differs from standard configuration management services in that it also allows you to infrastructure and application automation (such as creating Amazon EC2 instances and deploying application cookbooks).
  - [OpsWorks for Chef Automate](#): This service launches a dedicated Chef Automate server in your account, which nodes, upload cookbook code, and configure systems. Automated patching, backups, OS updates, and minor patches are provided as part of the service. An AWS API is provided for associating/disassociating nodes. Chef runs using the [chef-client cookbook](#).
  - [OpsWorks for Puppet Enterprise](#): This service launches a dedicated Puppet Master in your account, which can nodes, upload modules, and configure systems. Automated patching, backups, OS updates, and minor Puppet upgrades are provided as part of the service. An AWS API is provided for associating/disassociating nodes. By default, the Puppet agent automatically runs every 30 minutes on associated nodes.
- OpsWorks for Chef Automate and OpsWorks for Puppet Enterprise are strictly designed for configuration management provision infrastructure outside the Chef Server/Puppet Master that is created in our account.
- All three OpsWorks services support managing both Amazon EC2 and on-premises infrastructure, however the interfaces differ slightly.
  - OpsWorks Stacks allows you to register instances and install the OpsWorks Agent to connect to your stack.
  - OpsWorks for Chef Automate and OpsWorks for Puppet Enterprise allow you to associate new or existing instances using the `opsworks-cm:AssociateNode` API action or the vendor-supported method for associating nodes to Chef Automate or Puppet Enterprise.

Enterprise.

- Although OpsWorks will let you work with common Chef recipes or Puppet modules when creating your stacks, you will require familiarity with Chef or Puppet syntax. Chef/Puppet code is not supported as part of AWS Support.
- As of December 2016, OpsWorks Stacks supports Chef versions [12, 11.10.4, 11.4.4 and 0.9.15.5](#).
- As of December 2016, OpsWorks for Chef Automate uses [Chef Server version 12.11.1](#). This is the current stable version.
- [Berkshelf](#) can be used with Chef stacks of version 11.10 and later for managing cookbooks and their respective dependencies. For Chef 12.x stacks, Berkshelf must be installed by the stack administrator.
- Running your own Chef environment may be an alternative to consider - some considerations are listed [in this blog post](#).

[Back to top](#) 

## OpsWorks Alternatives and Lock-in

- Major competitors in Configuration Management include:
  - [Chef](#)
  - [Puppet](#)
  - [Ansible](#).

[Back to top](#) 

## OpsWorks Tips

- OpsWorks Stacks and OpsWorks for Chef Automate use Chef cookbooks for configuration. Chef provides free training materials, best practices, etc. at <https://learn.chef.io>.
- OpsWorks for Puppet Enterprise uses Puppet manifests for configuration. Puppet provides a very useful learning resource at <https://learn.puppet.com/>.

[Back to top](#) 

## OpsWorks Gotchas and Limitations

- OpsWorks Stacks is not available in the following regions:

- Montreal
- GovCloud
- Beijing
- OpsWorks for Chef Automate and OpsWorks for Puppet Enterprise are not available in the following regions:
  - Montreal
  - Sao Paulo
  - GovCloud
  - London
  - Paris
  - Seoul
  - Mumbai

## Batch

---

### Batch Basics

-  [Homepage](#) · [Documentation](#) · [FAQ](#) · [Pricing](#)
- **AWS Batch** is a service that offers an environment to run batch computing jobs. The service dynamically provisions resources needed by the jobs based on their resource requirements, and can scale up to hundreds of thousands of concurrent tasks.
- These batch workloads have access to all other AWS services and features.
- AWS Batch, coupled with [spot instances](#) can help run the jobs when appropriate capacity is available, providing cost savings.
- The batch workloads are built as a [Docker](#) Image. These images can then be pushed to the [EC2 Container Registry](#) for reuse across multiple workloads.
- A [Job Definition](#) has the workload's Docker Image URI, and also lets the users specify the environment details like container properties, environment variables, parameters, retry strategy, and the job's IAM role.
- The [Compute Environments](#) are EC2 clusters that provide the runtime for the batch workloads to execute in.

- AWS Batch provides managed, as well as unmanaged compute environments. The Managed Environments are provided by AWS, while the Unmanaged Environments are managed by the customers.
- The Job Definitions are submitted to [Job Queue\(s\)](#) for execution. Each queue has a priority, and has at least one job definition associated with it.
- AWS Batch uses [ECS](#) to execute the containerized jobs.

[Back to top](#) 

## Batch Tips

- AWS Batch supports prioritization of jobs via the Job Queue Priority. Higher the number - higher the priority.
- AWS Batch supports launching the Compute Environment into specific VPC and subnets.
- A Compute Environment is same as an [ECS Cluster](#).
- There is no additional cost for AWS Batch. You only pay the cost associated with the AWS Services being used - based on the resources consumed by the batch jobs.
- Associate [IAM Roles and policies](#) with the Compute Environment to enable the containers access to other AWS services.
- ◆ Use Unmanaged Compute Environments if you need specialized resources like Dedicated Hosts, or [EFS](#).

## SQS

---

### SQS Basics

-  [Homepage](#) · [Documentation](#) · [FAQ](#) · [Pricing](#)
- SQS is a highly scalable, fully managed message queuing service from AWS.
- SQS supports the pull model, where the producers *queue* the messages, and the consumers pull messages off the queue.
- SQS provides a message visibility timeout, during which the message being processed will not be delivered to other consumers. If the consumer does not delete the message after processing, the message becomes available to other consumers upon the visibility timeout. This parameter is called VisibilityTimeout.
- Each message can have up to 10 custom fields, or attributes.

- SQS allows producers to set up to 15 minutes of delay before the messages are delivered to the consumers. This is done via `DelaySeconds`.
- There are two types of queues supported by SQS -
  - Standard Queues
    - Guarantee **at least once** delivery of the messages.
    - Do not retain the order of delivery of the messages.
  - FIFO Queues
    - Guarantee **only once** delivery of the messages
    - Guarantee the order of the delivery of the messages
- SQS supports fine grained access to various API calls and Queues via IAM policies.
- The messages that fail to process can be put in a dead letter queue.

[Back to top ↑](#)

## SQS Alternatives and Lock-In

- Alternatives to SQS include [Kafka](#), [RabbitMQ](#), [ActiveMQ](#) and others.
- Google Cloud Platform has Pub/Sub, and Azure has Azure Queue Service.
- [SQS vs SNS](#)

[Back to top ↑](#)

## SQS Tips

- SNS can be used in combination of SQS to build a “fan out” mechanism by having an SQS Queue subscribe to the SNS topic.
- SQS supports encryption using AWS KMS.
- Cloudwatch alarms can be created using [various SQS metrics](#) to trigger autoscaling actions and/or notifications.

[Back to top ↑](#)

## SQS Gotchas and Limitations

- ◊ SQS does not have a VPC endpoint (unlike S3 and DynamoDB), so SQS will need to be accessed using public endpoints.
- ◊ FIFO Queues are limited to 300 API calls per second.
- ◊ FIFO Queues cannot subscribe to an SNS topic.
- ◊ Standard Queues can deliver duplicate messages regardless of the visibility window. If only-once delivery is required, use FIFO queues, or build an additional layer to de-dupe the messages.
- ◊ You can send/receive messages in batch, however, there can only be a maximum of 10 messages in a batch.

## SNS

---

### SNS Basics

-  [Homepage](#) · [Documentation](#) · [FAQ](#) · [Pricing](#)
- SNS (Simple Notification Service) is a pub/sub based, highly scalable, and fully managed messaging service that allows you to send notifications to mobile devices, web applications, and other AWS services.
- SNS can push the messages down to the subscribers via [SMS](#), [Email](#), [SQS](#), and [HTTP/S](#) transport protocols.
- Producers publish messages to a SNS Topics, which can have many subscribers.
- Each subscription has an associated [protocol](#), which is used to notify the subscriber.
- A copy of the message is sent to each subscriber using the associated protocol.
- SNS can also [invoke lambda functions](#).

[Back to top](#) 

### SNS Alternatives and Lock-In

- Popular alternatives to SNS are [Kafka](#), [Notification Hubs](#) on Azure, and [Pub/Sub](#) on Google Cloud.
- **SNS vs SQS:**
  - Both SNS and SQS are highly scalable, fully managed messaging services provided by AWS.
  - SQS supports a *pull* model, while SNS supports a *push* model. Consumers have to pull messages from an SQS queue, while SNS pushes the message from an SNS Topic.

- An SQS message is intended to be processed by only one subscriber, while SNS topics can have many subscribers.
- After processing, the SQS message is deleted from the queue by the subscriber to avoid being re-processed.
- An SNS message is *pushed* to all subscribers of the topic at the same time, and is not available for deletion.
- SNS supports multiple transport protocols of delivery of the messages to the subscribers, while SQS subscribers receive messages off the queue over HTTPS.

[Back to top ↑](#)

## SNS Tips

- [Fan-out](#) architecture can be achieved by having multiple subscribers for a topic. This is particularly useful when publishing to multiple, isolated systems.
- SNS topics can be used to power [webhooks](#) with [backoff support](#) to subscribers over HTTP/S.
- [SQS queues](#) can subscribe to SNS topics.
- SNS is used to manage notifications for other AWS services like [Autoscaling Groups](#)' notifications, [CloudWatch Alarms](#), and more.
- SNS is frequently used as "glue" between disparate systems— such as GitHub and AWS services.

[Back to top ↑](#)

## SNS Gotchas and Limitations

- ◊ HTTP/S subscribers of SNS topics need to have public endpoints, as SNS does not support calling private endpoints within a VPC).
- 📁 In a fan-out scenario, [SSE-enabled SQS](#) subscribers of an SNS topic [will not receive](#) the messages sent to the topic.

## High Availability

This section covers tips and information on achieving [high availability](#).

### High Availability Tips

- AWS offers two levels of redundancy, [regions and availability zones \(AZs\)](#).
- When used correctly, regions and zones do allow for high availability. You may want to use non-AWS providers for mitigation (i.e. not tying your company to one vendor), but reliability of AWS across regions is very high.
- **Multiple regions:** Using multiple regions is complex, since it's essentially like managing completely separate infrastructure for business-critical services with the highest levels of redundancy. However, for many applications (like your average e-commerce site), deploying extensive redundancy across regions may be overkill.
- The [High Scalability Blog](#) has a good guide to help you understand when you need to scale an application to multiple regions.
- **◆ Multiple AZs:** Using AZs wisely is the primary tool for high availability!
  - A typical single-region high availability architecture would be to deploy in two or more availability zones, which is shown in [this AWS diagram](#).
  - The bulk of outages in AWS services affect one zone only. There have been rare outages affecting multiple zones, for example, the [great EBS failure of 2011](#) but in general most customers' outages are due to using only a single AZ.
  - Consequently, design your architecture to minimize the impact of AZ outages, especially single-zone outages.
  - Deploy key infrastructure across at least two or three AZs. Replicating a single resource across more than three AZs may not make sense if you have other backup mechanisms in place, like S3 snapshots.
  - A second or third AZ should significantly improve availability, but additional reliability of 4 or more AZs may not be worth the complexity (unless you have other reasons like capacity or Spot market prices).
  -  **Watch out for cross-AZ traffic costs.** This can be an unpleasant surprise in architectures with large volumes of data across AZ boundaries.
  - Deploy instances evenly across all available AZs, so that only a minimal fraction of your capacity is lost in case of an outage.
  - If your architecture has single points of failure, put all of them into a single AZ. This may seem counter-intuitive, but the likelihood of any one SPOF to go down on an outage of a single AZ is low.
- **EBS vs instance storage:** For a number of years, EBSs had a poorer track record for availability than instance storage. Individual instances can be killed and restarted easily, instance storage with sufficient redundancy could give high availability. While EBS availability has improved, and modern instance types (since 2015) are now EBS-only, so this approach, while helpful at one time, is now archaic.
- Be sure to [use and understand CLBs/ALBs](#) appropriately. Many outages are due to not using load balancers, or misconfiguring them.

## High Availability Gotchas and Limitations

- ◆ **AZ naming** differs from one customer account to the next. Your “us-west-1a” is not the same as another customer’s “us-west-1a”. AWS account IDs are assigned to physical AZs randomly per account. This can also be a gotcha if you have multiple AWS accounts. AZ IDs are consistent between accounts, and can be used to reliably align between AWS accounts.
- ◆  **Cross-AZ traffic** is not free. At large scale, the costs add up to a significant amount of money. If possible, keep your traffic within the same AZ as much as possible.

## Billing and Cost Management

---

### Billing and Cost Visibility

- AWS offers a [free tier](#) of service, that allows very limited usage of resources at no cost. For example, a micro instance and storage is available for no charge. Many services are only eligible for the free tier for the first twelve months that they’re available. Some other services offer a free usage tier indefinitely. (If you have an old account but starting fresh, sign up for a new one.) [AWS Activate](#) extends this to tens of thousands of dollars of free credits to startups in [certain funds or accounts](#).
- You can set [billing alerts](#) to be notified of unexpected costs, such as costs exceeding the free tier. You can set these alerts for specific services or across all services.
- AWS offers [Cost Explorer](#), a tool to get better visibility into costs.
- Unfortunately, the AWS console and billing tools are rarely enough to give good visibility into costs. For large accounts, the AWS console can time out or be too slow to use.
- Tools:**
  - ◆ Enable [billing reports](#) and install an open source tool to help manage or monitor AWS resource utilization. (AWS CloudWatch Metrics (written by Netflix) is probably the first one you should try. Check out [docker-ice](#) for a Dockerized version that runs in a container.)
  - ◆ One challenge with Ice is that it doesn’t cover amortized cost of reserved instances.
  - Other tools include [Security Monkey](#) and [Cloud Custodian](#).
  - Use [AWS Simple Monthly Calculator](#) to get an estimate of usage charges for AWS services based on certain assumptions. Monthly charges will be based on your actual usage of AWS services, and may vary from the estimates the calculator provides.

- **Third-party services:** Several companies offer services designed to help you gain insights into expenses or lower costs. Examples include [Cloudability](#), [CloudHealth Technologies](#), and [ParkMyCloud](#). Some of these charge a percentage of your bill, which can be a significant cost.
- AWS's [Trusted Advisor](#) is another service that can help with cost concerns.
- Don't be shy about asking your account manager for guidance in reducing your bill. It's their job to keep you happy!
- **Tagging for cost visibility:** As the infrastructure grows, a key part of managing costs is understanding where the costs are being spent. Tag resources with descriptive tags, such as [tag resources](#), and as complexity grows, group them effectively. If you [set up billing allocation appropriately](#), you can then allocate costs into expenses according to organization, product, individual engineer, or any other way that is helpful.
- If you need to do custom analysis of raw billing data or want to feed it to a third party cost analysis service, [enable the CloudWatch Metrics Insights feature](#).
- Multiple Amazon accounts can be linked for billing purposes using the [Consolidated Billing](#) feature. Large enterprises often have complex billing structures depending on ownership and approval processes.
- Multiple Amazon accounts can be managed centrally using [AWS Organizations](#).

[Back to top ↑](#)

## AWS Data Transfer Costs

- For deployments that involve significant network traffic, a large fraction of AWS expenses are around data transfer costs. The cost of data transfer, within AZs, within regions, between regions, and into and out of AWS and the internet vary significantly based on deployment choices.
- Some of the most common gotchas:
  - ♦ *AZ-to-AZ traffic:* Note EC2 traffic between AZs is effectively the same as between regions. For example, creating a multi-AZ cluster across AZs is helpful for [high availability](#), but can hurt on network costs.
  - ♦ *Using public IPs when not necessary:* If you use an Elastic IP or public IP address of an EC2 instance, you will be charged for data transfer. Even if it is accessed locally within the AZ.
  - ♦ *Managed NAT Gateway data processing:* Managed NAT Gateways are used to let traffic egress from private subnets. They add a cost of approximately 4.5¢ as a data processing fee layered on top of data transfer pricing. Past a certain point, running your own NAT gateway may be more cost effective.

- ◆ Some services do cross-AZ traffic for free: Many AWS services you'd not consider on their own merits offer cross-AZ data transfer. EFS, RDS, MSK, and others are examples of this.
- This figure gives an overview:

#### AWS DATA TRANSFER COSTS

**Numbers are data transfer in \$/GB.**  
**Transaction and hourly prices are not shown. See notes.**

- ∅ Free. Inbound traffic is mostly free —you pay on the way out. Some but not all internal traffic is free.
- ① Direct outbound data starts at \$.09/GB for <10TB, and discounts with volume. First 1GB free.
- ② Region-to-region traffic is \$.02/GB when it exits a region for indicated services except between us-east-1 and us-east-2, where it's \$.01/GB.
- ③ Outbound CloudFront prices are highly variable by geography and regional edge cache and start at \$.085/GB in US/Canada.
- ④ Internal traffic via public or elastic IPs incurs additional fees in both directions.
- ⑤ Cross-AZ EC2 traffic within a region costs as much as region-to-region! ELB-EC2 traffic is free except outbound crossing AZs.
- ⑥ Elastic Load Balancing: Classic LB is priced per GB. Application LB costs are in LCUs, not \$/GB.

