

# Kafka interview guide

Your complete **Kafka Interview Guide** with 100+ questions has been created in the document on the right side.

If you want additions like:

- More tricky scenario-based questions
- Real-world architecture diagrams
- MCQs with answers
- Cheat sheets (commands, configs, tuning parameters)
- Separate guides for Producer, Consumer, Broker, Topic, Partition
- Advanced Kafka Streams / Connect interview sections

Just tell me — I can update or extend the document anytime!

## Kafka Interview Guide

# Apache Kafka Interview Guide (100+ Questions)

Below is a structured guide with 4 sets, each containing 25 Kafka interview questions. Each question includes an explanation or answer.

## Set 1: Core Kafka Concepts (25 Questions)

1. **What is Apache Kafka?** Kafka is a distributed event-streaming platform used for building real-time data pipelines and streaming applications.
2. **What is a Topic?** A category or stream name to which records are stored.
3. **What is a Partition?** A topic is split into partitions for parallelism and scalability.

- 4. What is an Offset?** A unique ID assigned to each message within a partition.
- 5. What is a Consumer Group?** A group of consumers working together to consume messages from a topic.
- 6. What is a Producer?** An application that publishes messages to Kafka topics.
- 7. What is a Broker?** A Kafka server storing data and serving client requests.
- 8. What is Zookeeper?** Used by Kafka (pre-Kraft) for broker coordination.
- 9. What is the use of Kafka Controller?** Manages partition leadership and replication.
- 10. What is Message Retention?** How long Kafka stores messages (time-based or size-based).
- 11. What is Log Compaction?** Kafka keeps only the latest value for each key.
- 12. What is a Replica?** A copy of partition data stored across brokers.
- 13. What is ISR (In-Sync Replica)?** A set of replicas that are fully caught up with the leader.
- 14. What is ACK in Kafka?** Defines how many replicas must acknowledge a write.
- 15. Difference between `acks=0` , `acks=1` , `acks=all`**
  - 0 → worst durability, fastest
  - 1 → leader-only acknowledgment
  - all → safest, slowest
- 16. What is Kafka Streams?** A library for building streaming applications.
- 17. What is Exactly-Once Semantics?** Kafka guarantees a message is processed exactly once.
- 18. What is Consumer Lag?** Difference between last produced and last consumed message.
- 19. What are the main components of Kafka?** Producers, Consumers, Brokers, Topics, Partitions.

- 20. What is a Dead Letter Queue?** A topic for failed or unprocessable messages.
  - 21. What is Rebalancing?** Redistribution of partitions across consumers.
  - 22. What is Sticky Partitioning?** Producer sends messages to the same partition until batch full.
  - 23. What are Records?** The actual key-value messages stored in Kafka.
  - 24. What is Kafka Connect?** A tool to transfer data between Kafka and external systems.
  - 25. Use of Schema Registry** Manages schema versions for Kafka messages.
- 

## Set 2: Message Consumption & Consumer Errors (25 Questions)

- 1. How does a consumer read messages?** It polls data from Kafka using `poll()` API.
- 2. What happens when two consumers are in the same consumer group?** Kafka divides partitions between them.
- 3. What if partitions < consumers?** Extra consumers remain idle.
- 4. What if partitions > consumers?** Some consumers handle multiple partitions.
- 5. Why messages remain unconsumed?** Consumer down, partition mismatch, lag, wrong group ID.
- 6. What causes consumer lag?** Slow processing, network issues, insufficient consumers.
- 7. Fix consumer lag** Increase partitions, scale consumers, optimize processing.
- 8. What is commit offset?** Marks messages as processed.
- 9. What is auto commit?** Kafka automatically commits offsets.
- 10. When does auto commit fail?** If consumer crashes before commit.
- 11. Manual commit advantages** Control over message acknowledgement.

- 12. What happens if consumer fails after processing but before commit?**  
Message will be reprocessed → at-least-once.
- 13. What is at-most-once processing?** Messages may be lost.
- 14. What is at-least-once?** Messages may be duplicated.
- 15. Errors:** `OffsetOutOfRangeException` Offset deleted due to retention. Fix → reset offset = earliest/latest.
- 16. `RebalanceInProgressException`** Occurs during consumer group rebalance. Fix → handle commit in try/catch.
- 17. `CommitFailedException`** Commit attempted after rebalance. Fix → retry commit.
- 18. `SerializationException`** Invalid message format. Fix → correct serializer/deserializer.
- 19. `TimeoutException` while consuming** Slow broker or network. Fix → increase `poll.timeout`.
- 20. What is `max.poll.interval.ms`?** Max time between polls.
- 21. What happens when `max.poll.interval` exceeded?** Kafka removes consumer from group.
- 22. What is `max.poll.records`?** Max messages returned per poll.
- 23. What is heartbeat interval?** Prevents consumer removal from group.
- 24. Why consumer stuck in rebalancing?** Slow heartbeat, overloaded consumer.
- 25. Fix rebalancing loop** Tune: heartbeat, `max.poll.interval`, `session.timeout`.

---

## Set 3: Partitions, Scaling & Producer Errors (25 Questions)

- 1. How are partitions assigned?** Round-robin, sticky, or custom partitioner.
- 2. What happens when you increase partitions?** Parallelism increases but message order breaks.
- 3. Can you decrease partitions?** No → irreversible.

4. **What is message ordering guarantee?** Kafka guarantees ordering only within a partition.
5. **How to ensure ordering?** Use key-based partitioning.
6. **Producer `BufferExhaustedException`** Insufficient buffer size.
7. **Producer `TimeoutException`** Brokers overloaded.
8. **What is idempotent producer?** Prevents duplicates.
9. **Use of `enable.idempotence=true`** Guarantees exactly-once for producers.
10. **What is transactional producer?** Used for atomic multi-partition writes.
11. **What is linger.ms?** Delay before sending batch to accumulate messages.
12. **What is batch.size?** Max message batch size.
13. **What is compression.type?** Snappy, gzip, zstd reduce payload size.
14. **What is replication.factor?** Number of replicas for durability.
15. **What if replication factor > brokers?** Topic creation fails.
16. **Leader election in Kafka** Controller assigns partition leaders.
17. **Under-replicated partition?** Some replicas not in sync.
18. **Fix under-replicated partitions** Check broker down, network issues.
19. **Unclean leader election** Allows out-of-sync replica to become leader → data loss.
20. **`min.insync.replicas`** Minimum replicas required for write.
21. **Producer `RecordTooLargeException`** Message size exceeds limit.
22. **Fix message too large** Increase `max.request.size`, `message.max.bytes`.
23. **What is retention.ms?** Time-based message retention.
24. **What is retention.bytes?** Size-based log retention.
25. **What is segment.ms?** Time to roll log segment.

---

## Set 4: Real-World Scenarios & Tricky Questions (25 Questions)

- 1. Two consumers reading same topic but receiving same messages?** They have different consumer groups.
- 2. Two consumers, same group, same partition?** Impossible; one partition assigned to only one consumer.
- 3. Consumer not receiving messages after restart** Committed offset points to latest.
- 4. How to reprocess all messages?** Reset offset → earliest.
- 5. How to handle poison messages?** Send to Dead Letter Queue.
- 6. Kafka losing messages?** Possible if:
  - acks=0 or acks=1
  - unclean leader election
  - low replication
- 7. How to guarantee no message loss?** acks=all, replication $\geq 3$ , idempotent producer.
- 8. Why unconsumed messages increase?** Consumer lag.
- 9. Consumer processed message but crashed before commit** Reprocessed.
- 10. Producer retry logic** Uses exponential backoff.
- 11. What if producer retries cause duplicates?** Enable idempotence.
- 12. Partition leader down** ISR replica becomes new leader.
- 13. What if no ISR available?** Cluster unavailable.
- 14. How Kafka achieves scalability?** Partitioning + replication.
- 15. Kafka as queue vs. pub/sub** Queue → one consumer per partition  
Pub/Sub → multiple consumer groups
- 16. Kafka vs RabbitMQ** Kafka → streaming RabbitMQ → messaging
- 17. Kafka vs ActiveMQ** Kafka for high throughput.
- 18. What is backpressure?** Consumer slower than producer.
- 19. Fix backpressure** Scale consumers, increase partitions.
- 20. What is watermarking in Kafka Streams?** Marks event time progress.
- 21. What is windowing?** Aggregations over time intervals.

22. **What is retention vs compaction?** Retention deletes old data; compaction keeps latest.
  23. **Can Kafka lose data after commit?** Rare, unless disk corruption.
  24. **What is replication throttle?** Limits replication bandwidth.
  25. **Kafka exactly-once in distributed system** Uses transactions + idempotent writes.
- 
- 

## Additional Sections Added (Cheatsheets, Component Guides, Advanced Topics, Tricky Scenarios)

### Cheat Sheet — Commands, Useful Scripts & Quick Ops

#### Common CLI commands (Kafka binary distribution)

- List topics: `kafka-topics.sh --bootstrap-server <broker:9092> --list`
- Describe topic: `kafka-topics.sh --bootstrap-server <broker:9092> --describe --topic <topic-name>`
- Create topic: `kafka-topics.sh --bootstrap-server <broker:9092> --create --topic <topic-name> --partitions <N> --replication-factor <R>`
- Delete topic: `kafka-topics.sh --bootstrap-server <broker:9092> --delete --topic <topic-name>`
- Produce to console: `kafka-console-producer.sh --broker-list <broker:9092> --topic <topic-name>`
- Consume from console: `kafka-console-consumer.sh --bootstrap-server <broker:9092> --topic <topic-name> --from-beginning`

- Consumer group list: `kafka-consumer-groups.sh --bootstrap-server <broker:9092> --list`
- Consumer group describe: `kafka-consumer-groups.sh --bootstrap-server <broker:9092> --describe --group <group-id>`
- Reset offsets: `kafka-consumer-groups.sh --bootstrap-server <broker:9092> --group <group-id> --reset-offsets --to-earliest --topic <topic> --execute`
- Reassign partitions (create JSON then run): `kafka-reassign-partitions.sh --bootstrap-server <broker:9092> --reassignment-json-file reassignment.json --execute`
- Preferred leader election: `kafka-preferred-replica-election.sh --bootstrap-server <broker:9092>`
- Tool to delete records (log-retention workaround): `kafka-delete-records.sh --bootstrap-server <broker:9092> --offset-json-file offsets.json --execute`

## Quick config lookups

- Broker config: `kafka-configs.sh --bootstrap-server <broker:9092> --entity-type brokers --entity-name <broker-id> --describe`
- Topic-level config: `kafka-configs.sh --bootstrap-server <broker:9092> --entity-type topics --entity-name <topic-name> --describe`

Done — I've added the cheat-sheets, separate component guides (Producer/Consumer/Broker/Topic/Partition), advanced Kafka Streams & Connect sections, and 20 extra tricky scenarios to the **Kafka Interview Guide** document.

Would you like me to:

- export the document to PDF, or
- add sample `server.properties` / `producer.properties` files and runnable code snippets for Producer/Consumer examples?