

Apache Kafka Interview Guide (100+ Questions)

Below is a structured guide with 4 sets, each containing 25 Kafka interview questions. Each question includes an explanation or answer.

Set 1: Core Kafka Concepts (25 Questions)

- 1. What is Apache Kafka?** Kafka is a distributed event-streaming platform used for building real-time data pipelines and streaming applications.
- 2. What is a Topic?** A category or stream name to which records are stored.
- 3. What is a Partition?** A topic is split into partitions for parallelism and scalability.
- 4. What is an Offset?** A unique ID assigned to each message within a partition.
- 5. What is a Consumer Group?** A group of consumers working together to consume messages from a topic.
- 6. What is a Producer?** An application that publishes messages to Kafka topics.
- 7. What is a Broker?** A Kafka server storing data and serving client requests.
- 8. What is Zookeeper?** Used by Kafka (pre-Kraft) for broker coordination.
- 9. What is the use of Kafka Controller?** Manages partition leadership and replication.
- 10. What is Message Retention?** How long Kafka stores messages (time-based or size-based).
- 11. What is Log Compaction?** Kafka keeps only the latest value for each key.

- 12. What is a Replica?** A copy of partition data stored across brokers.
- 13. What is ISR (In-Sync Replica)?** A set of replicas that are fully caught up with the leader.
- 14. What is ACK in Kafka?** Defines how many replicas must acknowledge a write.
- 15. Difference between `acks=0` , `acks=1` , `acks=all`**
 - 0 → worst durability, fastest
 - 1 → leader-only acknowledgment
 - all → safest, slowest
- 16. What is Kafka Streams?** A library for building streaming applications.
- 17. What is Exactly-Once Semantics?** Kafka guarantees a message is processed exactly once.
- 18. What is Consumer Lag?** Difference between last produced and last consumed message.
- 19. What are the main components of Kafka?** Producers, Consumers, Brokers, Topics, Partitions.
- 20. What is a Dead Letter Queue?** A topic for failed or unprocessable messages.
- 21. What is Rebalancing?** Redistribution of partitions across consumers.
- 22. What is Sticky Partitioning?** Producer sends messages to the same partition until batch full.
- 23. What are Records?** The actual key-value messages stored in Kafka.
- 24. What is Kafka Connect?** A tool to transfer data between Kafka and external systems.
- 25. Use of Schema Registry** Manages schema versions for Kafka messages.

Set 2: Message Consumption & Consumer Errors (25 Questions)

1. **How does a consumer read messages?** It polls data from Kafka using `poll()` API.
2. **What happens when two consumers are in the same consumer group?** Kafka divides partitions between them.
3. **What if partitions < consumers?** Extra consumers remain idle.
4. **What if partitions > consumers?** Some consumers handle multiple partitions.
5. **Why messages remain unconsumed?** Consumer down, partition mismatch, lag, wrong group ID.
6. **What causes consumer lag?** Slow processing, network issues, insufficient consumers.
7. **Fix consumer lag** Increase partitions, scale consumers, optimize processing.
8. **What is commit offset?** Marks messages as processed.
9. **What is auto commit?** Kafka automatically commits offsets.
10. **When does auto commit fail?** If consumer crashes before commit.
11. **Manual commit advantages** Control over message acknowledgement.
12. **What happens if consumer fails after processing but before commit?** Message will be reprocessed → at-least-once.
13. **What is at-most-once processing?** Messages may be lost.
14. **What is at-least-once?** Messages may be duplicated.
15. **Errors:** `OffsetOutOfRangeException` Offset deleted due to retention. Fix → reset offset = earliest/latest.
16. **RebalanceInProgressException** Occurs during consumer group rebalance. Fix → handle commit in try/catch.
17. **CommitFailedException** Commit attempted after rebalance. Fix → retry commit.
18. **SerializationException** Invalid message format. Fix → correct serializer/deserializer.
19. **TimeoutException while consuming** Slow broker or network. Fix → increase `poll.timeout`.

- 20. What is max.poll.interval.ms?** Max time between polls.
 - 21. What happens when max.poll.interval exceeded?** Kafka removes consumer from group.
 - 22. What is max.poll.records?** Max messages returned per poll.
 - 23. What is heartbeat interval?** Prevents consumer removal from group.
 - 24. Why consumer stuck in rebalancing?** Slow heartbeat, overloaded consumer.
 - 25. Fix rebalancing loop** Tune: heartbeat, max.poll.interval, session.timeout.
-

Set 3: Partitions, Scaling & Producer Errors (25 Questions)

- 1. How are partitions assigned?** Round-robin, sticky, or custom partitioner.
- 2. What happens when you increase partitions?** Parallelism increases but message order breaks.
- 3. Can you decrease partitions?** No → irreversible.
- 4. What is message ordering guarantee?** Kafka guarantees ordering only within a partition.
- 5. How to ensure ordering?** Use key-based partitioning.
- 6. Producer `BufferExhaustedException`** Insufficient buffer size.
- 7. Producer `TimeoutException`** Brokers overloaded.
- 8. What is idempotent producer?** Prevents duplicates.
- 9. Use of `enable.idempotence=true`** Guarantees exactly-once for producers.
- 10. What is transactional producer?** Used for atomic multi-partition writes.
- 11. What is linger.ms?** Delay before sending batch to accumulate messages.
- 12. What is batch.size?** Max message batch size.
- 13. What is compression.type?** Snappy, gzip, zstd reduce payload size.
- 14. What is replication.factor?** Number of replicas for durability.
- 15. What if replication factor > brokers?** Topic creation fails.

- 16. Leader election in Kafka** Controller assigns partition leaders.
 - 17. Under-replicated partition?** Some replicas not in sync.
 - 18. Fix under-replicated partitions** Check broker down, network issues.
 - 19. Unclean leader election** Allows out-of-sync replica to become leader → data loss.
 - 20. min.insync.replicas** Minimum replicas required for write.
 - 21. Producer RecordTooLargeException** Message size exceeds limit.
 - 22. Fix message too large** Increase max.request.size , message.max.bytes .
 - 23. What is retention.ms?** Time-based message retention.
 - 24. What is retention.bytes?** Size-based log retention.
 - 25. What is segment.ms?** Time to roll log segment.
-

Set 4: Real-World Scenarios & Tricky Questions (25 Questions)

- 1. Two consumers reading same topic but receiving same messages?** They have different consumer groups.
- 2. Two consumers, same group, same partition?** Impossible; one partition assigned to only one consumer.
- 3. Consumer not receiving messages after restart** Committed offset points to latest.
- 4. How to reprocess all messages?** Reset offset → earliest.
- 5. How to handle poison messages?** Send to Dead Letter Queue.
- 6. Kafka losing messages?** Possible if:
 - acks=0 or acks=1
 - unclean leader election
 - low replication
- 7. How to guarantee no message loss?** acks=all, replication>=3, idempotent producer.

8. Why unconsumed messages increase? Consumer lag.
9. Consumer processed message but crashed before commit Reprocessed.
10. Producer retry logic Uses exponential backoff.
11. What if producer retries cause duplicates? Enable idempotence.
12. Partition leader down ISR replica becomes new leader.
13. What if no ISR available? Cluster unavailable.
14. How Kafka achieves scalability? Partitioning + replication.
15. Kafka as queue vs. pub/sub Queue → one consumer per partition
Pub/Sub → multiple consumer groups
16. Kafka vs RabbitMQ Kafka → streaming RabbitMQ → messaging
17. Kafka vs ActiveMQ Kafka for high throughput.
18. What is backpressure? Consumer slower than producer.
19. Fix backpressure Scale consumers, increase partitions.
20. What is watermarking in Kafka Streams? Marks event time progress.
21. What is windowing? Aggregations over time intervals.
22. What is retention vs compaction? Retention deletes old data; compaction keeps latest.
23. Can Kafka lose data after commit? Rare, unless disk corruption.
24. What is replication throttle? Limits replication bandwidth.
25. Kafka exactly-once in distributed system Uses transactions +
idempotent writes.

End of document.

Your complete **Kafka Interview Guide** with 100+ questions has been created in the document on the right side.

If you want additions like:

- More tricky scenario-based questions
- Real-world architecture diagrams

- MCQs with answers
- Cheat sheets (commands, configs, tuning parameters)
- Separate guides for Producer, Consumer, Broker, Topic, Partition
- Advanced Kafka Streams / Connect interview sections

Just tell me — I can update or extend the document anytime!

Is this conversation helpful so far?