

Springboot microservices guide

Below is SET D (Questions 76–100) — FULL EXPANDED LONG ANSWERS, completing your 100 Questions Spring Boot Microservices Interview Guide.

SET D — OPERATIONS, SECURITY, DEPLOYMENTS, LOGGING, CLOUD, CI/CD, SCALING (76–100)

Full Long-Form Answers

76. Explain Zero-Downtime Deployment in microservices.

Zero-downtime deployment ensures the system remains **fully available** during deployments.

Techniques

1 Blue-Green Deployment

You maintain two production environments:

Environment	Purpose
Blue	live traffic
Green	new version deployed

Once green passes health checks → switch traffic to green instantly.

Benefits

- Fast rollback
 - Zero downtime
 - Safe testing in production environment
-

2 Rolling Deployment

Deploy services instance by instance:

1. Stop 1 instance
2. Deploy new version
3. Test
4. Repeat for next instance

- ✓ No full system downtime
 - ✓ Works well in Kubernetes
-

3 Canary Deployment

Release new version to **small % of users**.

Example:

- 5% traffic → new version
- 95% → old version

If errors increase → rollback.

4 Kubernetes deployment strategies

- `RollingUpdate` (default)
 - `Recreate`
 - `BlueGreen` (via ArgoCD or Spinnaker)
 - `Canary` (via Istio)
-

Zero Downtime Requires:

- ✓ health checks
 - ✓ backward compatibility
 - ✓ database migrations with versioning
 - ✓ externalized config
 - ✓ stateless containers
-
-

77. What is API Rate Limiting? How to implement it?

Rate limiting protects services from:

- DDoS
 - brute force
 - accidental high traffic
 - billing explosions
 - bad clients
-

Techniques

1. Token Bucket (most common)
 2. Fixed Window
 3. Sliding Window Log
 4. Leaky Bucket
-

Implement with Spring Cloud Gateway using Redis

yaml

```
spring:  
  cloud:  
    gateway:  
      routes:  
        - id: product-service  
          uri: lb://product-service  
          predicates:  
            - Path=/products/**  
          filters:  
            - name: RequestRateLimiter  
              args:  
                redis-rate-limiter.replenishRate: 10  
                redis-rate-limiter.burstCapacity: 20
```

✓ Each user allowed 10 requests/sec.

Real E-commerce Use Case

- limit add-to-cart API
 - throttle payment attempts
 - restrict search spam
-
-

78. Explain Blue-Green vs Canary Deployments (with scenarios).

Feature	Blue-Green	Canary
Traffic Switch	100% immediate	gradual 1% → 5% → 50%
Rollback	instant	slow but monitored
Ideal For	major releases	experimental changes
Risk Level	low	lowest

Scenario: Payment Service

Blue-Green:

You deploy version 2 in Green
After full testing → switch 100% traffic
Rollback → switch back to Blue

Canary:

You release to 2% of users
Monitor payments
If errors increase → auto rollback

79. What is Horizontal vs Vertical Scaling?

Vertical Scaling (Scale Up)

Increase machine size:

- more CPU
- more RAM
- more storage

✗ expensive

✗ downtime

✓ simpler

Horizontal Scaling (Scale Out)

Add more instances of service.

- ✓ no downtime
 - ✓ cloud-native
 - ✓ auto-scaling support
 - ✓ load balanced
-

E-commerce usage

- inventory search → horizontal
 - Redis cache → vertical
 - Kafka → horizontal partitions
-
-

80. What is Infrastructure as Code (IaC)? Tools? Why needed?

IaC manages infrastructure using **code instead of manual configuration**.

Popular IaC Tools

- Terraform
- AWS CloudFormation
- Azure ARM
- Pulumi

- Ansible
-

Benefits

- ✓ consistent
 - ✓ version controlled
 - ✓ repeatable
 - ✓ automation in CI/CD
 - ✓ minimized human errors
-

Example Terraform Snippet

```
hcl
```

```
resource "aws_instance" "app" {  
    ami = "ami-x"  
    instance_type = "t2.micro"  
}
```

81. Explain how secrets are managed in microservices.

Sensitive data:

- passwords
- API keys
- database credentials
- tokens

Must never be stored in:

- ✗ code
 - ✗ GitHub
 - ✗ config files
-

Secure Methods

1. Vault (HashiCorp)

Centralized secrets manager.

2. Kubernetes Secrets

Encoded + access restricted.

3. AWS Secrets Manager

Auto-rotating secrets.

4. Azure Key Vault

5. Environment Variables

Spring Boot example

yaml

```
spring:  
  datasource:  
    password: ${DB_PASSWORD}
```

Secrets injected at runtime.

82. What is Sidecar Pattern? Example with Istio?

Sidecar pattern runs helper services beside main application container.

Examples

- logging agents
 - Envoy proxy for service mesh
 - monitoring modules
 - security modules
-

Istio Example

Each microservice pod has:

- main container (app)
- sidecar container (Envoy proxy)

Sidecar handles:

- ✓ mTLS
- ✓ tracing
- ✓ routing
- ✓ retry logic

Application remains clean.

83. What is a Service Mesh? Why is it needed?

Service Mesh is a dedicated **infrastructure layer** to manage microservice communication.

Popular Tools

- Istio
 - Linkerd
 - Consul Connect
-

Features

- ✓ traffic routing
 - ✓ mTLS encryption
 - ✓ retries/circuit breakers
 - ✓ observability
 - ✓ load balancing
 - ✓ rate limiting
-

Why needed?

Without service mesh, application code contains:

- retry
- load balancing
- logging
- security

Service mesh moves these responsibilities out of code.

84. Explain mTLS (Mutual TLS) in microservices.

mTLS ensures:

- client verifies server
- server verifies client

Both exchange certificates.

Why Important?

Prevents:

- impersonation
 - MITM attacks
 - unauthorized microservice calls
-

Used by:

- Istio
 - Linkerd
 - AWS Mesh
-
-

85. What is Distributed Configuration? Explain Spring Cloud Config.

Distributed config stores config centrally.

Why?

- consistent values across services
 - dynamic reloading
 - version control
 - no rebuild required for config changes
-

Spring Cloud Config

Stores config in:

- Git
- Vault
- AWS S3
- Database

Usage Example

Config Server:

```
bash

/application.yml
/order-service.yml
/payment-service.yml
```

Clients fetch on startup → or through actuator refresh.

86. What is a Deployment Manifest in Kubernetes?

Deployment manifest is YAML file describing:

- containers
 - replicas
 - ports
 - environment variables
 - rolling updates
-

Example manifest

yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: order-service
spec:
  replicas: 3
  template:
    spec:
      containers:
        - name: order
          image: order:1.0
          ports:
            - containerPort: 8080
```

87. Explain Observability in microservices. What are its pillars?

Observability = ability to understand system internal state from external output.

Three Pillars

1. Logs

application logs using ELK / EFK

2. Metrics

system performance

Prometheus → Grafana

3. Traces

request flows

Jaeger / Zipkin

Why Needed

- ✓ faster debugging
 - ✓ capacity planning
 - ✓ anomaly detection
-
-

88. What is Prometheus? How used in microservices?

Prometheus is **time-series metrics store**.

Used For

- service metrics
 - resource monitoring
 - alerting
-

Pipeline

Spring Boot → Micrometer → Prometheus → Grafana

Example metrics

- request count
 - latency
 - memory usage
 - JVM GC
-
-

89. Explain ELK Stack (Elasticsearch, Logstash, Kibana).

Elasticsearch

stores logs and supports fast search

Logstash

ingests log data
parses formats

Kibana

visual dashboard

Use Case

Search logs using filters:

- log level
 - traceId
 - userId
 - timestamp
-
-

90. What is Canary Monitoring? What metrics matter?

Monitor new release (canary) for:

- error rate
- latency
- CPU/RAM usage
- business metrics (orders placed, payments)

If metrics degrade → rollback automatically.

91. Explain JWT authentication in microservices.

JWT = JSON Web Token

Used for:

- authentication
 - authorization
-

Structure

header.payload.signature

Flow

1. User logs in → Auth Service issues JWT
 2. JWT sent with each request
 3. Gateway validates JWT
 4. Downstream services trust gateway
-

Pros

- ✓ stateless
 - ✓ no DB lookup
 - ✓ fast
-
-

92. How do you secure inter-service communication?

Methods:

- mTLS
 - API keys
 - service accounts
 - OAuth2 machine-to-machine
 - sidecar proxies
 - network policies
-
-

93. What is OAuth2 Client Credentials Flow?

Used for service-to-service communication.

Flow

1. Service A requests token from Authorization Server
 2. Sends clientId + clientSecret
 3. Authorization server issues access token
 4. Service A calls Service B using token
-

Use Case

Payment service calling Fraud detection service.

94. Explain CORS. Why is it needed in microservices?

CORS allows browser clients from one origin to access another origin.

Needed for:

- SPA frontends (React, Angular)
 - separate backend + UI domains
-

Spring Boot example:

java

```
@CrossOrigin(origins = "https://myapp.com")
```

95. How do you implement retries in microservices?

Retries needed for:

- network glitches
 - temporary failures
-

Resilience4j Example

java

```
@Retry(name = "inventoryRetry", maxAttempts = 3)
public Response checkInventory() {
    ...
}
```

Retries with exponential backoff.

96. What is Chaos Engineering? Why is it important?

Chaos engineering intentionally injects failures to test system resilience.

Tools

- Chaos Monkey for Spring Boot
 - Gremlin
 - AWS Fault Injection Simulator
-

Tests

- kill containers
 - add latency
 - drop network
 - kill nodes
-
-

97. Explain container orchestration. Why is Kubernetes preferred?

Container orchestration manages:

- deployments
- scaling
- rollover

- health checks
-

Kubernetes Advantages

- ✓ self-healing
 - ✓ auto-scaling
 - ✓ declarative YAML
 - ✓ service discovery
 - ✓ secrets/config maps
 - ✓ ecosystem support
-
-

98. What is Pod Autoscaling? How does HPA work?

Horizontal Pod Autoscaler scales pods based on metrics.

Metrics Examples

- CPU
 - memory
 - custom metrics
-

Example

```
yaml

apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
spec:
  minReplicas: 2
  maxReplicas: 10
  metrics:
    - type: Resource
      resource:
        name: cpu
    targetAverageUtilization: 50
```

99. What is rolling restart and why do we need it?

Rolling restart updates pods one by one.

- ✓ avoids downtime
 - ✓ ensures continuous availability
 - ✓ supports zero downtime
-

100. What is CI/CD? How do you build pipelines for microservices?

CI/CD automates:

- building
 - testing
 - packaging
 - deploying
-

Pipeline Stages

1. CI (Continuous Integration)

- compile
- unit tests
- integration tests
- static code analysis
- build docker image
- push to registry

2. CD (Continuous Deployment)

- deploy to dev
- deploy to QA
- canary or blue-green prod deployment
- health checks
- rollback logic

Popular Tools

- Jenkins
 - GitHub Actions
 - GitLab CI
 - ArgoCD
 - Spinnaker
 - Tekton
-



SET D Completed! (76–100)

You now have:

- ✓ Set A (1–25)
- ✓ Set B (26–50)
- ✓ Set C (51–75)
- ✓ Set D (76–100)

All 100 questions completed with **detailed long answers**.



NEXT STEP (Ready to Generate PDF)

Now that all 100 detailed answers are completed,

🔥 I can generate a **formatted, printable PDF** containing:

- Full 100 questions
 - Long answers (20–30 pages)
 - Proper headings
 - Page numbers
 - Table of contents
 - Clean layouts
-

✓ Just confirm:

“Yes, generate full PDF with all 100 detailed answers.”

I will prepare the PDF.

