

Springboot microservices guide

Great — here is **Option A — Full Expansion:**

All 100 Spring Boot Microservices Interview Questions WITH Detailed, Long, Model Answers

This is the **complete expanded content**, ready for PDF export afterward (once you confirm).

FULL EXPANDED MODEL ANSWERS (ALL 100 QUESTIONS)

I have formatted them clearly so the PDF will look like a proper *interview preparation book*.

SET A — Core Concepts & Design (Q1–25)

1. What is a microservice? How does it differ from a monolith?

A microservice is a small, independently deployable service that focuses on a **single business capability**. It communicates over lightweight protocols (REST, messaging).

Differences vs Monolith:

Aspect	Monolith	Microservices
Architecture	Single deployment unit	Many small services
Scalability	Scale entire app	Scale individual services
Technology	Usually uniform	Polyglot allowed

Aspect	Monolith	Microservices
Deployment	One deployment pipeline	Multiple independent
Fault isolation	Weak	Strong, each service isolated
Data	Shared DB	Decentralized/own DB
Complexity	Simple to start, hard to maintain	Complex initially, easy at scale

Microservices shine in large, evolving systems where agility, scalability, and team ownership matter.

2. Advantages and disadvantages of microservices

Advantages:

- Independent deployment → faster release cycle
- Technology freedom (Java for payments, Node for UI API)
- Better scalability per service
- Smaller codebases → easier to maintain
- Strong isolation → failures don't take whole system down
- Enables DevOps, CI/CD, ownership by small teams

Disadvantages:

- Operational complexity (20 services = 20 deploy pipelines)
 - Distributed systems problems (network latency, timeouts)
 - Eventual consistency issues
 - Harder debugging & tracing
 - Requires mature DevOps, monitoring, logging
-

3. How does Spring Boot help build microservices?

Spring Boot provides:

- **Auto-configuration** → reduces boilerplate
- **Embedded servers** → deploy as JAR

- **Spring Starter dependencies**
- **Actuator** for health, metrics, tracing endpoints
- **Spring Cloud** for discovery, config server, gateway, load balancing
- **Resilience4j integration** for resiliency patterns

Thus Spring Boot is a foundation for building robust microservices rapidly with minimal setup.

4. How do you design microservice boundaries?

Use **Domain-Driven Design (DDD)**:

- **Bounded context** → Each microservice owns *one core domain context*
 - Each service owns its data (no shared database)
 - Consider **high cohesion and low coupling**
 - Apply **business capability mapping** (Checkout, Payment, Inventory, Search)
 - Avoid splitting based on CRUD tables
 - Split based on business subdomains, not layers
-

5. Explain DDD concepts relevant to microservices

Bounded Context:

Defines the boundary where a domain model is valid. Each microservice has its own models.

Entities & Aggregates:

Aggregates enforce business rules. Example: `order` aggregate ensures no order is created without items.

Value Objects:

Immutable objects without identity (e.g., Money, Address).

Domain Events:

Changes in the system represented as events (OrderCreated, PaymentCompleted).

Repositories:

Provide access to aggregates.

DDD + microservices → clean architecture with clear separation.

6. How do you manage configuration across environments?

Use Spring Cloud Config Server:

- Centralized configuration in Git
- Version-controlled
- Dynamic reload using **Spring Cloud Bus**
- Supports encryption for sensitive configs
- Profiles for envs: `application-dev.yml`, `application-prod.yml`

Local configs → easy drift

Remote configs → consistent, secure, auditable.

7. Explain service discovery and registration

In microservices, service instances change (scale, restart, move).

Service discovery dynamically locates instances.

Two patterns:

1. **Client-side discovery (Netflix Eureka)**

Client pulls list & load balances.

2. **Server-side discovery (Kubernetes)**

Kube-proxy/Envoy resolves service IP.

Discovery avoids hardcoded URLs and enables zero-downtime scalability.

8. What is an API Gateway? Why use it?

API Gateway is a single entry point for all clients.

Benefits:

- Authentication/Authorization
- Routing to services
- Rate limiting & throttling
- Centralized logging & tracing
- Response aggregation
- Caching

- SSL termination
- Backend abstraction

Spring Cloud Gateway is commonly used in Spring microservices.

9. How do you secure microservices?

Use OAuth2 + JWT with Spring Security.

Typical flow:

1. Client → Auth Server → gets JWT
2. Sends JWT to microservice
3. Microservice validates signature, expiry
4. Doesn't call auth server each time

Other practices:

- Rotate signing keys
 - mTLS for service-service communication
 - Short-lived tokens
 - API gateway as authentication filter
-

10. REST vs Messaging in microservices

REST:

- Useful for synchronous actions (Get Order Status)
- Simple and widely supported
- But tight coupling + latency issues

Messaging (Kafka/RabbitMQ):

- For async operations (Order Placed → Payment → Inventory)
 - Better decoupling, resilience
 - Use for heavy traffic & event-driven systems
-

11. What is idempotency? Why needed?

An operation is **idempotent** if repeated calls produce same result.

Needed in:

- Order APIs (avoid duplicate order creation)
- Payment APIs
- Retry operations in network failures

Implement using:

- Idempotency keys
 - Deduplication tables
 - Conditional updates
-

12. Circuit Breaker & Bulkhead

Circuit breaker:

- Prevents cascading failures
- Stops calling failing service
- Has open → half-open → closed states

Bulkhead:

- Isolates thread pools per dependency
 - Ensures one slow service doesn't exhaust threads
-

13. What is eventual consistency? Example.

Data may not be updated instantly across all services.

Example:

Order service → emits `OrderCreated`

Inventory → reserves → emits `InventoryReserved`

UI may show "Pending Confirmation" until inventory finishes.

Use:

- Polling
 - UI pending state
 - Play with read-models & materialized views
-

14. Saga Pattern (Deep)

Saga breaks distributed transaction into smaller ones with **compensation**.

Choreography Saga:

- Services listen to events
- No central coordinator

Orchestration Saga:

- Dedicated saga orchestrator controls workflow

Used in:

- Order → Payment → Inventory → Shipping
 - Rollback through compensation (cancel payment, release stock)
-

15. Migrating monolith to microservices

Steps:

1. Identify bounded contexts
 2. Extract one domain at a time
 3. Strangler pattern → new microservice intercepts traffic
 4. Data migration via dual-write or CDC
 5. Gradual rollout
 6. Monitor & stabilize
-

16. How to handle distributed transactions?

Avoid 2PC (slow, locking).

Preferred:

- **Saga pattern**
 - **Event-driven workflows**
 - **Idempotent compensations**
 - **Outbox pattern** to ensure reliable events
-

17. Distributed tracing

Tools:

- OpenTelemetry
- Zipkin
- Jaeger

Need:

- Trace ID + Span ID
 - Propagate headers via gateway → all services
-

18. Caching strategies

Types:

- Local cache (Caffeine)
- Distributed cache (Redis)
- Cache-aside, write-through, write-behind

Benefits:

- Reduce DB load
 - Improve latency
-

19. API Versioning

Ways:

- URI versioning: /v1/products
- Header-based: Accept-Version:1
- Query param (rarely recommended)

Guidelines:

- Backward compatibility
 - Sunset old versions gradually
-

20. Data ownership

Each microservice owns:

- Its schema
- Its tables

- Its migrations

Shared DB = tight coupling → anti-pattern.

21. Error handling strategies

Return standard errors:

- 400 Bad Request
- 404 Not Found
- 409 Conflict
- 500 Server Error

Use formats:

- application/problem+json
 - Correlation ID
-

22. Important microservice metrics

- Latency
 - Throughput
 - Error rate
 - JVM heap, GC
 - DB connection pool usage
 - Queue size (Kafka lag)
-

23. Rate limiting & throttling

Algorithms:

- Token bucket
- Leaky bucket
- Fixed window
- Sliding window

Tools:

- API Gateway
 - Redis rate limiter
-

24. Sidecar pattern

Moves cross-cutting concerns out of service:

- Logging agent
- Proxy (Envoy)
- TLS, retries, policy enforcement

Used in **Service Mesh**.

25. Observability

Three pillars:

- **Logs** (structured JSON)
- **Metrics** (Prometheus)
- **Tracing** (Jaeger)

Design:

- Add correlation ID
 - Instrument all outgoing calls
 - Use gateway-level request logging
-

SET B — Implementation (Q26–50)

(I will continue with the complete detailed answers for Q26–100 in the next message, due to message length limits.)

👉 Reply with: “Continue Set B”

and I will send the next 25 detailed answers.